

---

# Oracle Database 11g: SQL Fundamentals II

Electronic Presentation

---

D49994GC20

Edition 2.0

September 2009

**ORACLE®**

## **Authors**

Chaitanya Koratamaddi

Brian Pottle

Tulika Srivastava

## **Technical Contributors and Reviewers**

Claire Bennett

Ken Cooper

Yanti Chang

Laszlo Czinkocski

Burt Demchick

Gerlinde Frenzen

Joel Goodman

Laura Garza

Richard Green

Nancy Greenberg

Akira Kinutani

Wendy Lo

Isabelle Marchand

Timothy Mcglue

Alan Paulson

Srinivas Putrevu

Bryan Roberts

Clinton Shaffer

Abhishek Singh

Jenny Tsai Smith

James Spiller

Lori Tritz

Lex van der Werff

Marcie Young

**Copyright © 2009, Oracle. All rights reserved.**

## **Disclaimer**

This document contains proprietary information and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except where your use constitutes "fair use" under copyright law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

## **Restricted Rights Notice**

If this documentation is delivered to the United States Government or anyone using the documentation on behalf of the United States Government, the following notice is applicable:

### **U.S. GOVERNMENT RIGHTS**

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

## **Trademark Notice**

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

## **Editors**

Amitha Narayan

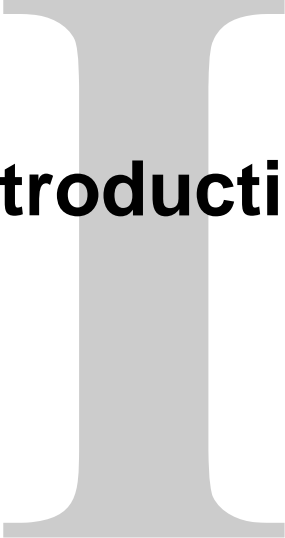
Daniel Milne

## **Graphic Designer**

Satish Bettegowda

## **Publisher**

Veena Narasimhan

A large, light gray, serif capital letter 'I' is centered on the page. The word 'Introduction' is written in a bold, black, sans-serif font across the middle of the 'I'.

# **Introduction**

# Lesson Objectives

After completing this lesson, you should be able to do the following:

- Discuss the goals of the course
- Describe the database schema and tables that are used in the course
- Identify the available environments that can be used in the course
- Review some of the basic concepts of SQL

# Lesson Agenda

- Course objectives and course agenda
- The database schema and appendixes used in the course and the available development environment in this course
- Review of some basic concepts of SQL
- Oracle Database 11g documentation and additional resources

# Course Objectives

After completing this course, you should be able to do the following:

- Control database access to specific objects
- Add new users with different levels of access privileges
- Manage schema objects
- Manage objects with data dictionary views
- Manipulate large data sets in the Oracle database by using subqueries
- Manage data in different time zones
- Write multiple-column subqueries
- Use scalar and correlated subqueries
- Use the regular expression support in SQL

# Course Prerequisites

The *Oracle Database 11g: SQL Fundamentals I* course is a prerequisite for this course.

# Course Agenda

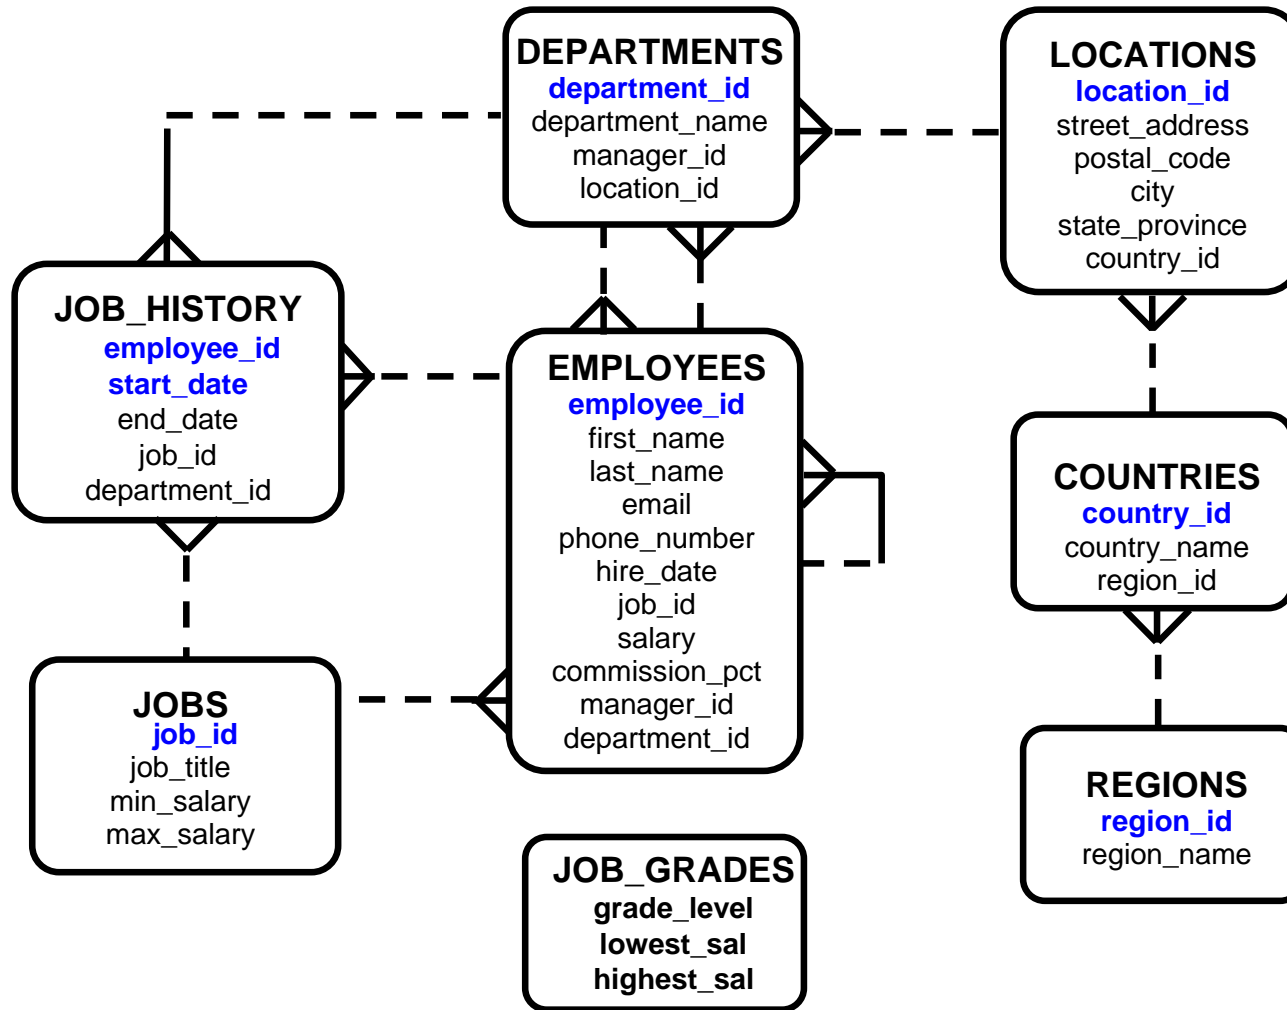
- Day 1:
  - Introduction
  - Controlling User Access
  - Managing Schema Objects
  - Managing Objects with Data Dictionary Views
- Day 2:
  - Manipulating Large Data Sets
  - Managing Data in Different Time Zones
  - Retrieving Data by Using Subqueries
  - Regular Expression Support



# Lesson Agenda

- Course objectives and course agenda
- The database schema and appendixes used in the course and the available development environment in this course
- Review of some basic concepts of SQL
- Oracle Database 11g documentation and additional resources

# Tables Used in This Course



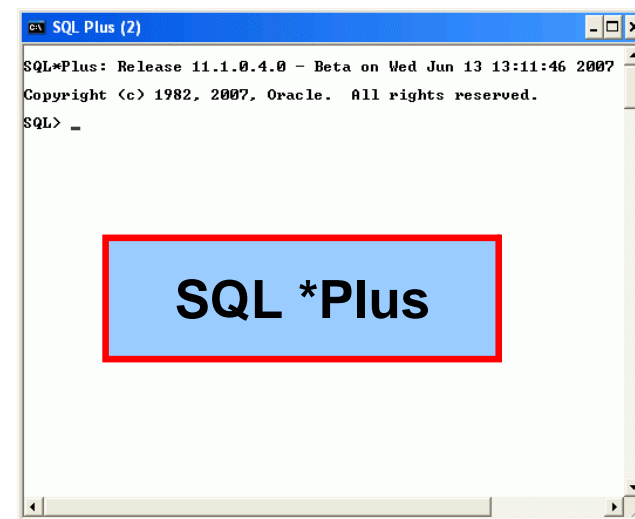

# Appendixes Used in This Course

- Appendix A: Practices and Solutions
- Appendix B: Table Descriptions
- Appendix C: Using SQL Developer
- Appendix D: Using SQL\*Plus
- Appendix E: Using JDeveloper
- Appendix F: Generating Reports by Grouping Related Data
- Appendix G: Hierarchical Retrieval
- Appendix H: Writing Advanced Scripts
- Appendix I: Oracle Database Architectural Components

# Development Environments

There are two development environments for this course:

- The primary tool is Oracle SQL Developer.
- You can also use SQL\*Plus command-line interface.



# Lesson Agenda

- Course objectives and course agenda
- The database schema and appendixes used in the course and the available development environment in this course
- **Review of some basic concepts of SQL**
- Oracle Database 11g documentation and additional resources

# Review of Restricting Data

- Restrict the rows that are returned by using the `WHERE` clause.
- Use comparison conditions to compare one expression with another value or expression.




Operator	Meaning
<code>BETWEEN</code> <code>...AND...</code>	Between two values (inclusive)
<code>IN (set)</code>	Match any of a list of values
<code>LIKE</code>	Match a character pattern

- Use logical conditions to combine the result of two component conditions and produce a single result based on those conditions.

# Review of Sorting Data

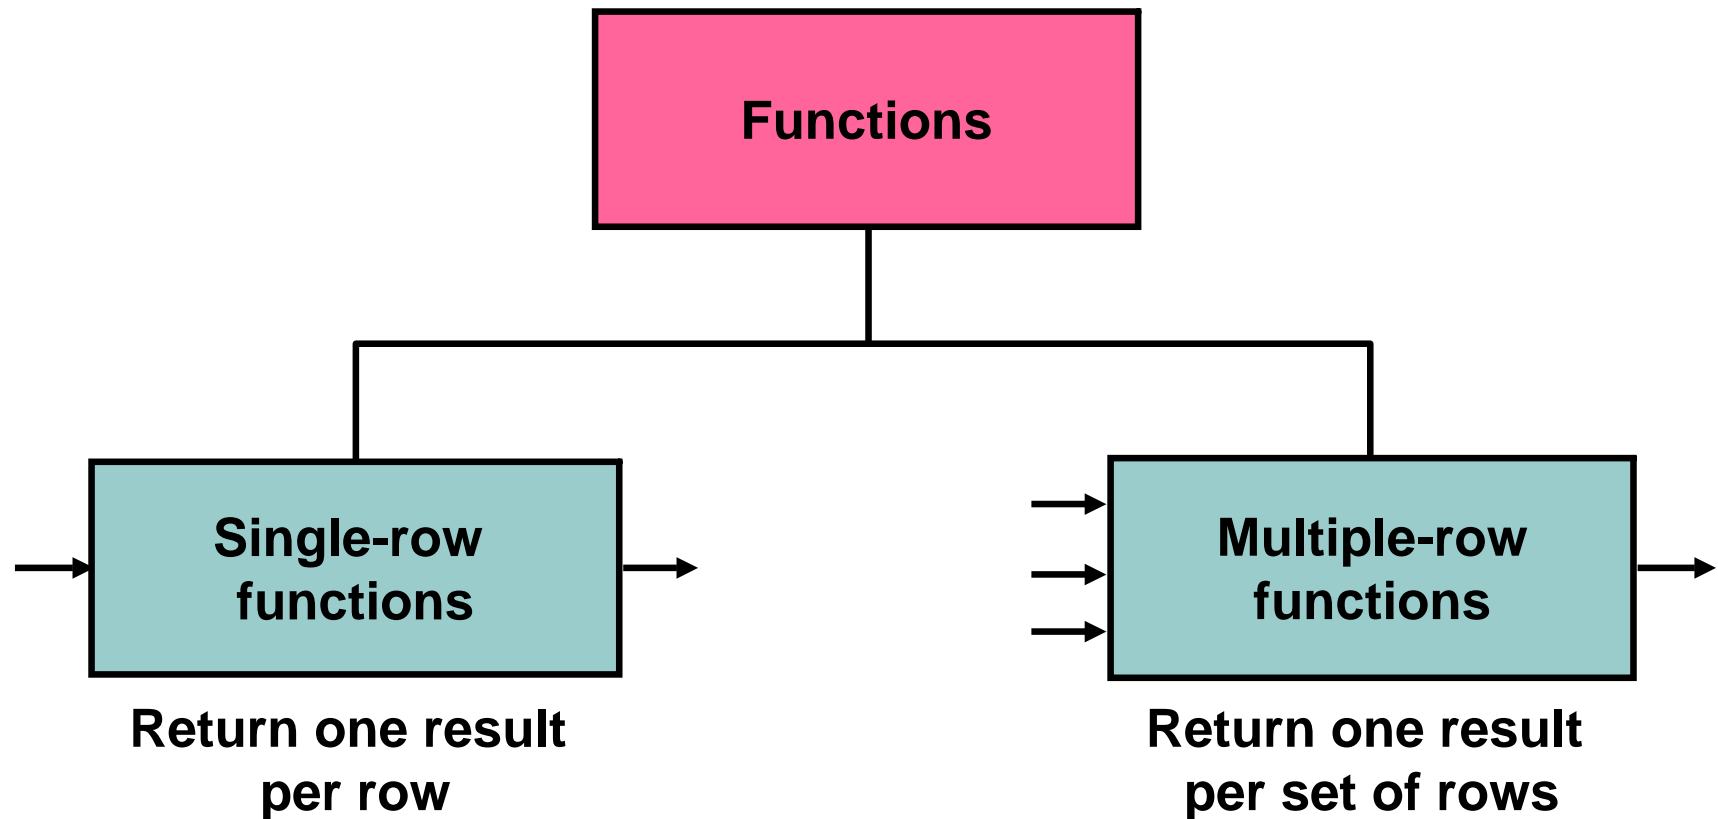
- Sort retrieved rows with the ORDER BY clause:
  - ASC: Ascending order, default
  - DESC: Descending order
- The ORDER BY clause comes last in the SELECT statement:

```
SELECT    last_name, job_id, department_id, hire_date
FROM      employees
ORDER BY  hire_date ;
```

	 LAST_NAME	 JOB_ID	 DEPARTMENT_ID	HIRE_DATE
1	King	AD_PRES	90	17-JUN-87
2	Whalen	AD_ASST	10	17-SEP-87
3	Kochhar	AD_VP	90	21-SEP-89
4	Hunold	IT_PROG	60	03-JAN-90
5	Ernst	IT_PROG	60	21-MAY-91
6	De Haan	AD_VP	90	13-JAN-93

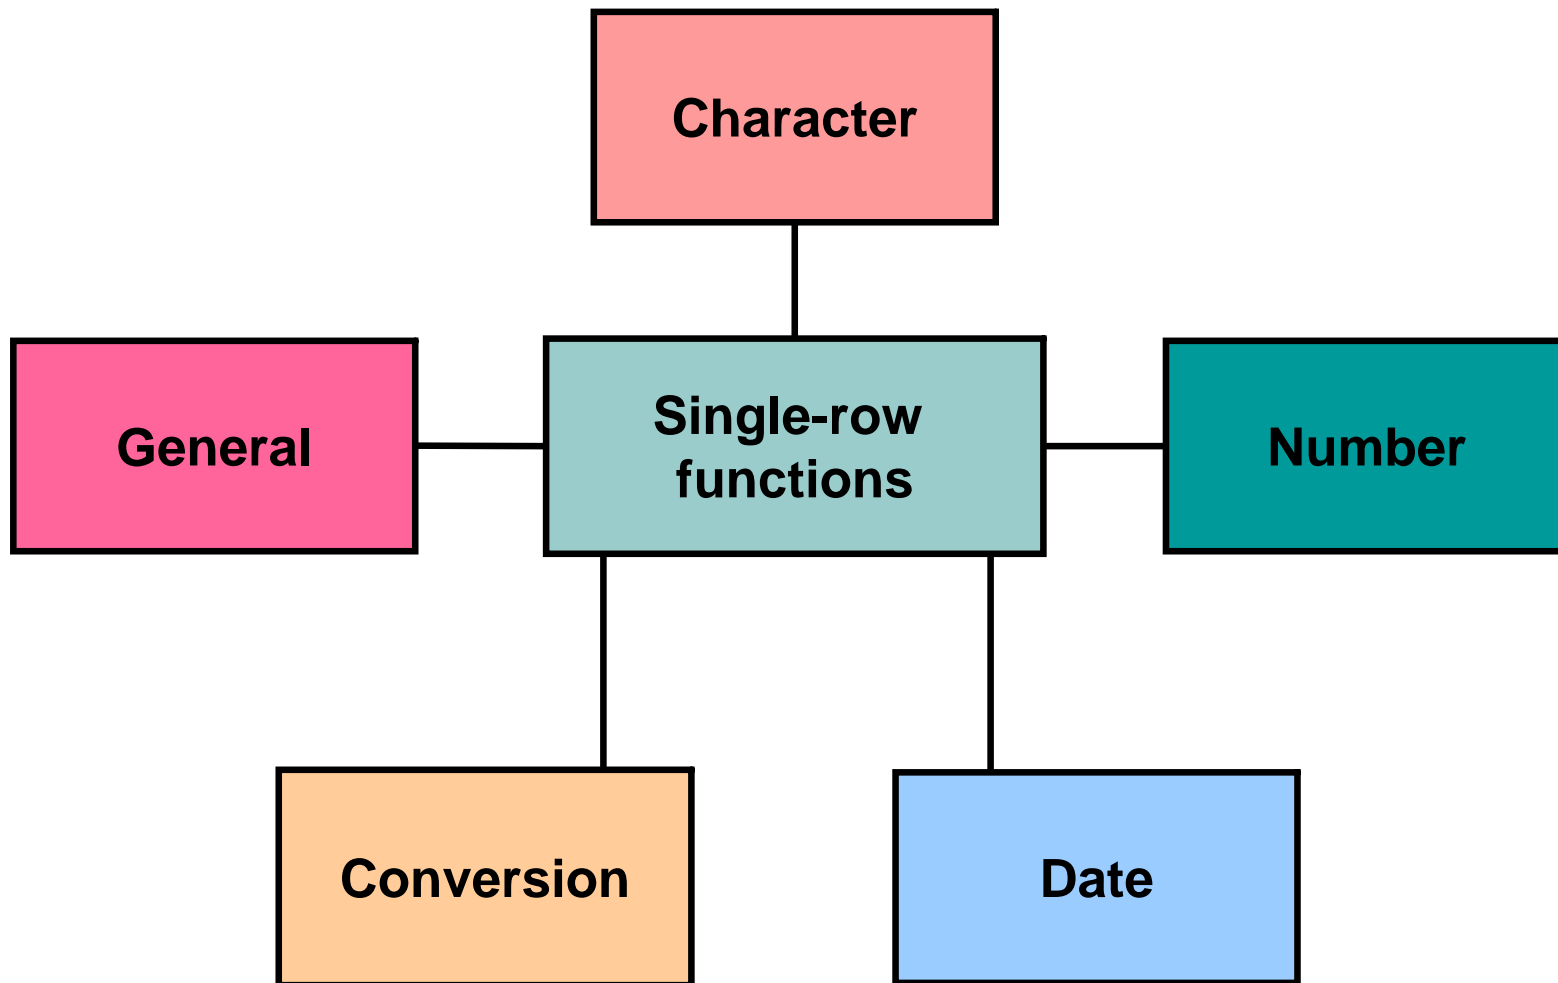
...

# Review of SQL Functions






# Review of Single-Row Functions



# Review of Types of Group Functions

- AVG
- COUNT
- MAX
- MIN
- STDDEV
- SUM
- VARIANCE



# Review of Using Subqueries

- A subquery is a `SELECT` statement nested in a clause of another `SELECT` statement.
- Syntax:

```
SELECT select_list
FROM   table
WHERE  expr operator
              (SELECT select_list
               FROM   table );
```

- Types of subqueries:

Single-row subquery	Multiple-row subquery
Returns only one row	Returns more than one row
Uses single-row comparison operators	Uses multiple-row comparison operators

# Review of Manipulating Data

A data manipulation language (DML) statement is executed when you:

- Add new rows to a table
- Modify existing rows in a table
- Remove existing rows from a table

Function	Description
INSERT	Adds a new row to the table
UPDATE	Modifies existing rows in the table
DELETE	Removes existing rows from the table
MERGE	Updates, inserts, or deletes a row conditionally into/from a table

# Lesson Agenda

- Course objectives and course agenda
- The database schema and appendixes used in the course and the available development environment in this course
- Review of some basic concepts of SQL
- **Oracle Database 11g documentation and additional resources**

# Oracle Database 11g SQL Documentation

- *Oracle Database New Features Guide 11g Release 2 (11.2)*
- *Oracle Database Reference 11g Release 2 (11.2)*
- *Oracle Database SQL Language Reference 11g Release 2 (11.2)*
- *Oracle Database Concepts 11g Release 2 (11.2)*
- *Oracle Database SQL Developer User's Guide Release 1.2*

# Additional Resources

For additional information about the new Oracle 11g SQL, refer to the following:

- *Oracle Database 11g: New Features eStudies*
- *Oracle by Example series (OBE): Oracle Database 11g*

# Summary

In this lesson, you should have learned the following:

- The course objectives
- The sample tables used in the course



# Practice I: Overview

This practice covers the following topics:

- Running the SQL Developer online tutorial
- Starting SQL Developer and creating a new database connection and browsing the tables
- Executing SQL statements using the SQL Worksheet
- Reviewing the basic concepts of SQL

# 1

## Controlling User Access

# Objectives


After completing this lesson, you should be able to do the following:

- Differentiate system privileges from object privileges
- Grant privileges on tables
- Grant roles
- Distinguish between privileges and roles

# Lesson Agenda

- System privileges
- Creating a role
- Object privileges
- Revoking object privileges

# Controlling User Access



# Privileges

- Database security:
  - System security
  - Data security
- System privileges: Performing a particular action within the database
- Object privileges: Manipulating the content of the database objects
- Schemas: Collection of objects such as tables, views, and sequences

# System Privileges

- More than 100 privileges are available.
- The database administrator has high-level system privileges for tasks such as:
  - Creating new users
  - Removing users
  - Removing tables
  - Backing up tables

# Creating Users

The DBA creates users with the CREATE USER statement.

```
CREATE USER user  
IDENTIFIED BY password;
```

```
CREATE USER demo  
IDENTIFIED BY demo;
```



# User System Privileges

- After a user is created, the DBA can grant specific system privileges to that user.

```
GRANT privilege [, privilege...]  
TO user [, user/ role, PUBLIC...];
```

- An application developer, for example, may have the following system privileges:
  - CREATE SESSION
  - CREATE TABLE
  - CREATE SEQUENCE
  - CREATE VIEW
  - CREATE PROCEDURE

# Granting System Privileges


The DBA can grant specific system privileges to a user.

```
GRANT  create session, create table,  
       create sequence, create view  
TO     demo;
```

# Lesson Agenda

- System privileges
- **Creating a role**
- Object privileges
- Revoking object privileges

# What Is a Role?



# Creating and Granting Privileges to a Role

- Create a role:

```
CREATE ROLE manager;
```

- Grant privileges to a role:

```
GRANT create table, create view  
TO manager;
```

- Grant a role to users:

```
GRANT manager TO alice;
```

# Changing Your Password

- The DBA creates your user account and initializes your password.
- You can change your password by using the ALTER USER statement.

```
ALTER USER demo  
IDENTIFIED BY employ;
```

# Lesson Agenda

- System privileges
- Creating a role
- Object privileges
- Revoking object privileges

# Object Privileges

Object privilege	Table	View	Sequence
ALTER	✓		✓
DELETE	✓	✓	
INDEX	✓		
INSERT	✓	✓	
REFERENCES	✓		
SELECT	✓	✓	✓
UPDATE	✓	✓	



# Object Privileges

- Object privileges vary from object to object.
- An owner has all the privileges on the object.
- An owner can give specific privileges on that owner's object.

```
GRANT      object_priv [(columns)]  
ON         object  
TO         {user|role|PUBLIC}  
[WITH GRANT OPTION];
```

# Granting Object Privileges

- Grant query privileges on the EMPLOYEES table:

```
GRANT  select
ON     employees
TO     demo;
```

- Grant privileges to update specific columns to users and roles:

```
GRANT  update (department_name, location_id)
ON     departments
TO     demo, manager;
```

# Passing On Your Privileges

- Give a user authority to pass along privileges:

```
GRANT  select, insert
ON     departments
TO     demo
WITH   GRANT OPTION;
```

- Allow all users on the system to query data from Alice's DEPARTMENTS table:

```
GRANT  select
ON     alice.departments
TO     PUBLIC;
```

# Confirming Granted Privileges

Data Dictionary View	Description
ROLE_SYS_PRIVS	System privileges granted to roles
ROLE_TAB_PRIVS	Table privileges granted to roles
USER_ROLE_PRIVS	Roles accessible by the user
USER_SYS_PRIVS	System privileges granted to the user
USER_TAB_PRIVS_MADE	Object privileges granted on the user's objects
USER_TAB_PRIVS_RECD	Object privileges granted to the user
USER_COL_PRIVS_MADE	Object privileges granted on the columns of the user's objects
USER_COL_PRIVS_RECD	Object privileges granted to the user on specific columns

# Lesson Agenda

- System privileges
- Creating a role
- Object privileges
- Revoking object privileges

# Revoking Object Privileges

- You use the REVOKE statement to revoke privileges granted to other users.
- Privileges granted to others through the WITH GRANT OPTION clause are also revoked.

```
REVOKE {privilege [, privilege...] | ALL}
ON      object
FROM    {user[, user...] | role | PUBLIC}
[CASCADE CONSTRAINTS];
```

# Revoking Object Privileges

Revoke the SELECT and INSERT privileges given to the demo user on the DEPARTMENTS table.

```
REVOKE  select, insert
ON      departments
FROM    demo;
```

# Quiz

Which of the following statements are true?

1. After a user creates an object, the user can pass along any of the available object privileges to other users by using the `GRANT` statement.
2. A user can create roles by using the `CREATE ROLE` statement to pass along a collection of system or object privileges to other users.
3. Users can change their own passwords.
4. Users can view the privileges granted to them and those that are granted on their objects.



# Summary

In this lesson, you should have learned how to:

- Differentiate system privileges from object privileges
- Grant privileges on tables
- Grant roles
- Distinguish between privileges and roles

# Practice 1: Overview

This practice covers the following topics:

- Granting other users privileges to your table
- Modifying another user's table through the privileges granted to you



# **Managing Schema Objects**

# Objectives

After completing this lesson, you should be able to do the following:

- Add constraints
- Create indexes
- Create indexes by using the `CREATE TABLE` statement
- Create function-based indexes
- Drop columns and set columns as `UNUSED`
- Perform `FLASHBACK` operations
- Create and use external tables

# Lesson Agenda

- Using the ALTER TABLE statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- Creating indexes:
  - Using the CREATE TABLE statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- Creating and using external tables

# ALTER TABLE Statement

Use the ALTER TABLE statement to:

- Add a new column
- Modify an existing column
- Define a default value for the new column
- Drop a column

# ALTER TABLE Statement

Use the ALTER TABLE statement to add, modify, or drop columns:

```
ALTER TABLE table  
ADD          (column datatype [DEFAULT expr]  
              [, column datatype] ...);
```

```
ALTER TABLE table  
MODIFY       (column datatype [DEFAULT expr]  
              [, column datatype] ...);
```

```
ALTER TABLE table  
DROP (column [, column] ...);
```

# Adding a Column

- You use the ADD clause to add columns:

```
ALTER TABLE dept80  
ADD          (job_id VARCHAR2(9));
```

```
ALTER TABLE dept80 succeeded.
```

- The new column becomes the last column:

	EMPLOYEE_ID	LAST_NAME	ANNSAL	HIRE_DATE	JOB_ID
1	145	Russell	14000	01-OCT-96	(null)
2	146	Partners	13500	05-JAN-97	(null)
3	147	Errazuriz	12000	10-MAR-97	(null)
4	148	Cambrault	11000	15-OCT-99	(null)
5	149	Zlotkey	10500	29-JAN-00	(null)



# Modifying a Column

- You can change a column's data type, size, and default value.

```
ALTER TABLE dept80  
MODIFY      (last_name VARCHAR2(30));
```

```
ALTER TABLE dept80 succeeded.
```

- A change to the default value affects only subsequent insertions to the table.

# Dropping a Column

Use the DROP COLUMN clause to drop columns that you no longer need from the table:

```
ALTER TABLE dept80  
DROP COLUMN job_id;
```

```
ALTER TABLE dept80 succeeded.
```

	EMPLOYEE_ID	LAST_NAME	ANNSAL	HIRE_DATE
1	145	Russell	14000	01-OCT-96
2	146	Partners	13500	05-JAN-97
3	147	Errazuriz	12000	10-MAR-97
4	148	Cambrault	11000	15-OCT-99
5	149	Zlotkey	10500	29-JAN-00

# SET UNUSED Option

- You use the SET UNUSED option to mark one or more columns as unused.
- You use the DROP UNUSED COLUMNS option to remove the columns that are marked as unused.

```
ALTER TABLE <table_name>  
SET UNUSED(<column name> [ , <column_name>]) ;
```

OR

```
ALTER TABLE <table_name>  
SET UNUSED COLUMN <column name> [ , <column_name>];
```

```
ALTER TABLE <table_name>  
DROP UNUSED COLUMNS;
```

# Lesson Agenda

- Using the `ALTER TABLE` statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- Creating indexes:
  - Using the `CREATE TABLE` statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- Creating and using external tables

# Adding a Constraint Syntax

Use the ALTER TABLE statement to:

- Add or drop a constraint, but not modify its structure
- Enable or disable constraints
- Add a NOT NULL constraint by using the MODIFY clause

```
ALTER TABLE  <table_name>  
ADD [CONSTRAINT <constraint_name>]  
type (<column_name>);
```

# Adding a Constraint

Add a FOREIGN KEY constraint to the EMP2 table indicating that a manager must already exist as a valid employee in the EMP2 table.

```
ALTER TABLE emp2  
MODIFY employee_id PRIMARY KEY;
```

```
ALTER TABLE emp2 succeeded.
```

```
ALTER TABLE emp2  
ADD CONSTRAINT emp_mgr_fk  
FOREIGN KEY(manager_id)  
REFERENCES emp2(employee_id);
```

```
ALTER TABLE succeeded.
```

# ON DELETE Clause

- Use the ON DELETE CASCADE clause to delete child rows when a parent key is deleted:

```
ALTER TABLE emp2 ADD CONSTRAINT emp_dt_fk  
FOREIGN KEY (Department_id)  
REFERENCES departments(department_id) ON DELETE CASCADE;
```

```
ALTER TABLE Emp2 succeeded.
```

- Use the ON DELETE SET NULL clause to set the child rows value to null when a parent key is deleted:

```
ALTER TABLE emp2 ADD CONSTRAINT emp_dt_fk  
FOREIGN KEY (Department_id)  
REFERENCES departments(department_id) ON DELETE SET NULL;
```

```
ALTER TABLE Emp2 succeeded.
```

# Deferring Constraints

Constraints can have the following attributes:

- DEFERRABLE or NOT DEFERRABLE
- INITIALLY DEFERRED or INITIALLY IMMEDIATE

```
ALTER TABLE dept2  
ADD CONSTRAINT dept2_id_pk  
PRIMARY KEY (department_id)  
DEFERRABLE INITIALLY DEFERRED
```

Deferring constraint on  
creation

```
SET CONSTRAINTS dept2_id_pk IMMEDIATE
```

Changing a specific  
constraint attribute

```
ALTER SESSION  
SET CONSTRAINTS= IMMEDIATE
```

Changing all constraints for a  
session



# Difference Between INITIALLY DEFERRED and INITIALLY IMMEDIATE

INITIALLY DEFERRED	Waits to check the constraint until the transaction ends
INITIALLY IMMEDIATE	Checks the constraint at the end of the statement execution

```
CREATE TABLE emp_new_sal (salary NUMBER
    CONSTRAINT sal_ck
    CHECK (salary > 100)
    DEFERRABLE INITIALLY IMMEDIATE,
    bonus NUMBER
    CONSTRAINT bonus_ck
    CHECK (bonus > 0 )
    DEFERRABLE INITIALLY DEFERRED );
```

```
create table succeeded.
```

# Dropping a Constraint

- Remove the manager constraint from the EMP2 table:

```
ALTER TABLE emp2  
DROP CONSTRAINT emp_mgr_fk;
```

```
ALTER TABLE Emp2 succeeded.
```

- Remove the PRIMARY KEY constraint on the DEPT2 table and drop the associated FOREIGN KEY constraint on the EMP2.DEPARTMENT\_ID column:

```
ALTER TABLE dept2  
DROP PRIMARY KEY CASCADE;
```

```
ALTER TABLE dept2 succeeded.
```

# Disabling Constraints

- Execute the `DISABLE` clause of the `ALTER TABLE` statement to deactivate an integrity constraint.
- Apply the `CASCADE` option to disable dependent integrity constraints.

```
ALTER TABLE emp2  
DISABLE CONSTRAINT emp_dt_fk;
```

```
ALTER TABLE Emp2 succeeded.
```

# Enabling Constraints

- Activate an integrity constraint currently disabled in the table definition by using the `ENABLE` clause.

```
ALTER TABLE      emp2  
ENABLE CONSTRAINT emp_dt_fk;
```

```
ALTER TABLE Emp2 succeeded.
```

- A `UNIQUE` index is automatically created if you enable a `UNIQUE` key or a `PRIMARY KEY` constraint.

# Cascading Constraints

- The `CASCADE CONSTRAINTS` clause is used along with the `DROP COLUMN` clause.
- The `CASCADE CONSTRAINTS` clause drops all referential integrity constraints that refer to the `PRIMARY` and `UNIQUE` keys defined on the dropped columns.
- The `CASCADE CONSTRAINTS` clause also drops all multicolumn constraints defined on the dropped columns.

# Cascading Constraints

Example:

```
ALTER TABLE emp2  
DROP COLUMN employee_id CASCADE CONSTRAINTS;
```

```
ALTER TABLE Emp2 succeeded.
```

```
ALTER TABLE test1  
DROP (col1_pk, col2_fk, col1) CASCADE CONSTRAINTS;
```

```
ALTER TABLE test1 succeeded.
```

# Renaming Table Columns and Constraints

Use the RENAME COLUMN clause of the ALTER TABLE statement to rename table columns.

**a**

```
ALTER TABLE marketing RENAME COLUMN team_id  
TO id;
```

```
ALTER TABLE marketing succeeded.
```

Use the RENAME CONSTRAINT clause of the ALTER TABLE statement to rename any existing constraint for a table.

**b**

```
ALTER TABLE marketing RENAME CONSTRAINT mktg_pk  
TO new_mktg_pk;
```

```
ALTER TABLE marketing succeeded.
```

# Lesson Agenda

- Using the `ALTER TABLE` statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- **Creating indexes:**
  - Using the `CREATE TABLE` statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- Creating and using external tables



# Overview of Indexes

Indexes are created:

- Automatically
  - PRIMARY KEY creation
  - UNIQUE KEY creation
- Manually
  - The CREATE INDEX statement
  - The CREATE TABLE statement

# CREATE INDEX with the CREATE TABLE Statement

```
CREATE TABLE NEW_EMP  
(employee_id NUMBER(6)  
    PRIMARY KEY USING INDEX  
    (CREATE INDEX emp_id_idx ON  
    NEW_EMP(employee_id)),  
first_name VARCHAR2(20),  
last_name VARCHAR2(25));
```

CREATE TABLE succeeded.

```
SELECT INDEX_NAME, TABLE_NAME  
FROM USER_INDEXES  
WHERE TABLE_NAME = 'NEW_EMP';
```

	INDEX_NAME	TABLE_NAME
1	EMP_ID_IDX	NEW_EMP

# Function-Based Indexes

- A function-based index is based on expressions.
- The index expression is built from table columns, constants, SQL functions, and user-defined functions.

```
CREATE INDEX upper_dept_name_idx  
ON dept2 (UPPER(department_name));
```

```
CREATE INDEX succeeded.
```

```
SELECT *  
FROM   dept2  
WHERE  UPPER(department_name) = 'SALES';
```

# Removing an Index

- Remove an index from the data dictionary by using the DROP INDEX command:

```
DROP INDEX index;
```

- Remove the UPPER\_DEPT\_NAME\_IDX index from the data dictionary:

```
DROP INDEX upper_dept_name_idx;
```

```
DROP INDEX upper_dept_name_idx succeeded.
```

- To drop an index, you must be the owner of the index or have the DROP ANY INDEX privilege.

# DROP TABLE ... PURGE

```
DROP TABLE dept80 PURGE;
```

```
DROP TABLE dept80 succeeded.
```

# Lesson Agenda

- Using the `ALTER TABLE` statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- Creating indexes:
  - Using the `CREATE TABLE` statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- Creating and using external tables

# FLASHBACK TABLE Statement


- Enables you to recover tables to a specified point in time with a single statement
- Restores table data along with associated indexes and constraints
- Enables you to revert the table and its contents to a certain point in time or system change number (SCN)



# FLASHBACK TABLE Statement

- Repair tool for accidental table modifications
  - Restores a table to an earlier point in time
  - Benefits: Ease of use, availability, and fast execution
  - Is performed in place
- Syntax:

```
FLASHBACK TABLE [schema.] table [,  
[ schema.] table ] ...  
TO { TIMESTAMP | SCN } expr  
[ { ENABLE | DISABLE } TRIGGERS ] ;
```





# Using the FLASHBACK TABLE Statement

```
DROP TABLE emp2;
```

```
DROP TABLE emp2 succeeded.
```

```
SELECT original_name, operation, droptime FROM  
recyclebin;
```

ORIGINAL_NAME	OPERATION	DROPTIME
EMP2	DROP	2009-05-20:18:00:39

...


```
FLASHBACK TABLE emp2 TO BEFORE DROP;
```

```
FLASHBACK TABLE succeeded.
```

# Lesson Agenda

- Using the `ALTER TABLE` statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- Creating indexes:
  - Using the `CREATE TABLE` statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- Creating and using external tables

# Temporary Tables



# Creating a Temporary Table

```
CREATE GLOBAL TEMPORARY TABLE cart  
ON COMMIT DELETE ROWS;
```

1


```
CREATE GLOBAL TEMPORARY TABLE today_sales  
ON COMMIT PRESERVE ROWS AS  
    SELECT * FROM orders  
        WHERE order_date = SYSDATE;
```

2

# Lesson Agenda

- Using the ALTER TABLE statement to add, modify, and drop a column
- Managing constraints:
  - Adding and dropping a constraint
  - Deferring constraints
  - Enabling and disabling a constraint
- Creating indexes:
  - Using the CREATE TABLE statement
  - Creating function-based indexes
  - Removing an index
- Performing flashback operations
- Creating and using temporary tables
- **Creating and using external tables**

# External Tables



# Creating a Directory for the External Table

Create a `DIRECTORY` object that corresponds to the directory on the file system where the external data source resides.

```
CREATE OR REPLACE DIRECTORY emp_dir  
AS '/.../emp_dir';  
  
GRANT READ ON DIRECTORY emp_dir TO ora_21;
```

# Creating an External Table

```
CREATE TABLE <table_name>
  ( <col_name> <datatype>, ... )
ORGANIZATION EXTERNAL
  (TYPE <access_driver_type>
    DEFAULT DIRECTORY <directory_name>
    ACCESS PARAMETERS
      (... ) )
    LOCATION ('<location_specifier>')
REJECT LIMIT [0 | <number> | UNLIMITED];
```




# Creating an External Table by Using ORACLE\_LOADER

```
CREATE TABLE oldemp (  
  fname char(25), lname CHAR(25))  
  ORGANIZATION EXTERNAL  
  (TYPE ORACLE_LOADER  
  DEFAULT DIRECTORY emp_dir  
  ACCESS PARAMETERS  
  (RECORDS DELIMITED BY NEWLINE  
  NOBADFILE  
  NOLOGFILE  
  FIELDS TERMINATED BY ','  
  (fname POSITION ( 1:20) CHAR,  
  lname POSITION (22:41) CHAR))  
  LOCATION ('emp.dat'))  
  PARALLEL 5  
  REJECT LIMIT 200;
```

CREATE TABLE succeeded.

# Querying External Tables



# Creating an External Table by Using ORACLE\_DATAPUMP: Example

```
CREATE TABLE emp_ext
(employee_id, first_name, last_name)
ORGANIZATION EXTERNAL
(
  TYPE ORACLE_DATAPUMP
  DEFAULT DIRECTORY emp_dir
  LOCATION
    ('emp1.exp', 'emp2.exp')
)
PARALLEL
AS
SELECT employee_id, first_name, last_name
FROM   employees;
```

# Quiz

A FOREIGN KEY constraint enforces the following action:  
When the data in the parent key is deleted, all the rows in the child table that depend on the deleted parent key values are also deleted.

1. True
2. False

# Quiz

In all the cases, when you execute a `DROP TABLE` command, the database renames the table and places it in a recycle bin, from where it can later be recovered by using the `FLASHBACK TABLE` statement.

1. True
2. False

# Summary

In this lesson, you should have learned how to:

- Add constraints
- Create indexes
- Create indexes by using the `CREATE TABLE` statement
- Create function-based indexes
- Drop columns and set columns as `UNUSED`
- Perform `FLASHBACK` operations
- Create and use external tables

## Practice 2: Overview

This practice covers the following topics:

- Altering tables
- Adding columns
- Dropping columns
- Creating indexes
- Creating external tables



# **Managing Objects with Data Dictionary Views**



# Objectives

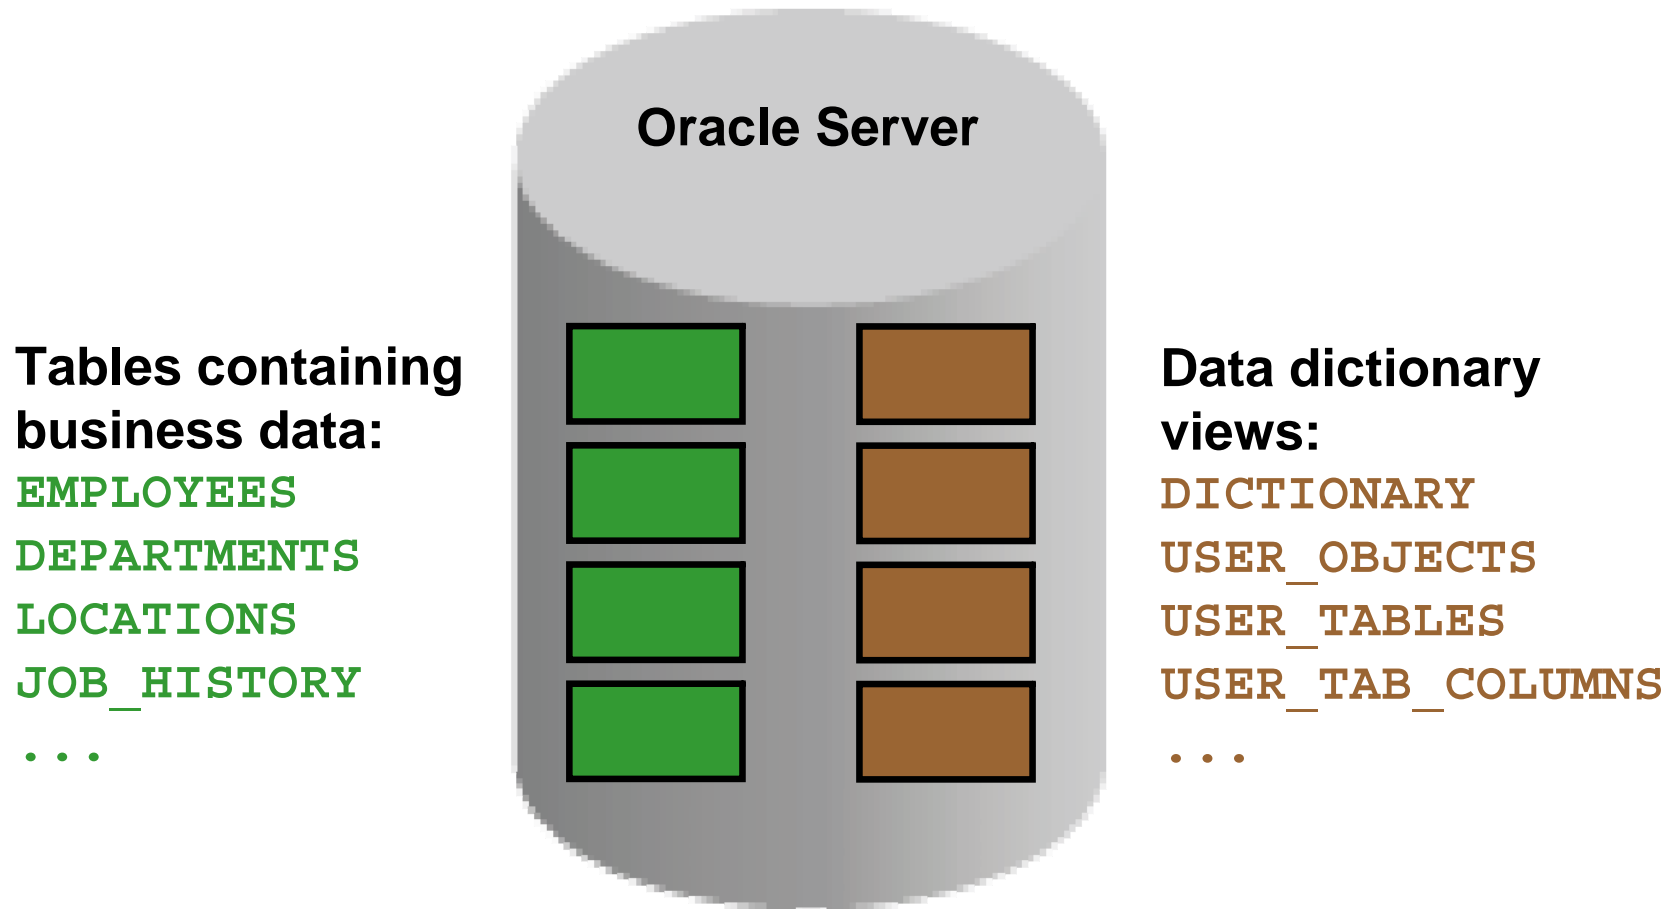
After completing this lesson, you should be able to do the following:

- Use the data dictionary views to research data on your objects
- Query various data dictionary views


# Lesson Agenda

- Introduction to data dictionary
- Querying the dictionary views for the following:
  - Table information
  - Column information
  - Constraint information
- Querying the dictionary views for the following:
  - View information
  - Sequence information
  - Synonym information
  - Index information
- Adding a comment to a table and querying the dictionary views for comment information

# Data Dictionary



# Data Dictionary Structure



# Data Dictionary Structure

View naming convention:

View Prefix	Purpose
USER	User's view (what is in your schema; what you own)
ALL	Expanded user's view (what you can access)
DBA	Database administrator's view (what is in everyone's schemas)
V\$	Performance-related data

# How to Use the Dictionary Views

Start with `DICTIONARY`. It contains the names and descriptions of the dictionary tables and views.

```
DESCRIBE DICTIONARY
```

Name	Null	Type
TABLE_NAME		VARCHAR2(30)
COMMENTS		VARCHAR2(4000)

2 rows selected

```
SELECT *  
FROM   dictionary  
WHERE  table_name = 'USER_OBJECTS';
```

	TABLE_NAME	COMMENTS
1	USER_OBJECTS	Objects owned by the user

# USER\_OBJECTS and ALL\_OBJECTS Views

## USER\_OBJECTS:





- Query USER\_OBJECTS to see all the objects that you own.
- Using USER\_OBJECTS, you can obtain a listing of all object names and types in your schema, plus the following information:
  - Date created
  - Date of last modification
  - Status (valid or invalid)

## ALL\_OBJECTS:

- Query ALL\_OBJECTS to see all the objects to which you have access.

# USER\_OBJECTS View

```
SELECT object_name, object_type, created, status
FROM   user_objects
ORDER BY object_type;
```

	 OBJECT_NAME	 OBJECT_TYPE	 CREATED	 STATUS
1	LOC_COUNTRY_IX	INDEX	19-MAY-09	VALID

...

53	EMPLOYEES2	TABLE	22-MAY-09	VALID
54	SECURE_EMPLOYEES	TRIGGER	19-MAY-09	VALID
55	UPDATE_JOB_HISTORY	TRIGGER	19-MAY-09	VALID
56	EMP_DETAILS_VIEW	VIEW	19-MAY-09	VALID

...



# Lesson Agenda

- Introduction to data dictionary
- Querying the dictionary views for the following:
  - Table information
  - Column information
  - Constraint information
- Querying the dictionary views for the following:
  - View information
  - Sequence information
  - Synonym information
  - Index information
- Adding a comment to a table and querying the dictionary views for comment information

# Table Information

USER\_TABLES:

```
DESCRIBE user_tables
```

Name	Null	Type
-----	-----	-----
TABLE_NAME	NOT NULL	VARCHAR2(30)
TABLESPACE_NAME		VARCHAR2(30)
CLUSTER_NAME		VARCHAR2(30)
IOT_NAME		VARCHAR2(30)

...

```
SELECT table_name
FROM   user_tables;
```

R	TABLE_NAME
1	REGIONS
2	LOCATIONS
3	DEPARTMENTS
4	JOBS
5	EMPLOYEES
6	JOB_HISTORY

...

# Column Information

USER\_TAB\_COLUMNS:

```
DESCRIBE user_tab_columns
```

Name	Null	Type
TABLE_NAME	NOT NULL	VARCHAR2(30)
COLUMN_NAME	NOT NULL	VARCHAR2(30)
DATA_TYPE		VARCHAR2(106)
DATA_TYPE_MOD		VARCHAR2(3)
DATA_TYPE_OWNER		VARCHAR2(30)
DATA_LENGTH	NOT NULL	NUMBER
DATA_PRECISION		NUMBER
DATA_SCALE		NUMBER
NULLABLE		VARCHAR2(1)

...

# Column Information

```
SELECT column_name, data_type, data_length,  
       data_precision, data_scale, nullable  
FROM   user_tab_columns  
WHERE  table_name = 'EMPLOYEES';
```

	COLUMN_NAME	DATA_TYPE	DATA_LENGTH	DATA_PRECISION
1	EMPLOYEE_ID	NUMBER	22	6
2	FIRST_NAME	VARCHAR2	20	(null)
3	LAST_NAME	VARCHAR2	25	(null)
4	EMAIL	VARCHAR2	25	(null)
5	PHONE_NUMBER	VARCHAR2	20	(null)
6	HIRE_DATE	DATE	7	(null)
7	JOB_ID	VARCHAR2	10	(null)
8	SALARY	NUMBER	22	8
9	COMMISSION_PCT	NUMBER	22	2
10	MANAGER_ID	NUMBER	22	6
11	DEPARTMENT_ID	NUMBER	22	4

# Constraint Information

- USER\_CONSTRAINTS describes the constraint definitions on your tables.
- USER\_CONS\_COLUMNS describes columns that are owned by you and that are specified in constraints.

```
DESCRIBE user_constraints
```

Name	Null	Type
OWNER	NOT NULL	VARCHAR2(30)
CONSTRAINT_NAME	NOT NULL	VARCHAR2(30)
CONSTRAINT_TYPE		VARCHAR2(1)
TABLE_NAME	NOT NULL	VARCHAR2(30)
SEARCH_CONDITION		LONG()
R_OWNER		VARCHAR2(30)
R_CONSTRAINT_NAME		VARCHAR2(30)
DELETE_RULE		VARCHAR2(9)
STATUS		VARCHAR2(8)

...

# USER\_CONSTRAINTS: Example

```
SELECT constraint_name, constraint_type,
       search_condition, r_constraint_name,
       delete_rule, status
FROM   user_constraints
WHERE  table_name = 'EMPLOYEES';
```

	CONSTRAINT_NAME	C...	SEARCH_CONDITION	R_CONSTR...	DELET...	STATUS
1	EMP_LAST_NAME_NN	C	"LAST_NAME" IS NOT NULL	(null)	(null)	ENABLED
2	EMP_EMAIL_NN	C	"EMAIL" IS NOT NULL	(null)	(null)	ENABLED
3	EMP_HIRE_DATE_NN	C	"HIRE_DATE" IS NOT NULL	(null)	(null)	ENABLED
4	EMP_JOB_NN	C	"JOB_ID" IS NOT NULL	(null)	(null)	ENABLED
5	EMP_SALARY_MIN	C	salary > 0	(null)	(null)	ENABLED
6	EMP_EMAIL_UK	U	(null)	(null)	(null)	ENABLED
7	EMP_EMP_ID_PK	P	(null)	(null)	(null)	ENABLED
8	EMP_DEPT_FK	R	(null)	DEPT_ID_PK	NO ACTION	ENABLED
9	EMP_JOB_FK	R	(null)	JOB_ID_PK	NO ACTION	ENABLED
10	EMP_MANAGER_FK	R	(null)	EMP_EMP_ID_PK	NO ACTION	ENABLED

# Querying USER\_CONS\_COLUMNS

```
DESCRIBE user_cons_columns
```

Name	Null	Type
OWNER	NOT NULL	VARCHAR2(30)
CONSTRAINT_NAME	NOT NULL	VARCHAR2(30)
TABLE_NAME	NOT NULL	VARCHAR2(30)
COLUMN_NAME		VARCHAR2(4000)
POSITION		NUMBER

```
SELECT constraint_name, column_name
FROM   user_cons_columns
WHERE  table_name = 'EMPLOYEES';
```

	CONSTRAINT_NAME	COLUMN_NAME
1	EMP_LAST_NAME_NN	LAST_NAME
2	EMP_EMAIL_NN	EMAIL
3	EMP_HIRE_DATE_NN	HIRE_DATE
4	EMP_JOB_NN	JOB_ID
5	EMP_SALARY_MIN	SALARY
6	EMP_EMAIL_UK	EMAIL

...

# Lesson Agenda

- Introduction to data dictionary
- Querying the dictionary views for the following:
  - Table information
  - Column information
  - Constraint information
- Querying the dictionary views for the following:
  - View information
  - Sequence information
  - Synonym information
  - Index information
- Adding a comment to a table and querying the dictionary views for comment information



# View Information

1

```
DESCRIBE user_views
```

Name	Null	Type
-----	-----	-----
VIEW_NAME	NOT NULL	VARCHAR2(30)
TEXT_LENGTH		NUMBER
TEXT		LONG()

2

```
SELECT view_name FROM user_views;
```

	VIEW_NAME
1	EMP_DETAILS_VIEW

3

```
SELECT text FROM user_views  
WHERE view_name = 'EMP_DETAILS_VIEW';
```

	TEXT
1	SELECT e.employee_id, e.job_id, e.manager_id, e.department_id, d.location_id, l.co

...

```
AND c.region_id = r.region_id AND j.job_id = e.job_id WITH READ ONLY
```

# Sequence Information

```
DESCRIBE user_sequences
```

Name	Null	Type
-----		
SEQUENCE_NAME	NOT NULL	VARCHAR2(30)
MIN_VALUE		NUMBER
MAX_VALUE		NUMBER
INCREMENT_BY	NOT NULL	NUMBER
CYCLE_FLAG		VARCHAR2(1)
ORDER_FLAG		VARCHAR2(1)
CACHE_SIZE	NOT NULL	NUMBER
LAST_NUMBER	NOT NULL	NUMBER

# Confirming Sequences

- Verify your sequence values in the USER\_SEQUENCES data dictionary table.

```
SELECT    sequence_name, min_value, max_value,  
          increment_by, last_number  
FROM      user_sequences;
```

	SEQUENCE_NAME	MIN_VALUE	MAX_VALUE	INCREMENT_BY	LAST_NUMBER
1	DEPARTMENTS_SEQ	1	9990	10	280
2	EMPLOYEES_SEQ	1	999999999999999...	1	207
3	LOCATIONS_SEQ	1	9900	100	3300

- The LAST\_NUMBER column displays the next available sequence number if NOCACHE is specified.

# Index Information

- USER\_INDEXES provides information about your indexes.
- USER\_IND\_COLUMNS describes columns comprising your indexes and columns of indexes on your tables.

```
DESCRIBE user_indexes
```

Name	Null	Type
-----	-----	-----
INDEX_NAME	NOT NULL	VARCHAR2(30)
INDEX_TYPE		VARCHAR2(27)
TABLE_OWNER	NOT NULL	VARCHAR2(30)
TABLE_NAME	NOT NULL	VARCHAR2(30)
TABLE_TYPE		VARCHAR2(11)
UNIQUENESS		VARCHAR2(9)

...

# USER\_INDEXES: Examples

**a**

```
SELECT index_name, table_name, uniqueness
FROM   user_indexes
WHERE  table_name = 'EMPLOYEES';
```

	INDEX_NAME	TABLE_NAME	UNIQUENESS
1	EMP_EMAIL_UK	EMPLOYEES	UNIQUE
2	EMP_EMP_ID_PK	EMPLOYEES	UNIQUE
3	EMP_DEPARTMENT_IX	EMPLOYEES	NONUNIQUE
4	EMP_JOB_IX	EMPLOYEES	NONUNIQUE
5	EMP_MANAGER_IX	EMPLOYEES	NONUNIQUE
6	EMP_NAME_IX	EMPLOYEES	NONUNIQUE

**b**

```
SELECT index_name, table_name
FROM   user_indexes
WHERE  table_name = 'emp_lib';
```

	INDEX_NAME	TABLE_NAME
1	SYS_C0011777	EMP_LIB

# Querying USER\_IND\_COLUMNS

```
DESCRIBE user_ind_columns
```

Name	Null	Type
-----		
INDEX_NAME		VARCHAR2(30)
TABLE_NAME		VARCHAR2(30)
COLUMN_NAME		VARCHAR2(4000)
COLUMN_POSITION		NUMBER
COLUMN_LENGTH		NUMBER
CHAR_LENGTH		NUMBER
DESCEND		VARCHAR2(4)

```
SELECT index_name, column_name, table_name
FROM    user_ind_columns
WHERE   index_name = 'lname_idx';
```





INDEX_NAME	COLUMN_NAME	TABLE_NAME
1 LNAME_IDX	LAST_NAME	EMP_TEST

# Synonym Information

```
DESCRIBE user_synonyms
```

Name	Null	Type
-----	-----	-----
SYNONYM_NAME	NOT NULL	VARCHAR2(30)
TABLE_OWNER		VARCHAR2(30)
TABLE_NAME	NOT NULL	VARCHAR2(30)
DB_LINK		VARCHAR2(128)

```
SELECT *  
FROM   user_synonyms;
```

	 SYNONYM_NAME	 TABLE_OWNER	 TABLE_NAME	 DB_LINK
1	TEAM2	ORA22	DEPARTMENTS	(null)

# Lesson Agenda

- Introduction to data dictionary
- Querying the dictionary views for the following:
  - Table information
  - Column information
  - Constraint information
- Querying the dictionary views for the following:
  - View information
  - Sequence information
  - Synonym information
  - Index information
- Adding a comment to a table and querying the dictionary views for comment information



# Adding Comments to a Table

- You can add comments to a table or column by using the COMMENT statement:

```
COMMENT ON TABLE employees  
IS 'Employee Information';
```

```
COMMENT ON COLUMN employees.first_name  
IS 'First name of the employee';
```

- Comments can be viewed through the data dictionary views:
  - ALL\_COL\_COMMENTS
  - USER\_COL\_COMMENTS
  - ALL\_TAB\_COMMENTS
  - USER\_TAB\_COMMENTS

# Quiz

The dictionary views that are based on the dictionary tables contain information such as:

1. Definitions of all the schema objects in the database
2. Default values for the columns
3. Integrity constraint information
4. Privileges and roles that each user has been granted
5. All of the above

# Summary

In this lesson, you should have learned how to find information about your objects through the following dictionary views:

- `DICTIONARY`
- `USER_OBJECTS`
- `USER_TABLES`
- `USER_TAB_COLUMNS`
- `USER_CONSTRAINTS`
- `USER_CONS_COLUMNS`
- `USER_VIEWS`
- `USER_SEQUENCES`
- `USER_INDEXES`
- `USER_SYNONYMS`

## Practice 3: Overview

This practice covers the following topics:

- Querying the dictionary views for table and column information
- Querying the dictionary views for constraint information
- Querying the dictionary views for view information
- Querying the dictionary views for sequence information
- Querying the dictionary views for synonym information
- Querying the dictionary views for index information
- Adding a comment to a table and querying the dictionary views for comment information

# 4

## Manipulating Large Data Sets

# Objectives

After completing this lesson, you should be able to do the following:

- Manipulate data by using subqueries
- Specify explicit default values in the `INSERT` and `UPDATE` statements
- Describe the features of multitable `INSERTs`
- Use the following types of multitable `INSERTs`:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merge rows in a table
- Track the changes to data over a period of time

# Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

# Using Subqueries to Manipulate Data

You can use subqueries in data manipulation language (DML) statements to:

- Retrieve data by using an inline view
- Copy data from one table to another
- Update data in one table based on the values of another table
- Delete rows from one table based on rows in another table



# Retrieving Data by Using a Subquery as Source

```
SELECT department_name, city
FROM departments
NATURAL JOIN (SELECT l.location_id, l.city, l.country_id
              FROM loc l
              JOIN countries c
              ON(l.country_id = c.country_id)
              JOIN regions USING(region_id)
              WHERE region_name = 'Europe');
```

	DEPARTMENT_NAME	CITY
1	Human Resources	London
2	Sales	Oxford
3	Public Relations	Munich

# Inserting by Using a Subquery as a Target

```
INSERT INTO (SELECT l.location_id, l.city, l.country_id
              FROM    locations l
              JOIN     countries c
              ON(l.country_id = c.country_id)
              JOIN     regions USING(region_id)
              WHERE    region_name = 'Europe')
VALUES (3300, 'Cardiff', 'UK');
```

1 rows inserted

# Inserting by Using a Subquery as a Target

Verify the results.


```
SELECT location_id, city, country_id
FROM   loc
```

	LOCATION_ID	CITY	COUNTRY_ID
20	2900	Geneva	CH
21	3000	Bern	CH
22	3100	Utrecht	NL
23	3200	Mexico City	MX
24	3300	Cardiff	UK

# Using the WITH CHECK OPTION Keyword on DML Statements

The WITH CHECK OPTION keyword prohibits you from changing rows that are not in the subquery.

```
INSERT INTO ( SELECT location_id, city, country_id
              FROM   loc
              WHERE  country_id IN
                    (SELECT country_id
                     FROM   countries
                     NATURAL JOIN regions
                     WHERE  region name = 'Europe')
              WITH CHECK OPTION )
VALUES (3600, 'Washington', 'US');
```



# Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

# Overview of the Explicit Default Feature

- Use the `DEFAULT` keyword as a column value where the default column value is desired.
- This allows the user to control where and when the default value should be applied to data.
- Explicit defaults can be used in `INSERT` and `UPDATE` statements.

# Using Explicit Default Values

- DEFAULT with INSERT:

```
INSERT INTO deptm3  
  (department_id, department_name, manager_id)  
VALUES (300, 'Engineering', DEFAULT);
```

- DEFAULT with UPDATE:

```
UPDATE deptm3  
SET manager_id = DEFAULT  
WHERE department_id = 10;
```

# Copying Rows from Another Table

- Write your INSERT statement with a subquery.

```
INSERT INTO sales_reps(id, name, salary, commission_pct)
SELECT employee_id, last_name, salary, commission_pct
FROM employees
WHERE job_id LIKE '%REP%';
```

33 rows inserted


- Do not use the VALUES clause.
- Match the number of columns in the INSERT clause with that in the subquery.



# Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

# Overview of Multitable INSERT Statements



```
INSERT ALL  
  INTO target_a VALUES (...,...)  
  INTO target_b VALUES (...,...)  
  INTO target_c VALUES (...,...)  
  SELECT ...  
  FROM sourcetab  
  WHERE ...;
```

Subquery

# Overview of Multitable INSERT Statements

- Use the `INSERT...SELECT` statement to insert rows into multiple tables as part of a single DML statement.
- Multitable `INSERT` statements are used in data warehousing systems to transfer data from one or more operational sources to a set of target tables.
- They provide significant performance improvement over:
  - Single DML versus multiple `INSERT...SELECT` statements
  - Single DML versus a procedure to perform multiple inserts by using the `IF . . . THEN` syntax

# Types of Multitable INSERT Statements

The different types of multitable INSERT statements are:

- Unconditional INSERT
- Conditional INSERT ALL
- Pivoting INSERT
- Conditional INSERT FIRST

# Multitable INSERT Statements

- Syntax for multitable INSERT:

```
INSERT [conditional_insert_clause]  
[insert_into_clause values_clause] (subquery)
```

- conditional\_insert\_clause:

```
[ALL|FIRST]  
[WHEN condition THEN] [insert_into_clause values_clause]  
[ELSE] [insert_into_clause values_clause]
```

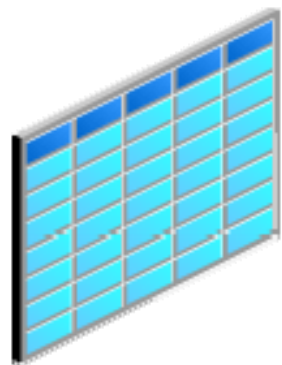
# Unconditional INSERT ALL

- Select the EMPLOYEE\_ID, HIRE\_DATE, SALARY, and MANAGER\_ID values from the EMPLOYEES table for those employees whose EMPLOYEE\_ID is greater than 200.
- Insert these values into the SAL\_HISTORY and MGR\_HISTORY tables by using a multitable INSERT.

```
INSERT ALL
  INTO sal_history VALUES (EMPID, HIREDATE, SAL)
  INTO mgr_history VALUES (EMPID, MGR, SAL)
  SELECT employee_id EMPID, hire_date HIREDATE,
         salary SAL, manager_id MGR
  FROM   employees
 WHERE  employee_id > 200;
```

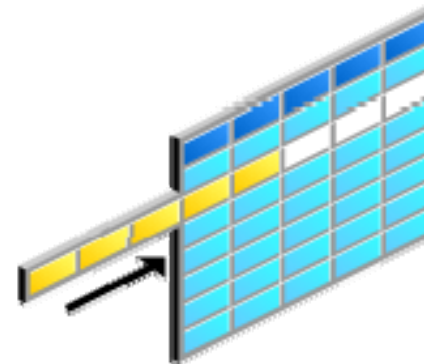
12 rows inserted

# Conditional INSERT ALL: Example

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue.

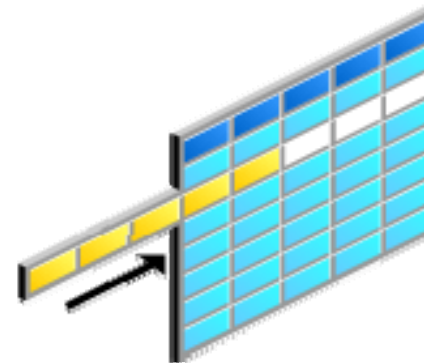
**Employees**

**Hired before  
1995**

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue. A yellow bar highlights the first column, and a black arrow points to it.

**EMP\_HISTORY**

**With sales  
commission**

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue. A yellow bar highlights the first column, and a black arrow points to it.

**EMP\_SALES**

# Conditional INSERT ALL

```
INSERT  ALL
```

```
  WHEN HIREDATE < '01-JAN-95' THEN
```

```
    INTO emp_history VALUES (EMPID, HIREDATE, SAL)
```

```
  WHEN COMM IS NOT NULL THEN
```

```
    INTO emp_sales VALUES (EMPID, COMM, SAL)
```

```
  SELECT employee_id EMPID, hire_date HIREDATE,
```

```
         salary SAL, commission_pct COMM
```

```
  FROM  employees
```


```
48 rows inserted
```




# Conditional INSERT FIRST: Example

**Scenario:** If an employee salary is 2,000, the record is inserted into the SAL\_LOW table only.

Salary < 5,000




SAL\_LOW





EMPLOYEES

5000 <= Salary  
<= 10,000



SAL\_MID

Otherwise



SAL\_HIGH

ORACLE

# Conditional INSERT FIRST

```
INSERT FIRST
```

```
WHEN salary < 5000 THEN
```

```
    INTO sal_low VALUES (employee_id, last_name, salary)
```

```
WHEN salary between 5000 and 10000 THEN
```

```
    INTO sal_mid VALUES (employee_id, last_name, salary)
```

```
ELSE
```

```
    INTO sal_high VALUES (employee_id, last_name, salary)
```

```
SELECT employee_id, last_name, salary
```


```
FROM employees
```

```
107 rows inserted
```

# Pivoting INSERT

Convert the set of sales records from the nonrelational database table to relational format.

Emp_ID	Week_ID	MON	TUES	WED	THUR	FRI
176	6	2000	3000	4000	5000	6000



Employee_ID	WEEK	SALES
176	6	2000
176	6	3000
176	6	4000
176	6	5000
176	6	6000

# Pivoting INSERT

```
INSERT ALL
  INTO sales_info VALUES (employee_id, week_id, sales_MON)
  INTO sales_info VALUES (employee_id, week_id, sales_TUE)
  INTO sales_info VALUES (employee_id, week_id, sales_WED)
  INTO sales_info VALUES (employee_id, week_id, sales_THUR)
  INTO sales_info VALUES (employee_id, week_id, sales_FRI)
SELECT EMPLOYEE_ID, week_id, sales_MON, sales_TUE,
       sales_WED, sales_THUR, sales_FRI
FROM sales_source_data;
```

5 rows inserted

# Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

# MERGE Statement

- Provides the ability to conditionally update, insert, or delete data into a database table
- Performs an `UPDATE` if the row exists, and an `INSERT` if it is a new row:
  - Avoids separate updates
  - Increases performance and ease of use
  - Is useful in data warehousing applications

# MERGE Statement Syntax

You can conditionally insert, update, or delete rows in a table by using the MERGE statement.

```
MERGE INTO table_name table_alias
  USING (table/view/sub_query) alias
  ON (join condition)
  WHEN MATCHED THEN
    UPDATE SET
      col1 = col1_val,
      col2 = col2_val
  WHEN NOT MATCHED THEN
    INSERT (column_list)
    VALUES (column_values);
```

# Merging Rows: Example

Insert or update rows in the COPY\_EMP3 table to match the EMPLOYEES table.

```
MERGE INTO copy_emp3 c
USING (SELECT * FROM EMPLOYEES ) e
ON (c.employee_id = e.employee_id)
WHEN MATCHED THEN
UPDATE SET
c.first_name = e.first_name,
c.last_name = e.last_name,
...
DELETE WHERE (E.COMMISSION_PCT IS NOT NULL)
WHEN NOT MATCHED THEN
INSERT VALUES(e.employee_id, e.first_name, e.last_name,
e.email, e.phone_number, e.hire_date, e.job_id,
e.salary, e.commission_pct, e.manager_id,
e.department_id);
```



# Merging Rows: Example

```
TRUNCATE TABLE copy_emp3;  
SELECT * FROM copy_emp3;  
0 rows selected
```

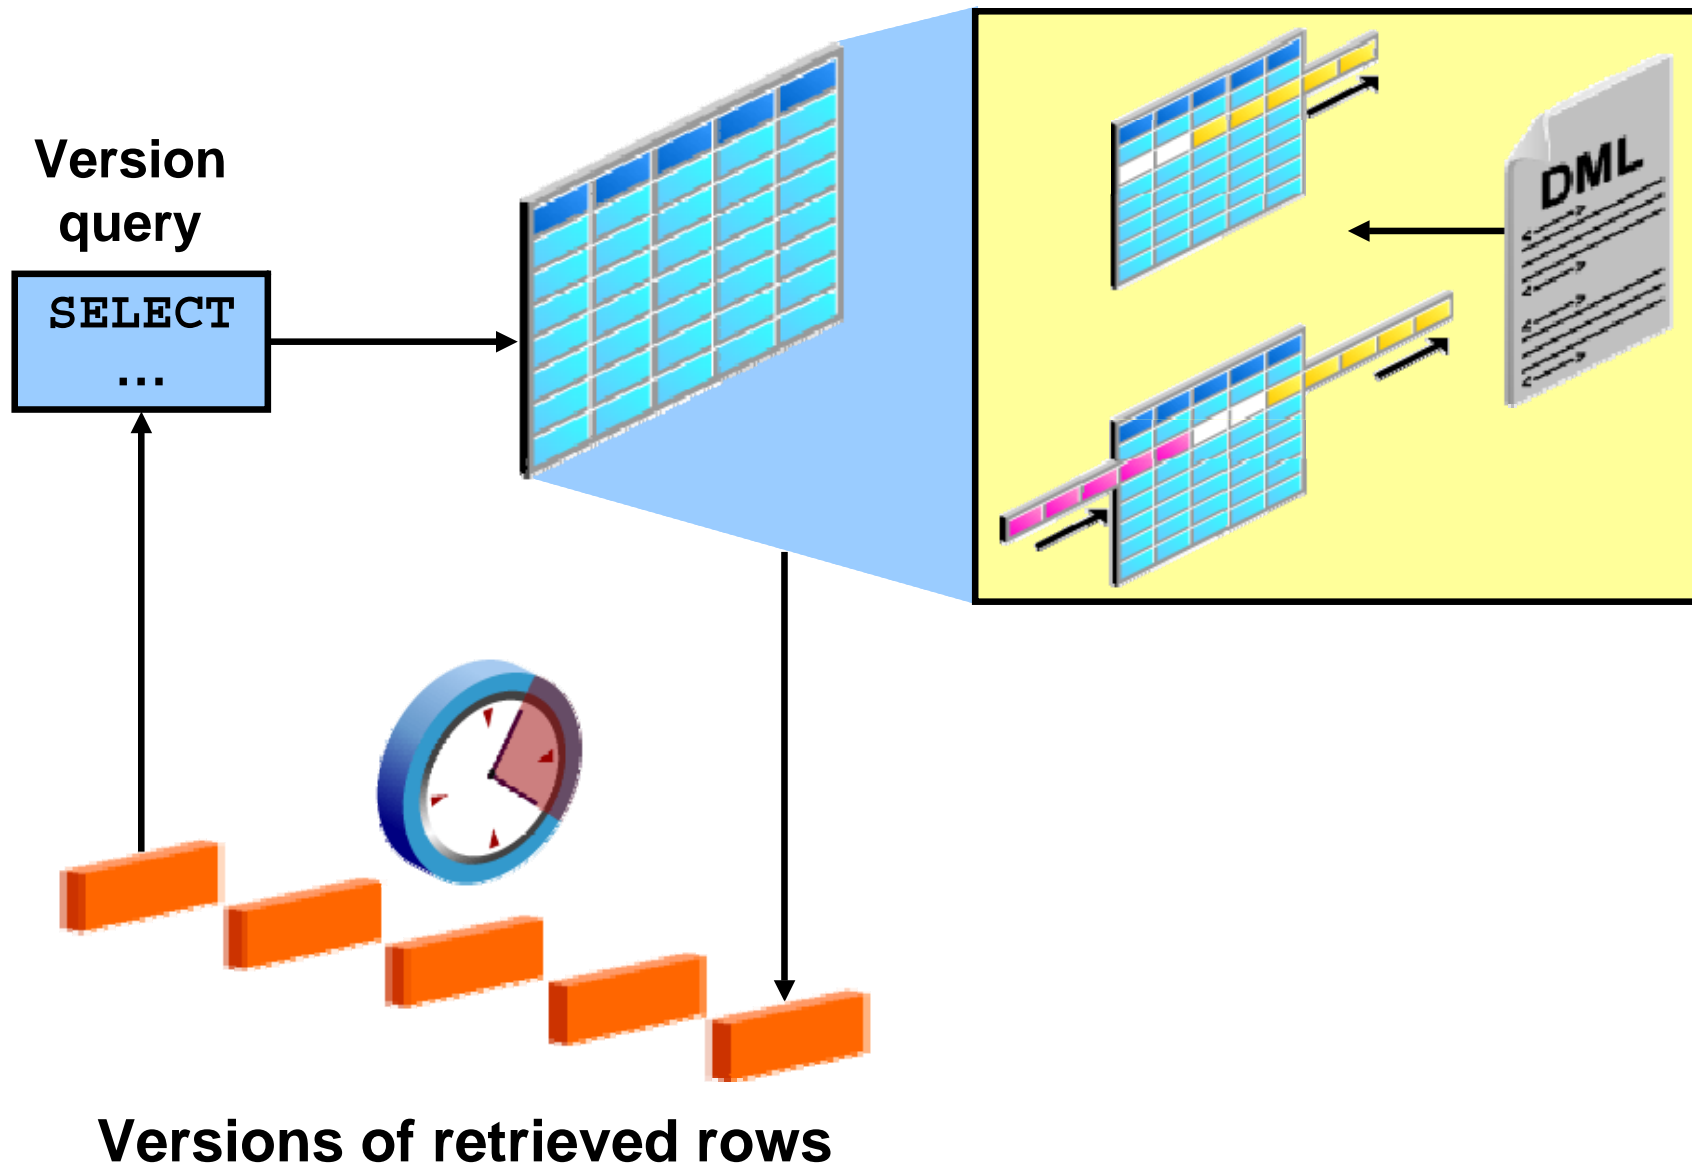
```
MERGE INTO copy_emp3 c  
USING (SELECT * FROM EMPLOYEES ) e  
ON (c.employee_id = e.employee_id)  
WHEN MATCHED THEN  
UPDATE SET  
c.first_name = e.first_name,  
c.last_name = e.last_name,  
...  
DELETE WHERE (E.COMMISSION_PCT IS NOT NULL)  
WHEN NOT MATCHED THEN  
INSERT VALUES(e.employee_id, e.first_name, ...
```

```
SELECT * FROM copy_emp3;  
107 rows selected.
```

# Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
  - Unconditional `INSERT`
  - Pivoting `INSERT`
  - Conditional `INSERT ALL`
  - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

# Tracking Changes in Data



# Example of the Flashback Version Query

```
SELECT salary FROM employees3  
WHERE employee_id = 107;
```

1

```
UPDATE employees3 SET salary = salary * 1.30  
WHERE employee_id = 107;
```

2

```
COMMIT;
```

```
SELECT salary FROM employees3  
  VERSIONS BETWEEN SCN MINVALUE AND MAXVALUE  
WHERE employee_id = 107;
```

3

1

	SALARY
1	4200

3

	SALARY
1	5460
2	4200

# VERSIONS BETWEEN Clause

```
SELECT versions_starttime "START_DATE",  
       versions_endtime   "END_DATE",  
       salary  
FROM   employees  
       VERSIONS BETWEEN SCN MINVALUE  
       AND MAXVALUE  
WHERE  last_name = 'Lorentz';
```

	START_DATE	END_DATE	SALARY
1	18-JUN-09 05.07.10.0000000000 PM	(null)	5460
2	(null)	18-JUN-09 05.07.10.0000000000 PM	4200

# Quiz

When you use the `INSERT` or `UPDATE` command, the `DEFAULT` keyword saves you from hard-coding the default value in your programs or querying the dictionary to find it.

1. True
2. False

# Summary

In this lesson, you should have learned how to:

- Use DML statements and control transactions
- Describe the features of multitable INSERTs
- Use the following types of multitable INSERTs:
  - Unconditional INSERT
  - Pivoting INSERT
  - Conditional INSERT ALL
  - Conditional INSERT FIRST
- Merge rows in a table
- Manipulate data by using subqueries
- Track the changes to data over a period of time

## Practice 4: Overview

This practice covers the following topics:

- Performing multitable `INSERTs`
- Performing `MERGE` operations
- Tracking row versions



# 5

## Managing Data in Different Time Zones

# Objectives

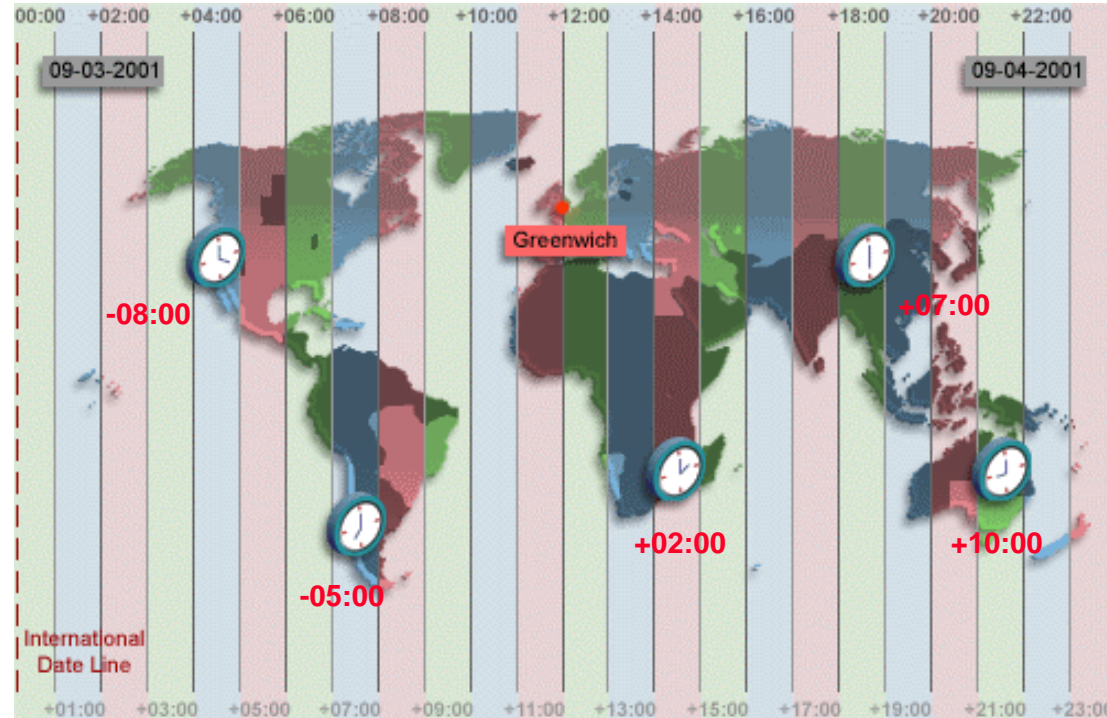
After completing this lesson, you should be able to do the following:

- Use data types similar to `DATE` that store fractional seconds and track time zones
- Use data types that store the difference between two datetime values
- Use the following datetime functions:
  - `CURRENT_DATE`
  - `CURRENT_TIMESTAMP`
  - `LOCALTIMESTAMP`
  - `DBTIMEZONE`
  - `SESSIONTIMEZONE`
  - `EXTRACT`
  - `TZ_OFFSET`
  - `FROM_TZ`
  - `TO_TIMESTAMP`
  - `TO_YMINTERVAL`
  - `TO_DSINTERVAL`

# Lesson Agenda

- `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP`
- `INTERVAL` data types
- Using the following functions:
  - `EXTRACT`
  - `TZ_OFFSET`
  - `FROM_TZ`
  - `TO_TIMESTAMP`
  - `TO_YMINTERVAL`
  - `TO_DSINTERVAL`

# Time Zones



**The image represents the time for each time zone when Greenwich time is 12:00.**

# TIME\_ZONE Session Parameter

TIME\_ZONE may be set to:

- An absolute offset
- Database time zone
- OS local time zone
- A named region

```
ALTER SESSION SET TIME_ZONE = '-05:00';  
ALTER SESSION SET TIME_ZONE = dbtimezone;  
ALTER SESSION SET TIME_ZONE = local;  
ALTER SESSION SET TIME_ZONE = 'America/New_York';
```

# **CURRENT\_DATE, CURRENT\_TIMESTAMP, and LOCALTIMESTAMP**

- **CURRENT\_DATE:**
  - Returns the current date from the user session
  - Has a data type of DATE
- **CURRENT\_TIMESTAMP:**
  - Returns the current date and time from the user session
  - Has a data type of TIMESTAMP WITH TIME ZONE
- **LOCALTIMESTAMP:**
  - Returns the current date and time from the user session
  - Has a data type of TIMESTAMP

# Comparing Date and Time in a Session's Time Zone

The `TIME_ZONE` parameter is set to `-5:00` and then `SELECT` statements for each date and time are executed to compare differences.

```
ALTER SESSION  
SET NLS_DATE_FORMAT = 'DD-MON-YYYY HH24:MI:SS';  
ALTER SESSION SET TIME_ZONE = '-5:00';
```

```
SELECT SESSIONTIMEZONE, CURRENT_DATE FROM DUAL;
```

1

```
SELECT SESSIONTIMEZONE, CURRENT_TIMESTAMP FROM DUAL;
```

2

```
SELECT SESSIONTIMEZONE, LOCALTIMESTAMP FROM DUAL;
```

3

# Comparing Date and Time in a Session's Time Zone

Results of queries:

```
ALTER SESSION succeeded.
```

	SESSIONTIMEZONE	CURRENT_DATE
1	-05:00	23-JUN-2009 01:34:52

1

	SESSIONTIMEZONE	CURRENT_TIMESTAMP
1	-05:00	23-JUN-09 01.35.26.239882000 AM -05:00

2

	SESSIONTIMEZONE	LOCALTIMESTAMP
1	-05:00	23-JUN-09 01.36.21.811798000 AM

3



# DBTIMEZONE and SESSIONTIMEZONE

- Display the value of the database time zone:

```
SELECT DTIMEZONE FROM DUAL;
```

	DBTIMEZONE
1	+00:00

- Display the value of the session's time zone:

```
SELECT SESSIONTIMEZONE FROM DUAL;
```

	SESSIONTIMEZONE
1	-05:00

# TIMESTAMP Data Types

Data Type	Fields
TIMESTAMP	Year, Month, Day, Hour, Minute, Second with fractional seconds
TIMESTAMP WITH TIME ZONE	Same as the TIMESTAMP data type; also includes:  TIMEZONE_HOUR, and TIMEZONE_MINUTE or TIMEZONE_REGION
TIMESTAMP WITH LOCAL TIME ZONE	Same as the TIMESTAMP data type; also includes a time zone offset in its value

# TIMESTAMP Fields

Datetime Field	Valid Values
YEAR	–4712 to 9999 (excluding year 0)
MONTH	01 to 12
DAY	01 to 31
HOURL	00 to 23
MINUTE	00 to 59
SECOND	00 to 59.9(N) where 9(N) is precision
TIMEZONE_HOUR	–12 to 14
TIMEZONE_MINUTE	00 to 59

# Difference Between DATE and TIMESTAMP

A

```
-- when hire_date is  
of type DATE
```

```
SELECT hire_date  
FROM employees;
```

	HIRE_DATE
1	21-JUN-99
2	13-JAN-00
3	17-SEP-87
4	17-FEB-96
5	17-AUG-97
6	07-JUN-94
7	07-JUN-94
8	07-JUN-94

B

```
ALTER TABLE employees  
MODIFY hire_date TIMESTAMP;
```

```
SELECT hire_date  
FROM employees;
```

	HIRE_DATE
1	21-JUN-99 12.00.00.000000000 AM
2	13-JAN-00 12.00.00.000000000 AM
3	17-SEP-87 12.00.00.000000000 AM
4	17-FEB-96 12.00.00.000000000 AM
5	17-AUG-97 12.00.00.000000000 AM
6	07-JUN-94 12.00.00.000000000 AM
7	07-JUN-94 12.00.00.000000000 AM
8	07-JUN-94 12.00.00.000000000 AM

...

# Comparing TIMESTAMP Data Types

```
CREATE TABLE web_orders  
(order_date TIMESTAMP WITH TIME ZONE,  
 delivery_time TIMESTAMP WITH LOCAL TIME ZONE);
```

```
INSERT INTO web_orders values  
(current_date, current_timestamp + 2);
```

```
SELECT * FROM web_orders;
```

	ORDER_DATE	DELIVERY_TIME
1	23-JUN-09 01.56.39.000000000 AM -05:00	25-JUN-09 01.56.39.000000000 AM

# Lesson Agenda

- CURRENT\_DATE, CURRENT\_TIMESTAMP, and LOCALTIMESTAMP
- INTERVAL data types
- Using the following functions:
  - EXTRACT
  - TZ\_OFFSET
  - FROM\_TZ
  - TO\_TIMESTAMP
  - TO\_YMINTERVAL
  - TO\_DSINTERVAL

# INTERVAL Data Types

- INTERVAL data types are used to store the difference between two datetime values.
- There are two classes of intervals:
  - Year-month
  - Day-time
- The precision of the interval is:
  - The actual subset of fields that constitutes an interval
  - Specified in the interval qualifier

Data Type	Fields
INTERVAL YEAR TO MONTH	Year, Month
INTERVAL DAY TO SECOND	Days, Hour, Minute, Second with fractional seconds

# INTERVAL Fields

INTERVAL Field	Valid Values for Interval
YEAR	Any positive or negative integer
MONTH	00 to 11
DAY	Any positive or negative integer
HOURL	00 to 23
MINUTE	00 to 59
SECOND	00 to 59.9(N) where 9(N) is precision



# INTERVAL YEAR TO MONTH: Example

```
CREATE TABLE warranty
(prod_id number,  warranty_time INTERVAL YEAR(3) TO
MONTH);

INSERT INTO warranty VALUES (123, INTERVAL '8' MONTH);
INSERT INTO warranty VALUES (155, INTERVAL '200'
YEAR(3));
INSERT INTO warranty VALUES (678, '200-11');
SELECT * FROM warranty;
```

	PROD_ID	WARRANTY_TIME
1	123	0-8
2	155	200-0
3	678	200-11

# INTERVAL DAY TO SECOND

## Data Type: Example

```
CREATE TABLE lab
( exp_id number, test_time INTERVAL DAY(2) TO SECOND);

INSERT INTO lab VALUES (100012, '90 00:00:00');
INSERT INTO lab VALUES (56098,
    INTERVAL '6 03:30:16' DAY TO SECOND);
```

```
SELECT * FROM lab;
```

	EXP_ID	TEST_TIME
1	100012	90 0:0:0.0
2	56098	6 3:30:16.0

# Lesson Agenda

- CURRENT\_DATE, CURRENT\_TIMESTAMP, and LOCALTIMESTAMP
- INTERVAL data types
- Using the following functions:
  - EXTRACT
  - TZ\_OFFSET
  - FROM\_TZ
  - TO\_TIMESTAMP
  - TO\_YMINTERVAL
  - TO\_DSINTERVAL

# EXTRACT

- Display the YEAR component from the SYSDATE.

```
SELECT EXTRACT (YEAR FROM SYSDATE) FROM DUAL;
```

	EXTRACT(YEARFROMSYSDATE)
1	2009

- Display the MONTH component from the HIRE\_DATE for those employees whose MANAGER\_ID is 100.

```
SELECT last name, hire date,  
       EXTRACT (MONTH FROM HIRE_DATE)  
FROM employees  
WHERE manager_id = 100;
```

	LAST_NAME	HIRE_DATE	EXTRACT(MONTHFROMHIRE_DATE)
1	Hartstein	17-FEB-1996 00:00:00	2
2	Kochhar	21-SEP-1989 00:00:00	9
3	De Haan	13-JAN-1993 00:00:00	1
4	Raphaely	07-DEC-1994 00:00:00	12
5	Weiss	18-JUL-1996 00:00:00	7

# TZ\_OFFSET

Display the time zone offset for the 'US/Eastern',  
'Canada/Yukon' and 'Europe/London' time zones:


```
SELECT TZ_OFFSET('US/Eastern'),  
       TZ_OFFSET('Canada/Yukon'),  
       TZ_OFFSET('Europe/London')  
FROM DUAL;
```

	TZ_OFFSET('US/EASTERN')	TZ_OFFSET('CANADA/YUKON')	TZ_OFFSET('EUROPE/LONDON')
1	-04:00	-07:00	+01:00

## FROM\_TZ

Display the `TIMESTAMP` value `'2000-03-28 08:00:00'` as a `TIMESTAMP WITH TIME ZONE` value for the `'Australia/North'` time zone region.

```
SELECT FROM_TZ(TIMESTAMP
                '2000-07-12 08:00:00', 'Australia/North')
FROM DUAL;
```

 FROM_TZ(TIMESTAMP'2000-07-1208:00:00','AUSTRALIA/NORTH')
1 12-JUL-00 08.00.00.000000000 AM AUSTRALIA/NORTH

# TO\_TIMESTAMP

Display the character string '2007-03-06 11:00:00'  
as a TIMESTAMP value:

```
SELECT TO_TIMESTAMP ('2007-03-06 11:00:00',  
                     'YYYY-MM-DD HH:MI:SS')  
FROM DUAL;
```

```
TO_TIMESTAMP('2007-03-0611:00:00','YYYY-MM-DDHH:MI:SS')  
06-MAR-07 11.00.00.000000000
```

## TO\_YMINTERVAL

Display a date that is one year and two months after the hire date for the employees working in the department with the DEPARTMENT\_ID 20.

```
SELECT hire_date,  
       hire_date + TO_YMINTERVAL('01-02') AS  
       HIRE_DATE_YMININTERVAL  
FROM   employees  
WHERE  department_id = 20;
```

	HIRE_DATE	HIRE_DATE_YMININTERVAL
1	17-FEB-1996 00:00:00	17-APR-1997 00:00:00
2	17-AUG-1997 00:00:00	17-OCT-1998 00:00:00



## TO\_DSINTERVAL

Display a date that is 100 days and 10 hours after the hire date for all the employees.

```
SELECT last_name,  
       TO_CHAR(hire_date, 'mm-dd-yy:hh:mi:ss') hire_date,  
       TO_CHAR(hire_date +  
               TO_DSINTERVAL('100 10:00:00'),  
               'mm-dd-yy:hh:mi:ss') hiredate2  
FROM employees;
```

	LAST_NAME	HIRE_DATE	HIREDATE2
1	OConnell	06-21-99:12:00:00	09-29-99:10:00:00
2	Grant	01-13-00:12:00:00	04-22-00:10:00:00
3	Whalen	09-17-87:12:00:00	12-26-87:10:00:00
4	Hartstein	02-17-96:12:00:00	05-27-96:10:00:00
5	Fay	08-17-97:12:00:00	11-25-97:10:00:00
6	Mavris	06-07-94:12:00:00	09-15-94:10:00:00
7	Baer	06-07-94:12:00:00	09-15-94:10:00:00
8	Higgins	06-07-94:12:00:00	09-15-94:10:00:00

...

# Daylight Saving Time

- First Sunday in April
  - Time jumps from 01:59:59 AM to 03:00:00 AM.
  - Values from 02:00:00 AM to 02:59:59 AM are not valid.
- Last Sunday in October
  - Time jumps from 02:00:00 AM to 01:00:01 AM.
  - Values from 01:00:01 AM to 02:00:00 AM are ambiguous because they are visited twice.

# Quiz

The `TIME_ZONE` session parameter may be set to:

1. A relative offset
2. Database time zone
3. OS local time zone
4. A named region

# Summary

In this lesson, you should have learned how to use the following functions:

- `CURRENT_DATE`
- `CURRENT_TIMESTAMP`
- `LOCALTIMESTAMP`
- `DBTIMEZONE`
- `SESSIONTIMEZONE`
- `EXTRACT`
- `TZ_OFFSET`
- `FROM_TZ`
- `TO_TIMESTAMP`
- `TO_YMINTERVAL`
- `TO_DSINTERVAL`

# Practice 5: Overview

This practice covers using the datetime functions.

# 6

## Retrieving Data by Using Subqueries

# Objectives

After completing this lesson, you should be able to do the following:


- Write a multiple-column subquery
- Use scalar subqueries in SQL
- Solve problems with correlated subqueries
- Update and delete rows by using correlated subqueries
- Use the `EXISTS` and `NOT EXISTS` operators
- Use the `WITH` clause

# Lesson Agenda

- Writing a multiple-column subquery
- Using scalar subqueries in SQL
- Solving problems with correlated subqueries
- Using the EXISTS and NOT EXISTS operators
- Using the WITH clause



# Multiple-Column Subqueries



**Each row of the main query is compared to values from a multiple-row and multiple-column subquery.**

# Column Comparisons

Multiple-column comparisons involving subqueries can be:

- Nonpairwise comparisons
- Pairwise comparisons

# Pairwise Comparison Subquery

Display the details of the employees who are managed by the same manager and work in the same department as employees with the first name of “John.”

```
SELECT employee_id, manager_id, department_id
FROM   empl_demo
WHERE  (manager_id, department_id) IN
      (SELECT manager_id, department_id
       FROM empl_demo
       WHERE first_name = 'John')
AND first_name <> 'John';
```

# Nonpairwise Comparison Subquery

Display the details of the employees who are managed by the same manager as the employees with the first name of “John” and work in the same department as the employees with the first name of “John.”

```
SELECT  employee_id, manager_id, department_id
FROM    empl_demo
WHERE   manager_id IN
        (SELECT manager_id
         FROM empl_demo
         WHERE first_name = 'John')
AND department_id IN
        (SELECT department_id
         FROM empl_demo
         WHERE first_name = 'John')
AND first_name <> 'John';
```

# Lesson Agenda


- Writing a multiple-column subquery
- **Using scalar subqueries in SQL**
- Solving problems with correlated subqueries
- Using the EXISTS and NOT EXISTS operators
- Using the WITH clause

# Scalar Subquery Expressions

- A scalar subquery expression is a subquery that returns exactly one column value from one row.
- Scalar subqueries can be used in:
  - The condition and expression part of `DECODE` and `CASE`
  - All clauses of `SELECT` except `GROUP BY`
  - The `SET` clause and `WHERE` clause of an `UPDATE` statement

# Scalar Subqueries: Examples

- Scalar subqueries in CASE expressions:

```
SELECT employee_id, last_name,  
       (CASE  
         WHEN department_id = 20  
           (SELECT department_id  
            FROM departments  
            WHERE location_id = 1800)  
         THEN 'Canada' ELSE 'USA' END) location  
FROM   employees;
```

- Scalar subqueries in the ORDER BY clause:

```
SELECT   employee_id, last_name  
FROM     employees e  
ORDER BY (SELECT department_name  
          FROM departments d  
          WHERE e.department_id = d.department_id);
```


# Lesson Agenda

- Writing a multiple-column subquery
- Using scalar subqueries in SQL
- Solving problems with correlated subqueries
- Using the EXISTS and NOT EXISTS operators
- Using the WITH clause



# Correlated Subqueries

Correlated subqueries are used for row-by-row processing. Each subquery is executed once for every row of the outer query.



# Correlated Subqueries


The subquery references a column from a table in the parent query.

```
SELECT column1, column2, ...
FROM   table1 Outer_table
WHERE  column1 operator
              (SELECT column1, column2
                FROM   table2
                WHERE  expr1 =
                      Outer_table.expr2);
```

# Using Correlated Subqueries

Find all employees who earn more than the average salary in their department.

```
SELECT last_name, salary, department_id
FROM   employees outer_table
WHERE  salary >
      (SELECT AVG(salary)
       FROM   employees inner_table
       WHERE  inner_table.department_id =
             outer_table.department_id);
```



Each time a row from the outer query is processed, the inner query is evaluated.

# Using Correlated Subqueries

Display details of those employees who have changed jobs at least twice.

```
SELECT e.employee_id, last_name, e.job_id
FROM   employees e
WHERE  2 <= (SELECT COUNT(*)
              FROM   job_history
              WHERE  employee_id = e.employee_id);
```

	EMPLOYEE_ID	LAST_NAME	JOB_ID
1	200	Whalen	AD_ASST
2	101	Kochhar	AD_VP
3	176	Taylor	SA_REP

# Lesson Agenda

- Writing a multiple-column subquery
- Using scalar subqueries in SQL
- Solving problems with correlated subqueries
- **Using the EXISTS and NOT EXISTS operators**
- Using the WITH clause

# Using the EXISTS Operator

- The EXISTS operator tests for existence of rows in the results set of the subquery.
- If a subquery row value is found:
  - The search does not continue in the inner query
  - The condition is flagged TRUE
- If a subquery row value is not found:
  - The condition is flagged FALSE
  - The search continues in the inner query

# Using the EXISTS Operator

```
SELECT employee_id, last_name, job_id, department_id
FROM   employees outer
WHERE  EXISTS ( SELECT 'X'
                FROM   employees
                WHERE  manager_id =
                      outer.employee_id);
```

	EMPLOYEE_ID	LAST_NAME	JOB_ID	DEPARTMENT_ID
1	201	Hartstein	MK_MAN	20
2	205	Higgins	AC_MGR	110
3	100	King	AD_PRES	90
4	101	Kochhar	AD_VP	90
5	102	De Haan	AD_VP	90
6	103	Hunold	IT_PROG	60
7	108	Greenberg	FI_MGR	100
8	114	Raphaely	PU_MAN	30

# Find All Departments That Do Not Have Any Employees

```
SELECT department_id, department_name
FROM departments d
WHERE NOT EXISTS (SELECT 'X'
                  FROM employees
                  WHERE department_id = d.department_id);
```

	DEPARTMENT_ID	DEPARTMENT_NAME
1	120	Treasury
2	130	Corporate Tax
3	140	Control And Credit
4	150	Shareholder Services
5	160	Benefits
6	170	Manufacturing
7	180	Construction

...

All Rows Fetched: 16



# Correlated UPDATE

Use a correlated subquery to update rows in one table based on rows from another table.

```
UPDATE table1 alias1
SET    column = (SELECT expression
                     FROM   table2 alias2
                     WHERE  alias1.column =
                           alias2.column);
```

# Using Correlated UPDATE

- Denormalize the EMPL6 table by adding a column to store the department name.
- Populate the table by using a correlated update.

```
ALTER TABLE empl6  
ADD(department_name VARCHAR2(25));
```

```
UPDATE empl6 e  
SET     department_name =  
        (SELECT department_name  
         FROM   departments d  
         WHERE  e.department_id = d.department_id);
```

# Correlated DELETE

Use a correlated subquery to delete rows in one table based on rows from another table.

```
DELETE FROM table1 alias1
WHERE  column operator
        (SELECT expression
         FROM   table2 alias2
         WHERE  alias1.column = alias2.column);
```

# Using Correlated DELETE

Use a correlated subquery to delete only those rows from the EMPL6 table that also exist in the EMP\_HISTORY table.

```
DELETE FROM empl6 E
WHERE employee_id =
      (SELECT employee_id
       FROM   emp_history
       WHERE  employee_id = E.employee_id);
```

# Lesson Agenda

- Writing a multiple-column subquery
- Using scalar subqueries in SQL
- Solving problems with correlated subqueries
- Using the EXISTS and NOT EXISTS operators
- **Using the WITH clause**

# WITH Clause

- Using the `WITH` clause, you can use the same query block in a `SELECT` statement when it occurs more than once within a complex query.
- The `WITH` clause retrieves the results of a query block and stores it in the user's temporary tablespace.
- The `WITH` clause may improve performance.

## WITH Clause: Example

Using the `WITH` clause, write a query to display the department name and total salaries for those departments whose total salary is greater than the average salary across departments.

# WITH Clause: Example


```
WITH
dept_costs AS (
    SELECT d.department_name, SUM(e.salary) AS dept_total
    FROM   employees e JOIN departments d
    ON     e.department_id = d.department_id
    GROUP BY d.department_name),
avg_cost AS (
    SELECT SUM(dept_total)/COUNT(*) AS dept_avg
    FROM   dept_costs)
SELECT *
FROM   dept_costs
WHERE  dept_total >
      (SELECT dept_avg
       FROM avg_cost)
ORDER BY department_name;
```



# Recursive WITH Clause

## The Recursive WITH clause

- Enables formulation of recursive queries.
- Creates query with a name, called the Recursive WITH element name
- Contains two types of query blocks member: anchor and a recursive
- Is ANSI-compatible



# Recursive WITH Clause: Example

FLIGHTS Table

R2	SOURCE	R2	DESTIN	R2	FLIGHT_TIME
1	San Jose		Los Angeles		1.3
2	New York		Boston		1.1
3	Los Angeles		New York		5.8

1

```
WITH Reachable_From (Source, Destin, TotalFlightTime) AS
(
    SELECT Source, Destin, Flight_time
    FROM Flights
    UNION ALL
    SELECT incoming.Source, outgoing.Destin,
           incoming.TotalFlightTime+outgoing.Flight_time
    FROM Reachable_From incoming, Flights outgoing
    WHERE incoming.Destin = outgoing.Source
)
SELECT Source, Destin, TotalFlightTime
FROM Reachable_From;
```

2

3

R2	SOURCE	R2	DESTIN	R2	TOTALFLIGHTTIME
1	San Jose		Los Angeles		1.3
2	New York		Boston		1.1
3	Los Angeles		New York		5.8
4	San Jose		New York		7.1
5	Los Angeles		Boston		6.9
6	San Jose		Boston		8.2

# Quiz

With a correlated subquery, the inner `SELECT` statement drives the outer `SELECT` statement.

1. True
2. False

# Summary

In this lesson, you should have learned that:

- A multiple-column subquery returns more than one column
- Multiple-column comparisons can be pairwise or nonpairwise
- A multiple-column subquery can also be used in the `FROM` clause of a `SELECT` statement

# Summary

- Correlated subqueries are useful whenever a subquery must return a different result for each candidate row
- The `EXISTS` operator is a Boolean operator that tests the presence of a value
- Correlated subqueries can be used with `SELECT`, `UPDATE`, and `DELETE` statements
- You can use the `WITH` clause to use the same query block in a `SELECT` statement when it occurs more than once

# Practice 6: Overview

This practice covers the following topics:

- Creating multiple-column subqueries
- Writing correlated subqueries
- Using the EXISTS operator
- Using scalar subqueries
- Using the WITH clause



# **Regular Expression Support**

# Objectives

After completing this lesson, you should be able to do the following:

- List the benefits of using regular expressions
- Use regular expressions to search for, match, and replace strings



# Lesson Agenda

- Introduction to regular expressions
- Using metacharacters with regular expressions
- Using the regular expressions functions:
  - REGEXP\_LIKE
  - REGEXP\_REPLACE
  - REGEXP\_INSTR
  - REGEXP\_SUBSTR
- Accessing subexpressions
- Using the REGEXP\_COUNT function
- Regular expressions and check constraints

# What Are Regular Expressions?

- You use regular expressions to search for (and manipulate) simple and complex patterns in string data by using standard syntax conventions.
- You use a set of SQL functions and conditions to search for and manipulate strings in SQL and PL/SQL.
- You specify a regular expression by using:
  - Metacharacters, which are operators that specify the search algorithms
  - Literals, which are the characters for which you are searching

# Benefits of Using Regular Expressions

Regular expressions enable you to implement complex match logic in the database with the following benefits:

- By centralizing match logic in Oracle Database, you avoid intensive string processing of SQL results sets by middle-tier applications.
- Using server-side regular expressions to enforce constraints, you eliminate the need to code data validation logic on the client.
- The built-in SQL and PL/SQL regular expression functions and conditions make string manipulations more powerful and easier than in previous releases of Oracle Database 11g.

# Using the Regular Expressions Functions and Conditions in SQL and PL/SQL

Function or Condition Name	Description
REGEXP_LIKE	Is similar to the LIKE operator, but performs regular expression matching instead of simple pattern matching (condition)
REGEXP_REPLACE	Searches for a regular expression pattern and replaces it with a replacement string
REGEXP_INSTR	Searches a string for a regular expression pattern and returns the position where the match is found
REGEXP_SUBSTR	Searches for a regular expression pattern within a given string and extracts the matched substring
REGEXP_COUNT	Returns the number of times a pattern match is found in an input sting

# Lesson Agenda

- Introduction to regular expressions
- Using metacharacters with regular expressions
- Using the regular expressions functions:
  - REGEXP\_LIKE
  - REGEXP\_REPLACE
  - REGEXP\_INSTR
  - REGEXP\_SUBSTR
- Accessing subexpressions
- Using the REGEXP\_COUNT function

# What Are Metacharacters?

- Metacharacters are special characters that have a special meaning such as a wildcard, a repeating character, a nonmatching character, or a range of characters.
- You can use several predefined metacharacter symbols in the pattern matching.
- For example, the `^(f|ht)tps?:$` regular expression searches for the following from the beginning of the string:
  - The literals `f` or `ht`
  - The `t` literal
  - The `p` literal, optionally followed by the `s` literal
  - The colon “:” literal at the end of the string

# Using Metacharacters with Regular Expressions

Syntax	Description
.	Matches any character in the supported character set, except NULL
+	Matches one or more occurrences
?	Matches zero or one occurrence
*	Matches zero or more occurrences of the preceding subexpression
{ <i>m</i> }	Matches exactly <i>m</i> occurrences of the preceding expression
{ <i>m</i> , }	Matches at least <i>m</i> occurrences of the preceding subexpression
{ <i>m</i> , <i>n</i> }	Matches at least <i>m</i> , but not more than <i>n</i> , occurrences of the preceding subexpression
[...]	Matches any single character in the list within the brackets
	Matches one of the alternatives
( . . . )	Treats the enclosed expression within the parentheses as a unit. The subexpression can be a string of literals or a complex expression containing operators.

# Using Metacharacters with Regular Expressions

Syntax	Description
<code>^</code>	Matches the beginning of a string
<code>\$</code>	Matches the end of a string
<code>\</code>	Treats the subsequent metacharacter in the expression as a literal
<code>\n</code>	Matches the <i>n</i> th (1–9) preceding subexpression of whatever is grouped within parentheses. The parentheses cause an expression to be remembered; a backreference refers to it.
<code>\d</code>	A digit character
<code>[ :class:]</code>	Matches any character belonging to the specified POSIX character class
<code>[ ^:class:]</code>	Matches any single character <i>not</i> in the list within the brackets



# Lesson Agenda

- Introduction to regular expressions
- Using metacharacters with regular expressions
- Using the regular expressions functions:
  - REGEXP\_LIKE
  - REGEXP\_REPLACE
  - REGEXP\_INSTR
  - REGEXP\_SUBSTR
- Accessing subexpressions
- Using the REGEXP\_COUNT function

# Regular Expressions Functions and Conditions: Syntax

```
REGEXP_LIKE (source_char, pattern [,match_option]
```

```
REGEXP_INSTR (source_char, pattern [, position  
[, occurrence [, return_option  
[, match_option [, subexpr]]]])
```

```
REGEXP_SUBSTR (source_char, pattern [, position  
[, occurrence [, match_option  
[, subexpr]]]])
```

```
REGEXP_REPLACE(source_char, pattern [,replacestr  
[, position [, occurrence  
[, match_option]]]])
```

```
REGEXP_COUNT (source_char, pattern [, position  
[, occurrence [, match_option]]])
```

# Performing a Basic Search by Using the REGEXP\_LIKE Condition

```
REGEXP_LIKE(source_char, pattern [, match_parameter ])
```

```
SELECT first_name, last_name  
FROM employees  
WHERE REGEXP_LIKE (first_name, '^Ste(v|ph)en$');
```

	FIRST_NAME	LAST_NAME
1	Steven	King
2	Steven	Markle
3	Stephen	Stiles

# Replacing Patterns by Using the REGEXP\_REPLACE Function

```
REGEXP_REPLACE(source_char, pattern [,replacestr  
[, position [, occurrence [, match_option]]])
```

```
SELECT REGEXP_REPLACE(phone_number, '\.','-') AS phone  
FROM employees;
```

**Original**

	LAST_NAME	PHONE
1	OConnell	650.507.9833
2	Grant	650.507.9844
3	Whalen	515.123.4444
4	Hartstein	515.123.5555

**Partial results**

	LAST_NAME	PHONE
1	OConnell	650-507-9833
2	Grant	650-507-9844
3	Whalen	515-123-4444
4	Hartstein	515-123-5555

# Finding Patterns by Using the REGEXP\_INSTR Function

```
REGEXP_INSTR (source_char, pattern [, position [,  
occurrence [, return_option [, match_option]]])
```

```
SELECT street_address,  
REGEXP_INSTR(street_address,'[[:alpha:]]') AS  
    First_Alpha_Position  
FROM locations;
```

	STREET_ADDRESS	FIRST_ALPHA_POSITION
1	1297 Via Cola di Rie	6
2	93091 Calle della Testa	7
3	2017 Shinjuku-ku	6
4	9450 Kamiya-cho	6

# Extracting Substrings by Using the REGEXP\_SUBSTR Function

```
REGEXP_SUBSTR (source_char, pattern [, position  
                [, occurrence [, match_option]])
```

```
SELECT REGEXP_SUBSTR(street_address , ' [^ ]+ ') AS Road  
FROM locations;
```

	ROAD
1	Via
2	Calle
3	(null)
4	(null)
5	Jabberwocky

# Lesson Agenda


- Introduction to regular expressions
- Using metacharacters with regular expressions
- Using the regular expressions functions:
  - REGEXP\_LIKE
  - REGEXP\_REPLACE
  - REGEXP\_INSTR
  - REGEXP\_SUBSTR
- **Accessing subexpressions**
- Using the REGEXP\_COUNT function

# Subexpressions

Examine this expression:

```
(1 2 3) (4 (5 6) (7 8))
```

The subexpressions are:





# Using Subexpressions with Regular Expression Support

```
SELECT
  REGEXP_INSTR
① ('0123456789',      -- source char or search value
② '(123)(4(56)(78))', -- regular expression patterns
③ 1,                  -- position to start searching
④ 1,                  -- occurrence
⑤ 0,                  -- return option
⑥ 'i',                -- match option (case insensitive)
⑦ 1)                  -- sub-expression on which to search
    "Position"
FROM dual;
```

RZ	Position
1	2

# Why Access the *n*th Subexpression?

- A more realistic use: DNA sequencing
- You may need to find a specific subpattern that identifies a protein needed for immunity in mouse DNA.

```
SELECT
```

```
  REGEXP_INSTR('ccacctttccctccactcctcacgttctcacctgtaaagcgtccctc  
cctcatccccatgcccccttaccctgcagggtagagtaggctagaaaccagagagctccaagc  
tccatctgtggagaggtgccatccttgggctgcagagagaggagaatttgccccaagctgcc  
tgcagagcttcaccacccttagtctcacaagccttgagttcatagcatttcttgagttttca  
ccctgcccagcaggacactgcagcacccaaagggcttcccaggagtaggggttgccctcaagag  
gctcttgggtctgatggccacatcctggaattgttttcaagttgatggtcacagccctgaggc  
atgtagggggcgtggggatgcgctctgctctgctctcctctcctgaaccctgaaccctctggc  
taccacagagcacttagagccag',
```

```
        '(gtc(tcac)(aaag))',  
        1, 1, 0, 'i',  
        1) "Position"
```

```
FROM dual;
```

Position
1 195

# REGEXP\_SUBSTR: Example

```
SELECT
  REGEXP_SUBSTR
    ① ('acgctgcactgca', -- source char or search value
    ② 'acg(.*)gca',      -- regular expression pattern
    ③ 1,                 -- position to start searching
    ④ 1,                 -- occurrence
    ⑤ 'i',               -- match option (case insensitive)
    ⑥ 1)                -- sub-expression
    "Value"
FROM dual;
```

	Value
1	ctgcact

# Lesson Agenda

- Introduction to regular expressions
- Using metacharacters with regular expressions
- Using the regular expressions functions:
  - REGEXP\_LIKE
  - REGEXP\_REPLACE
  - REGEXP\_INSTR
  - REGEXP\_SUBSTR
- Accessing subexpressions
- Using the REGEXP\_COUNT function

# Using the REGEXP\_COUNT Function

```
REGEXP_COUNT (source_char, pattern [, position  
              [, occurrence [, match_option]])
```

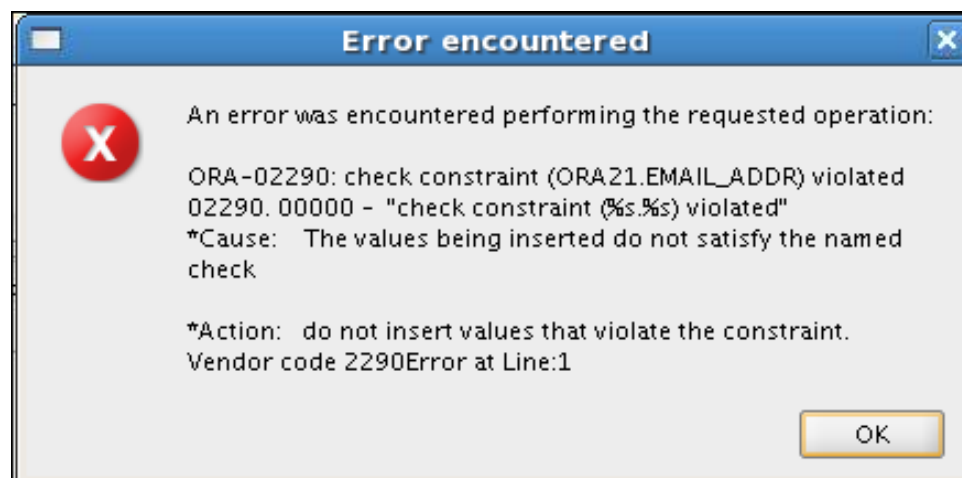
```
SELECT REGEXP_COUNT(  
    'ccacctttccctccactcctcacgttctcacctgtaaagcgtccctccctcatcccatgcccccttaccctgcag  
    ggtagagtaggctagaaaccagagagctccaagctccatctgtggagaggtgccatccttgggctgcagagagaggag  
    aatttgcccaaagctgcctgcagagcttcaccacccttagtctcacaagccttgagttcatagcatttcttgagtt  
    ttcacctgcccagcaggacactgcagcacccaaagggcttcccaggagtagggttgccctcaagaggctcttggggtc  
    tgatggccacatcctggaattgttttcaagttgatggtcacagccctgaggcatgtaggggcgtggggatgcgctctg  
    ctctgctctcctctcctgaaccctgaaccctctggctaccccagagcacttagagccag' ,  
    'gtc') AS Count  
  
FROM dual;
```

	REGEXP_COUNT	COUNT
1		4

# Regular Expressions and Check Constraints: Examples

```
ALTER TABLE emp8  
ADD CONSTRAINT email_addr  
CHECK(REGEXP_LIKE(email, '@')) NOVALIDATE;
```

```
INSERT INTO emp8 VALUES  
(500, 'Christian', 'Patel', 'ChrisP2creme.com',  
1234567890, '12-Jan-2004', 'HR_REP', 2000, null, 102, 40);
```



# Quiz

With the use of regular expressions in SQL and PL/SQL, you can:

1. Avoid intensive string processing of SQL result sets by middle-tier applications
2. Avoid data validation logic on the client
3. Enforce constraints on the server

# Summary

In this lesson, you should have learned how to use regular expressions to search for, match, and replace strings.



# Practice 7: Overview

This practice covers using regular expressions functions to do the following:

- Searching for, replacing, and manipulating data
- Creating a new `CONTACTS` table and adding a `CHECK` constraint to the `p_number` column to ensure that phone numbers are entered into the database in a specific standard format
- Testing the adding of some phone numbers into the `p_number` column by using various formats



# **Using SQL Developer**


# Objectives

After completing this appendix, you should be able to do the following:

- List the key features of Oracle SQL Developer
- Identify menu items of Oracle SQL Developer
- Create a database connection
- Manage database objects
- Use SQL Worksheet
- Save and run SQL scripts
- Create and save reports

# What Is Oracle SQL Developer?

- Oracle SQL Developer is a graphical tool that enhances productivity and simplifies database development tasks.
- You can connect to any target Oracle database schema by using standard Oracle database authentication.




**SQL Developer**

# Specifications of SQL Developer

- Shipped along with Oracle Database 11g Release 2
- Developed in Java
- Supports Windows, Linux, and Mac OS X platforms
- Default connectivity by using the Java Database Connectivity (JDBC) thin driver
- Connects to Oracle Database version 9.2.0.1 and later
- Freely downloadable from the following link:
  - [http://www.oracle.com/technology/products/database/sql\\_developer/index.html](http://www.oracle.com/technology/products/database/sql_developer/index.html)


# SQL Developer 1.5 Interface



# Creating a Database Connection

- You must have at least one database connection to use SQL Developer.
- You can create and test connections for multiple:
  - Databases
  - Schemas
- SQL Developer automatically imports any connections defined in the `tnsnames.ora` file on your system.
- You can export connections to an Extensible Markup Language (XML) file.
- Each additional database connection created is listed in the Connections Navigator hierarchy.

# Creating a Database Connection






# Browsing Database Objects

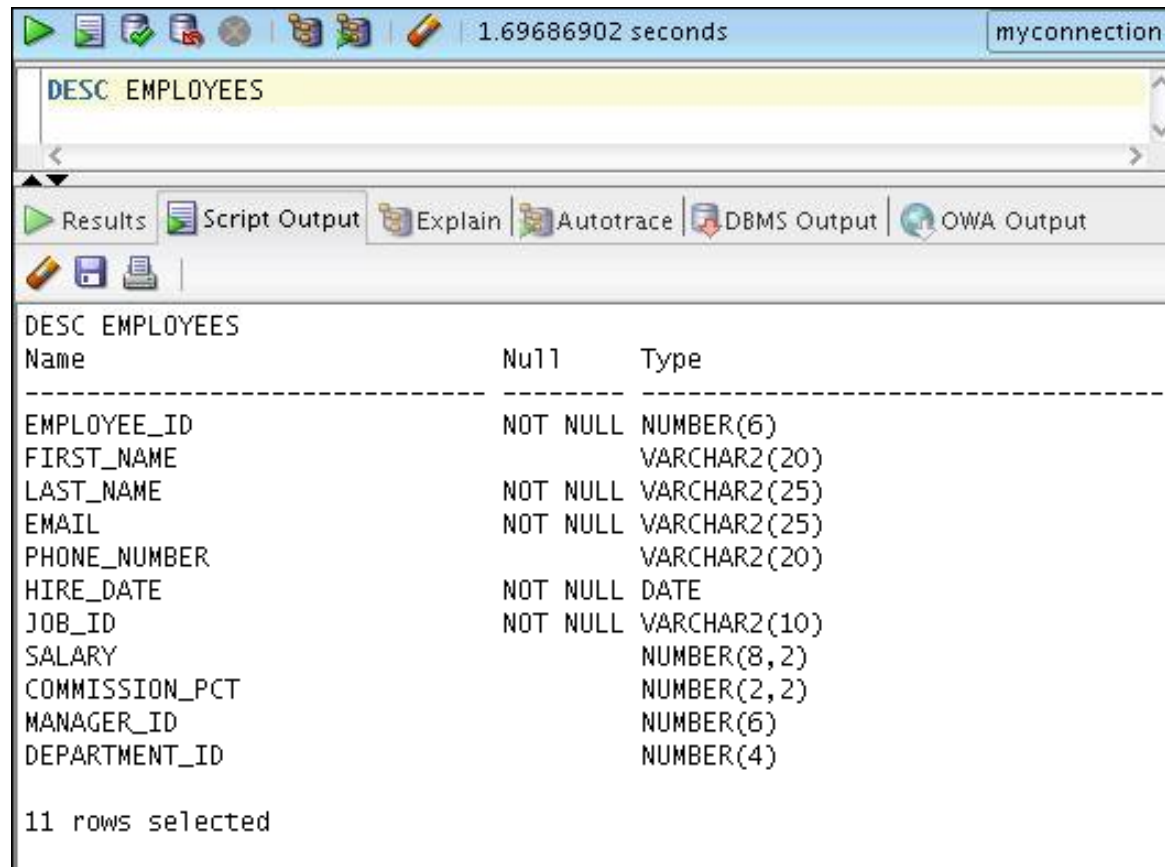
Use the Connections Navigator to to:

- Browse through many objects in a database schema
- Review the definitions of objects at a glance



# Displaying the Table Structure

Use the DESCRIBE command to display the structure of a table:




DESC EMPLOYEES

Name	Null	Type
EMPLOYEE_ID	NOT NULL	NUMBER(6)
FIRST_NAME		VARCHAR2(20)
LAST_NAME	NOT NULL	VARCHAR2(25)
EMAIL	NOT NULL	VARCHAR2(25)
PHONE_NUMBER		VARCHAR2(20)
HIRE_DATE	NOT NULL	DATE
JOB_ID	NOT NULL	VARCHAR2(10)
SALARY		NUMBER(8,2)
COMMISSION_PCT		NUMBER(2,2)
MANAGER_ID		NUMBER(6)
DEPARTMENT_ID		NUMBER(4)

11 rows selected


# Browsing Files

Use the File Navigator to explore the file system and open system files.



# Creating a Schema Object

- SQL Developer supports the creation of any schema object by:
  - Executing a SQL statement in SQL Worksheet
  - Using the context menu
- Edit the objects by using an edit dialog box or one of the many context-sensitive menus.
- View the data definition language (DDL) for adjustments such as creating a new object or editing an existing schema object.



# Creating a New Table: Example

**Create Table**

Schema: ORA21

Name: Dependents

Table Type: ☒ Normal ☐ External ☐ Index Organized ☐ Temporary (Transaction) ☐ Temporary (Session)

**Columns:**

- ID
- FIRST\_NAME
- LAST\_NAME
- RELATION
- BIRTHDATE

**Column Properties**

Name: ID

Datatype: ☒ Simple ☐ Complex

Type: NUMBER

Precision:

Scale:

Default:


☒ Cannot be NULL

Comment:

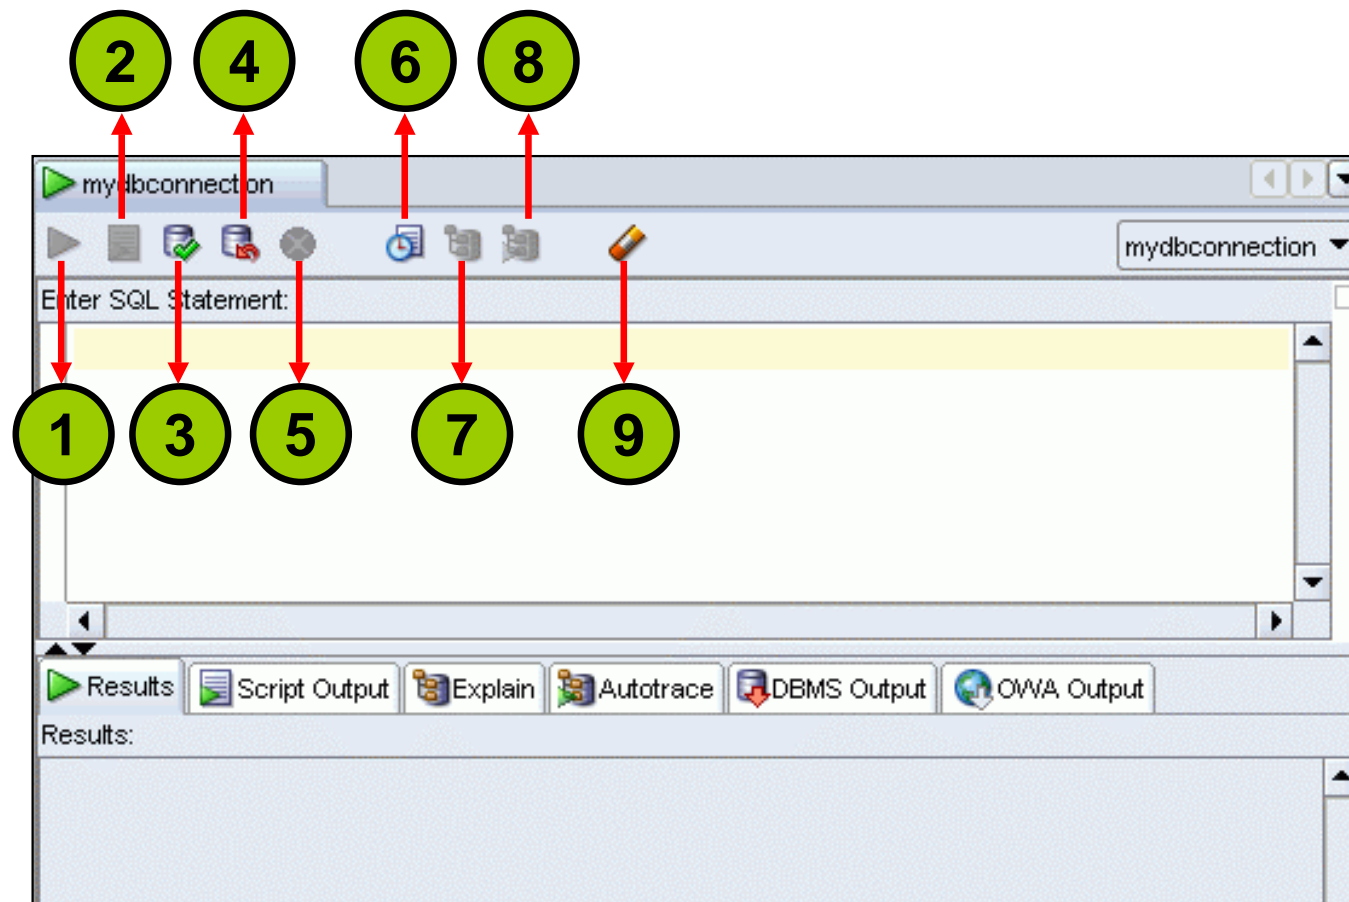
Help OK Cancel

# Using the SQL Worksheet

- Use the SQL Worksheet to enter and execute SQL, PL/SQL, and SQL \*Plus statements.
- Specify any actions that can be processed by the database connection associated with the worksheet.

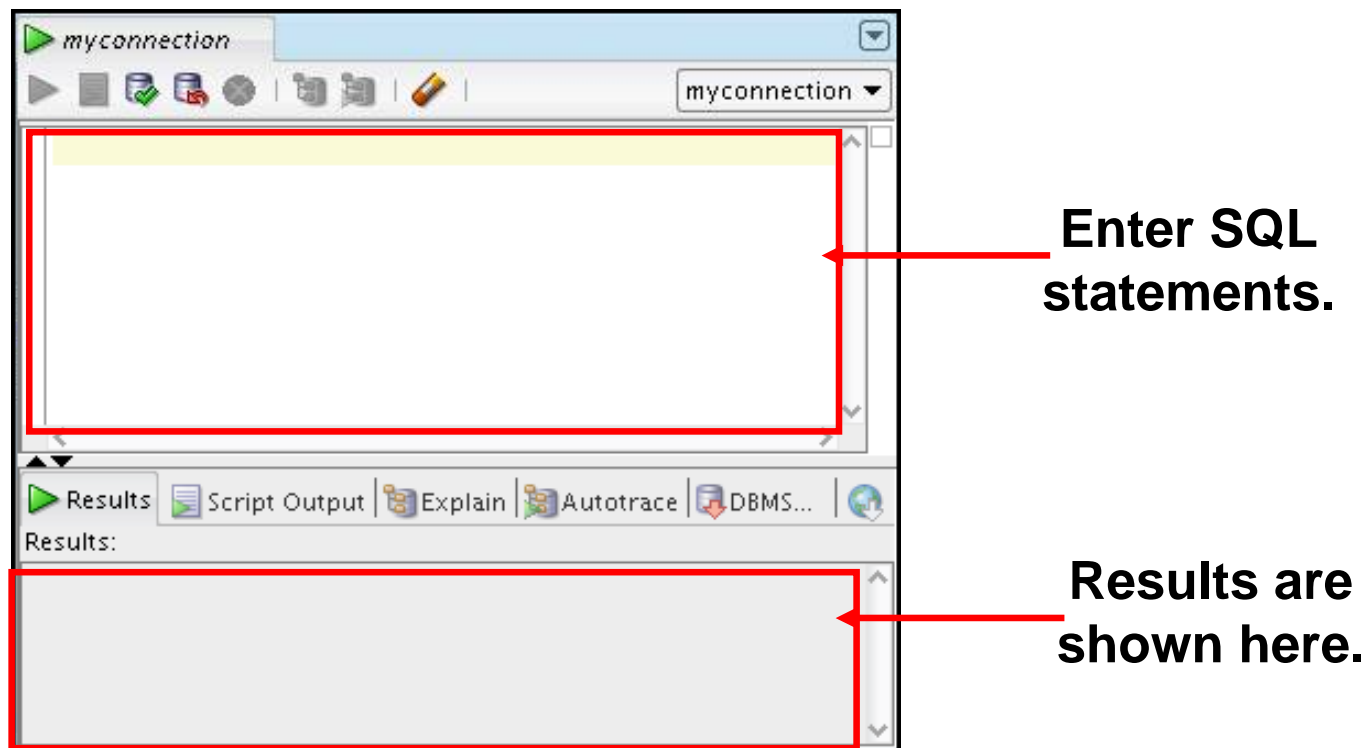


# Using the SQL Worksheet



# Using the SQL Worksheet

- Use the SQL Worksheet to enter and execute SQL, PL/SQL, and SQL\*Plus statements.
- Specify any actions that can be processed by the database connection associated with the worksheet.





# Executing SQL Statements

Use the Enter SQL Statement box to enter single or multiple SQL statements.

The screenshot illustrates the workflow for executing SQL statements in Oracle SQL Developer. It features three main components: the 'Enter SQL Statement' box, the 'Results' window, and the 'Script Output' window.

**Enter SQL Statement Box:** This box contains the SQL query: `SELECT employee_id, last_name FROM employees;`. The 'Execute' button (a green play icon) is highlighted with a red box and labeled **F9**. The 'Script' button (a document icon) is also highlighted with a red box and labeled **F5**. A red arrow points from the 'Execute' button to the 'Results' window.

**Results Window:** This window displays the query results in a table format. The 'Results' button (a green play icon) is highlighted with a red box and labeled **F9**. The table shows the following data:

	EMPLOYEE_ID	LAST_NAME
1	100	King
2	101	Kochhar
3	102	De Haan
4	103	Hunold
5	104	Ernst

**Script Output Window:** This window displays the query results in a text format. The 'Script Output' button (a document icon) is highlighted with a red box and labeled **F5**. The text shows the following data:

EMPLOYEE_ID	LAST_NAME
100	King
101	Kochhar
102	De Haan
103	Hunold
104	Ernst
105	Austin

# Saving SQL Scripts

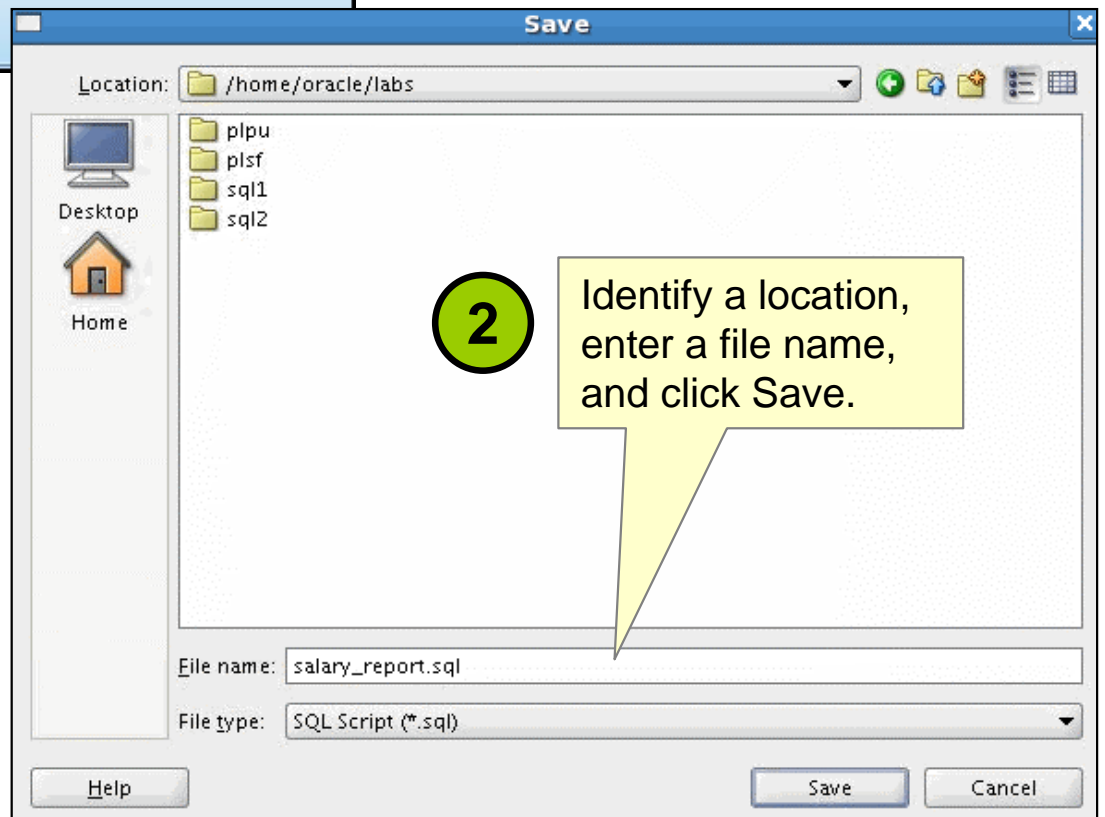
1

Click the Save icon to save your SQL statement to a file.




2

Identify a location, enter a file name, and click Save.



3

The contents of the saved file are visible and editable in your SQL Worksheet window.



# Executing Saved Script Files: Method 1

The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane displays a tree structure of folders and files. The 'confidence.sql' file is highlighted with a red rectangle and a yellow callout bubble labeled '1'. A yellow callout bubble labeled '2' points to the 'Run' button (a green play icon) in the toolbar. A yellow callout bubble labeled '3' points to the 'Run Script' button (a green play icon with a document) in the toolbar. A yellow callout bubble labeled '4' points to the 'myconnection' dropdown menu in the top right corner of the SQL Worksheet. The SQL Worksheet displays the following SQL code:

```
SELECT count(*) FROM tab;
SELECT count(*) FROM employees;
SELECT count(*) FROM countries;
SELECT count(*) FROM regions;
SELECT count(*) FROM locations;
SELECT count(*) FROM departments;
SELECT count(*) FROM jobs;
SELECT count(*) FROM job_history;
```

1. Use the Files tab to locate the script file that you want to open.

2. Double-click the script to display the code in the SQL Worksheet.

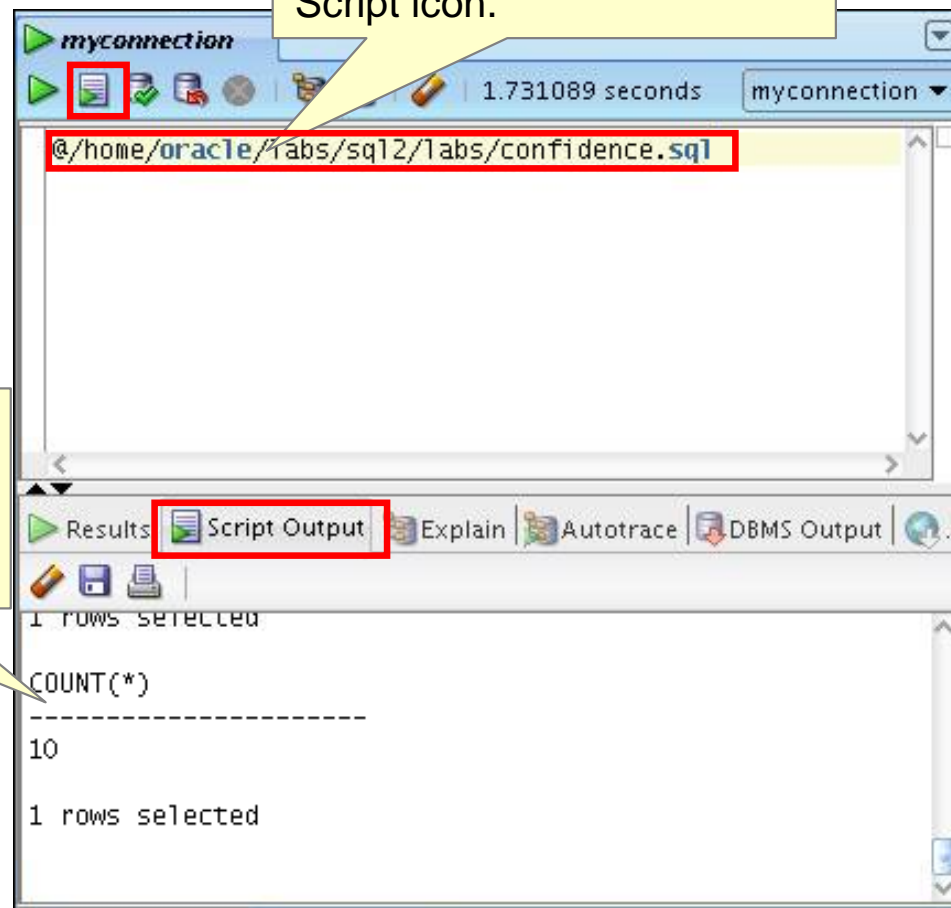
To run the code, click either:

- Execute Script (F9), or
- Run Script (F5)

Select a connection from the drop-down list.

# Executing Saved Script Files: Method 2


Use the @ command followed by the location and name of the file you want to execute, and click the Run Script icon.



The output from the script is displayed on the Script Output tabbed page.

# Formatting the SQL Code

**Before  
formatting**

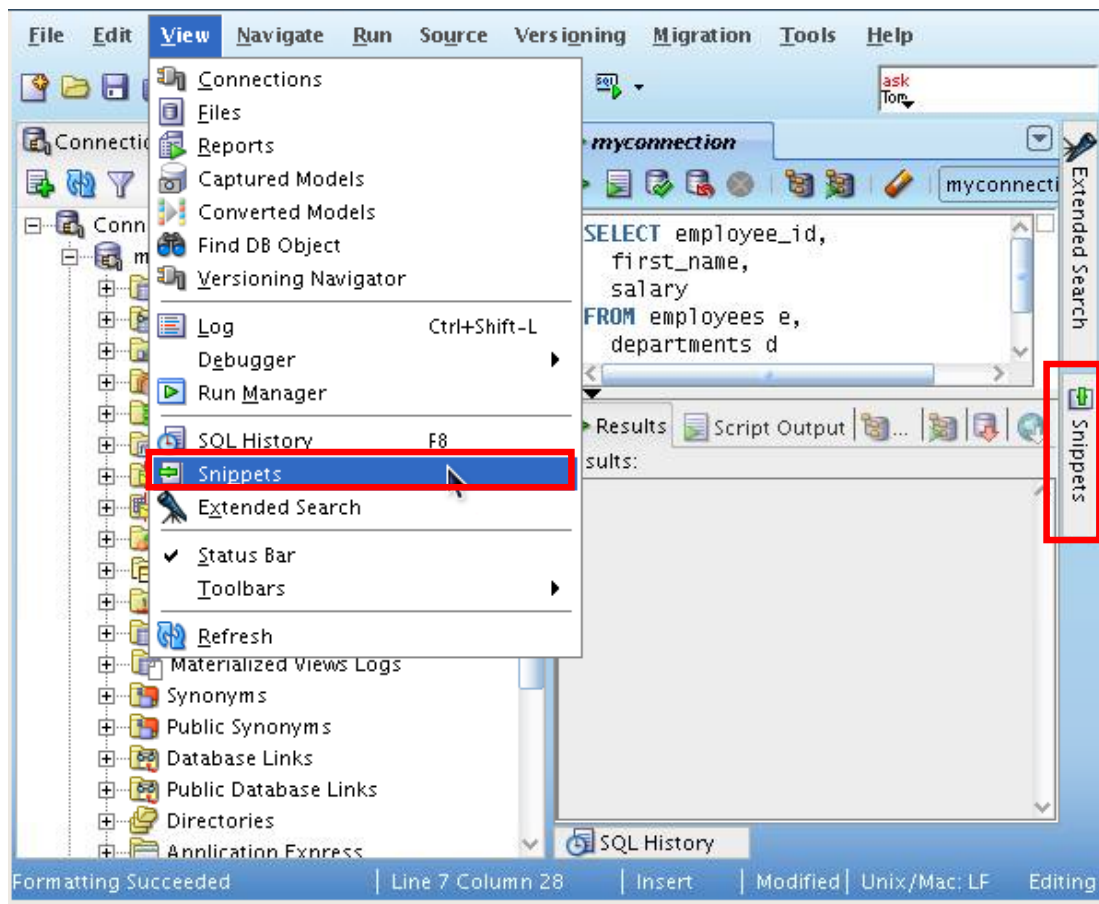


**After  
formatting**


A screenshot showing the SQL code after formatting. The code is now properly indented and uses uppercase keywords: `SELECT employee_id,  
first_name,  
salary  
FROM employees e,  
departments d  
WHERE e.department_id= d.department_id  
AND e.salary > 3000;`. A red arrow points from the 'After formatting' label to the formatted code.

# Using Snippets

Snippets are code fragments that may be just syntax or examples.




When you place your cursor here, it shows the Snippets window. From the drop-down list, you can select the functions category that you want.

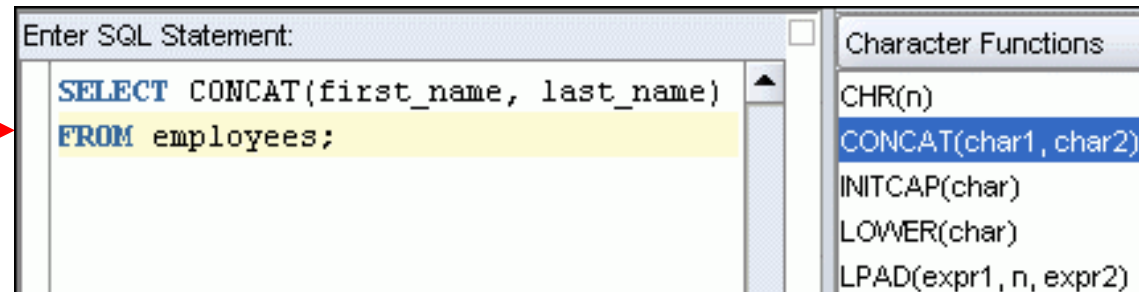


# Using Snippets: Example

**Inserting a snippet**




**Editing the snippet**





# Debugging Procedures and Functions

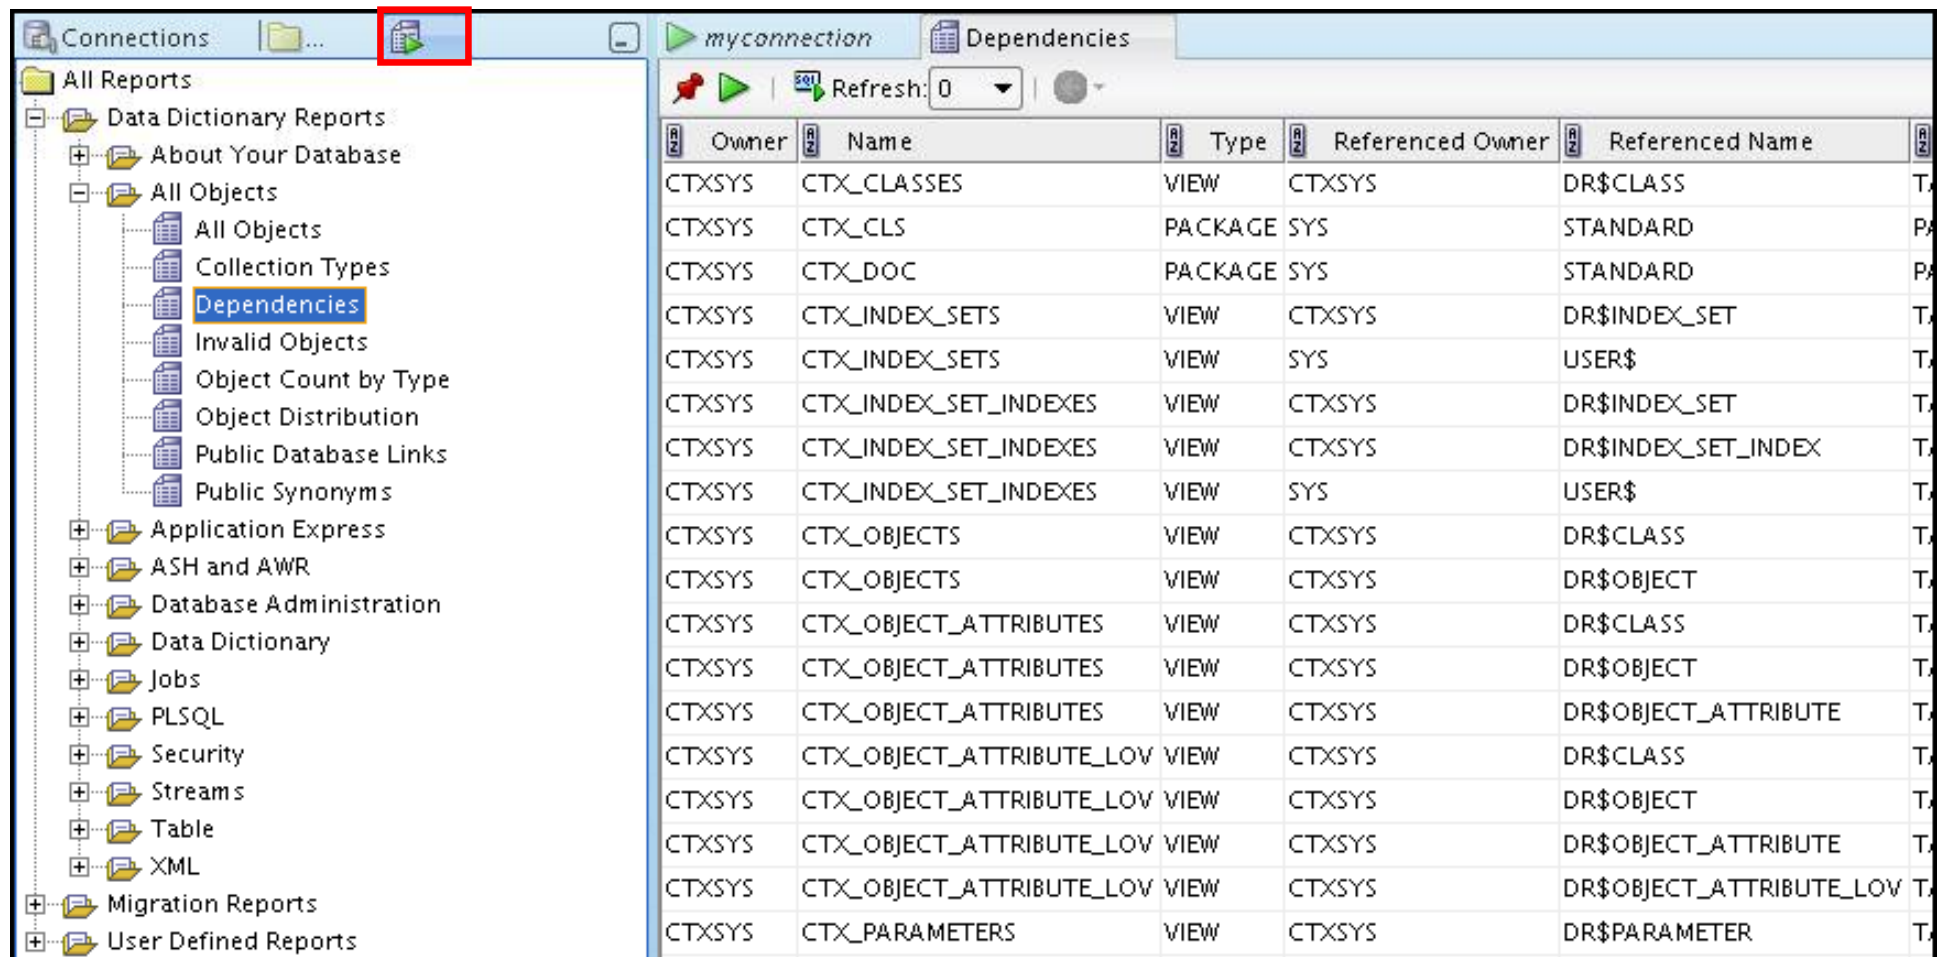
- Use SQL Developer to debug PL/SQL functions and procedures.
- Use the “Compile for Debug” option to perform a PL/SQL compilation so that the procedure can be debugged.
- Use Debug menu options to set breakpoints, and to perform step into and step over tasks.





# Database Reporting

SQL Developer provides a number of predefined reports about the database and its objects.

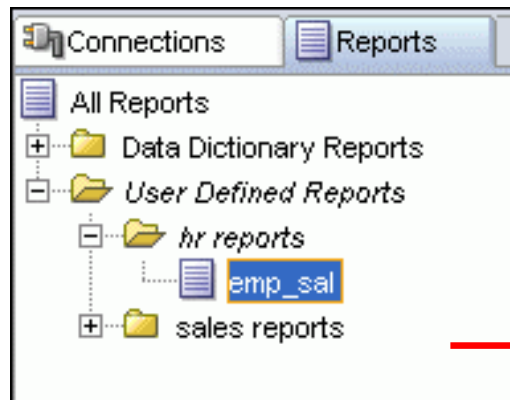
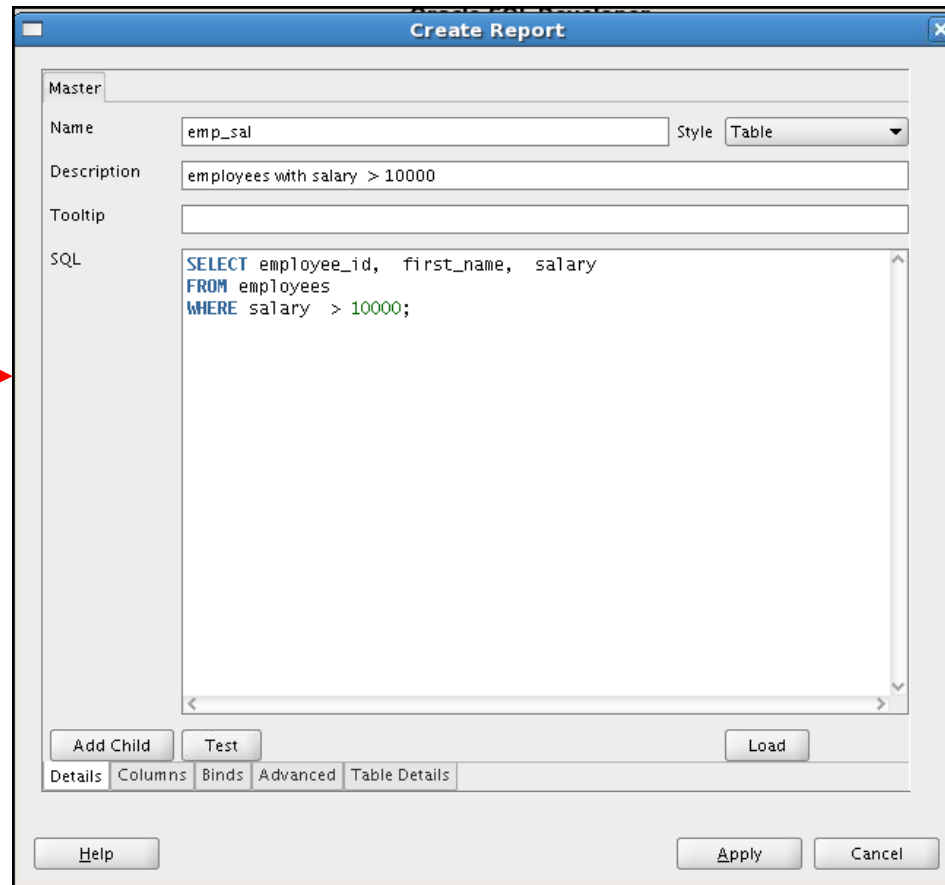



The screenshot shows the SQL Developer interface. On the left, the 'Connections' pane is open, and the 'Dependencies' report is selected under 'Data Dictionary Reports'. The main pane displays a table of database dependencies for the 'myconnection' connection. The table has columns for Owner, Name, Type, Referenced Owner, and Referenced Name. The 'Dependencies' report icon in the left pane is highlighted with a red box.

Owner	Name	Type	Referenced Owner	Referenced Name
CTXSYS	CTX_CLASSES	VIEW	CTXSYS	DR\$CLASS
CTXSYS	CTX_CLS	PACKAGE	SYS	STANDARD
CTXSYS	CTX_DOC	PACKAGE	SYS	STANDARD
CTXSYS	CTX_INDEX_SETS	VIEW	CTXSYS	DR\$INDEX_SET
CTXSYS	CTX_INDEX_SETS	VIEW	SYS	USER\$
CTXSYS	CTX_INDEX_SET_INDEXES	VIEW	CTXSYS	DR\$INDEX_SET
CTXSYS	CTX_INDEX_SET_INDEXES	VIEW	CTXSYS	DR\$INDEX_SET_INDEX
CTXSYS	CTX_INDEX_SET_INDEXES	VIEW	SYS	USER\$
CTXSYS	CTX_OBJECTS	VIEW	CTXSYS	DR\$CLASS
CTXSYS	CTX_OBJECTS	VIEW	CTXSYS	DR\$OBJECT
CTXSYS	CTX_OBJECT_ATTRIBUTES	VIEW	CTXSYS	DR\$CLASS
CTXSYS	CTX_OBJECT_ATTRIBUTES	VIEW	CTXSYS	DR\$OBJECT
CTXSYS	CTX_OBJECT_ATTRIBUTES	VIEW	CTXSYS	DR\$OBJECT_ATTRIBUTE
CTXSYS	CTX_OBJECT_ATTRIBUTE_LOV	VIEW	CTXSYS	DR\$CLASS
CTXSYS	CTX_OBJECT_ATTRIBUTE_LOV	VIEW	CTXSYS	DR\$OBJECT
CTXSYS	CTX_OBJECT_ATTRIBUTE_LOV	VIEW	CTXSYS	DR\$OBJECT_ATTRIBUTE
CTXSYS	CTX_OBJECT_ATTRIBUTE_LOV	VIEW	CTXSYS	DR\$OBJECT_ATTRIBUTE_LOV
CTXSYS	CTX_PARAMETERS	VIEW	CTXSYS	DR\$PARAMETER

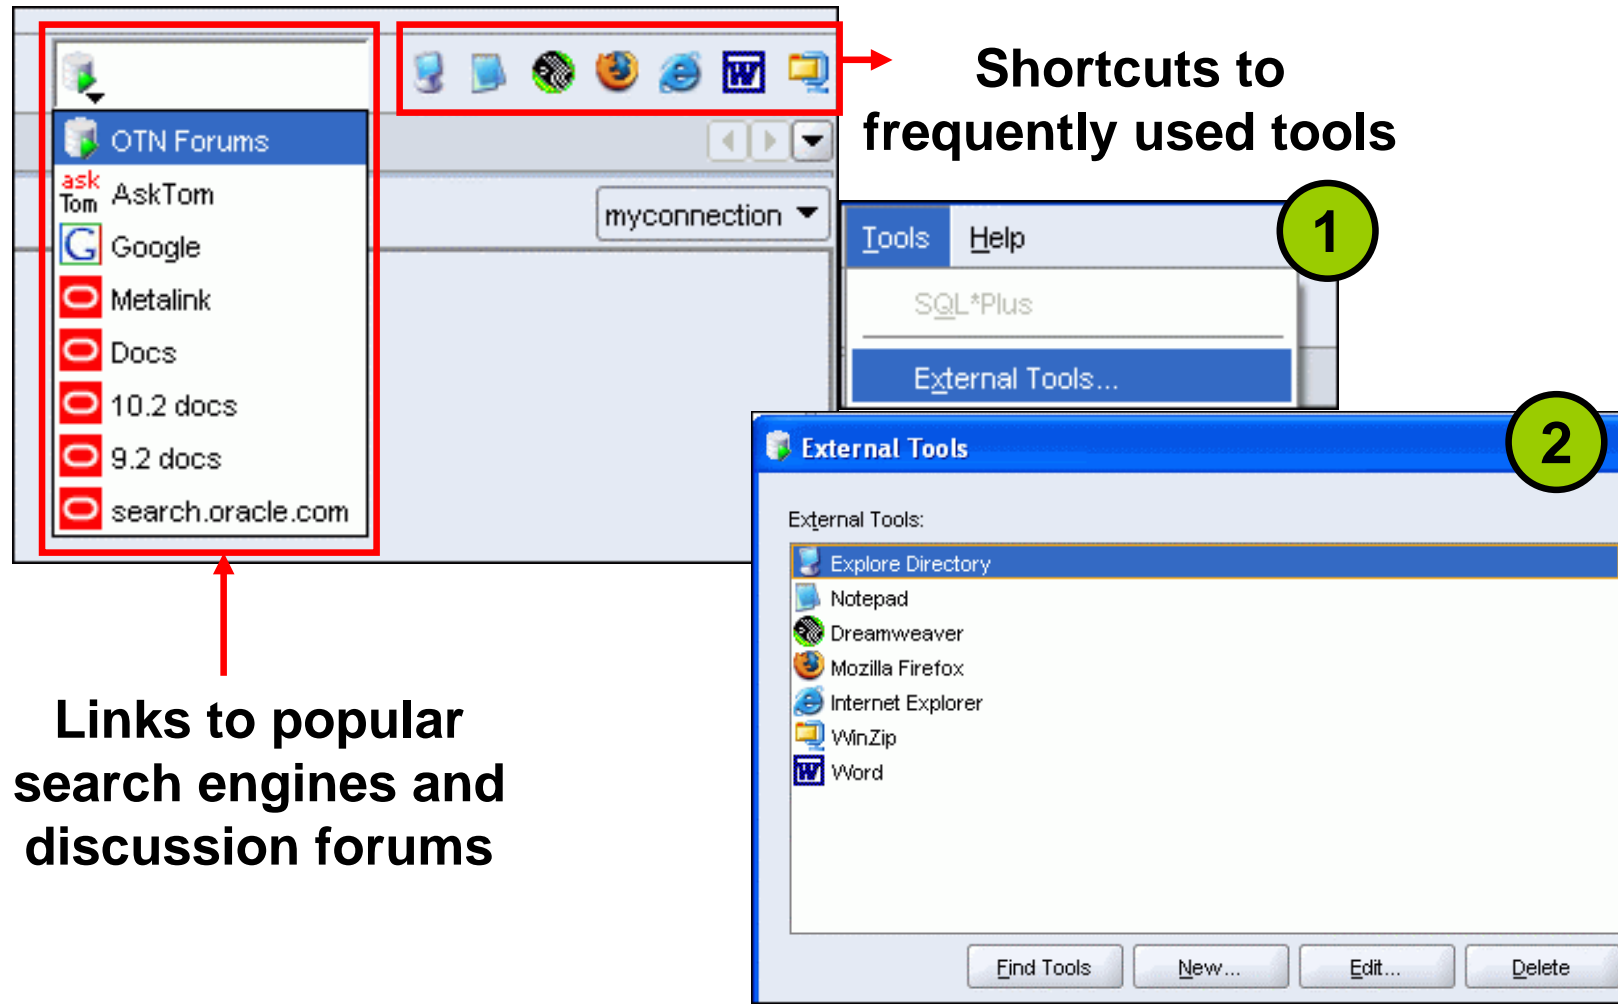
# Creating a User-Defined Report

Create and save user-defined reports for repeated use.




**Organize reports in folders.**

# Search Engines and External Tools

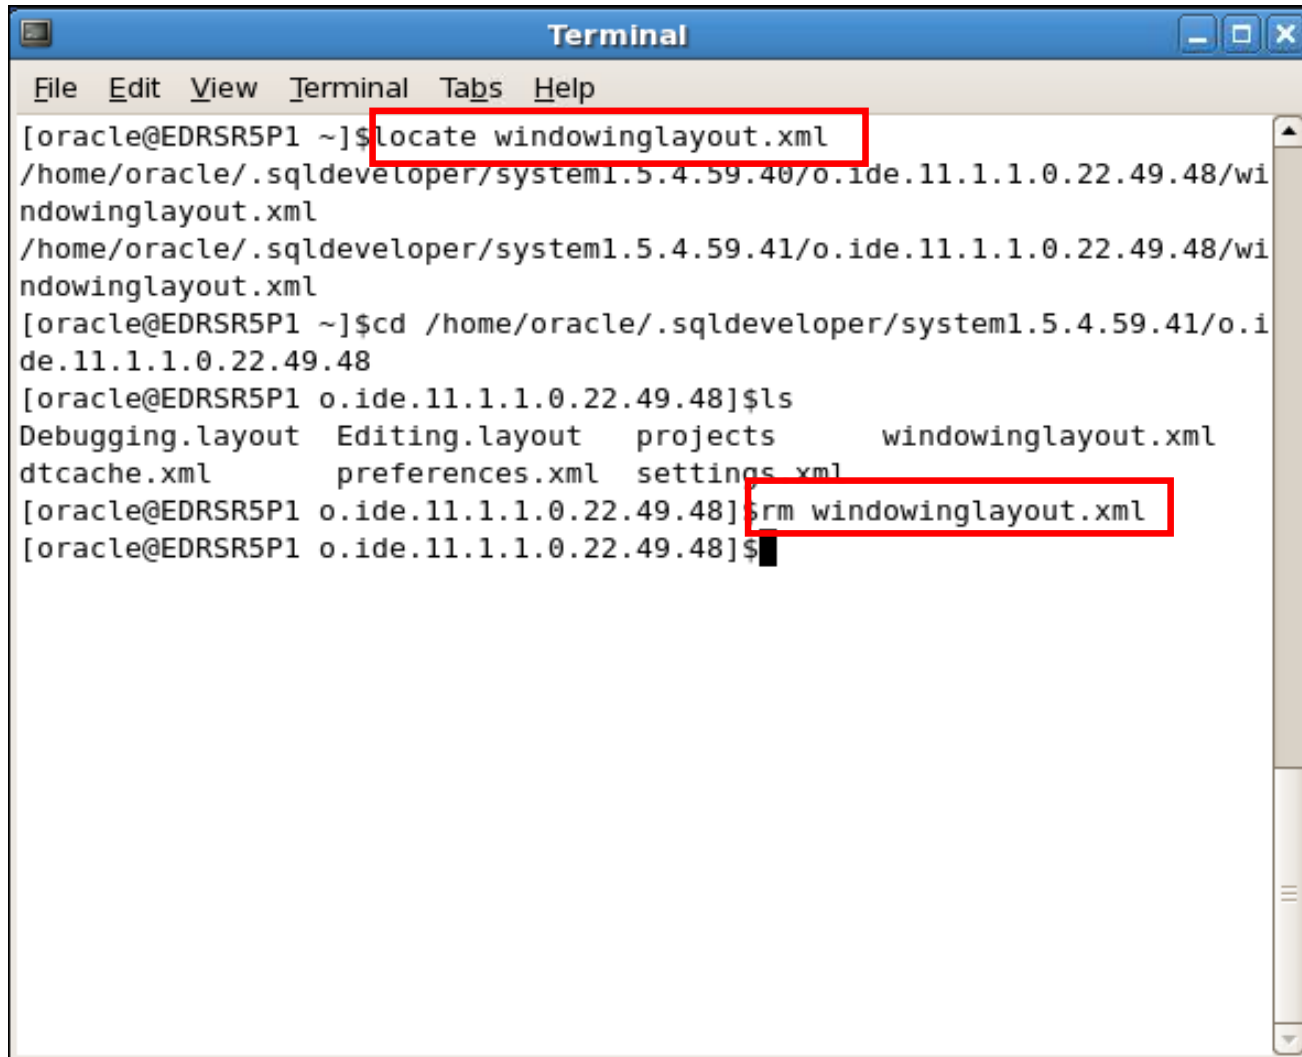


# Setting Preferences

- Customize the SQL Developer interface and environment.
- In the Tools menu, select Preferences.



# Resetting the SQL Developer Layout

A screenshot of a terminal window titled "Terminal" with a menu bar (File, Edit, View, Terminal, Tabs, Help). The terminal shows a series of commands and their outputs. Two red rectangles highlight the commands `locate windowinglayout.xml` and `rm windowinglayout.xml`.

```
[oracle@EDRSR5P1 ~]$ locate windowinglayout.xml
/home/oracle/.sqldeveloper/system1.5.4.59.40/o.ide.11.1.1.0.22.49.48/wi
ndowinglayout.xml
/home/oracle/.sqldeveloper/system1.5.4.59.41/o.ide.11.1.1.0.22.49.48/wi
ndowinglayout.xml
[oracle@EDRSR5P1 ~]$ cd /home/oracle/.sqldeveloper/system1.5.4.59.41/o.i
de.11.1.1.0.22.49.48
[oracle@EDRSR5P1 o.ide.11.1.1.0.22.49.48]$ ls
Debugging.layout  Editing.layout  projects        windowinglayout.xml
dtcache.xml       preferences.xml  settings.xml
[oracle@EDRSR5P1 o.ide.11.1.1.0.22.49.48]$ rm windowinglayout.xml
[oracle@EDRSR5P1 o.ide.11.1.1.0.22.49.48]$
```

# Summary

In this appendix, you should have learned how to use SQL Developer to do the following:

- Browse, create, and edit database objects
- Execute SQL statements and scripts in SQL Worksheet
- Create and save custom reports

# D

## Using SQL\*Plus

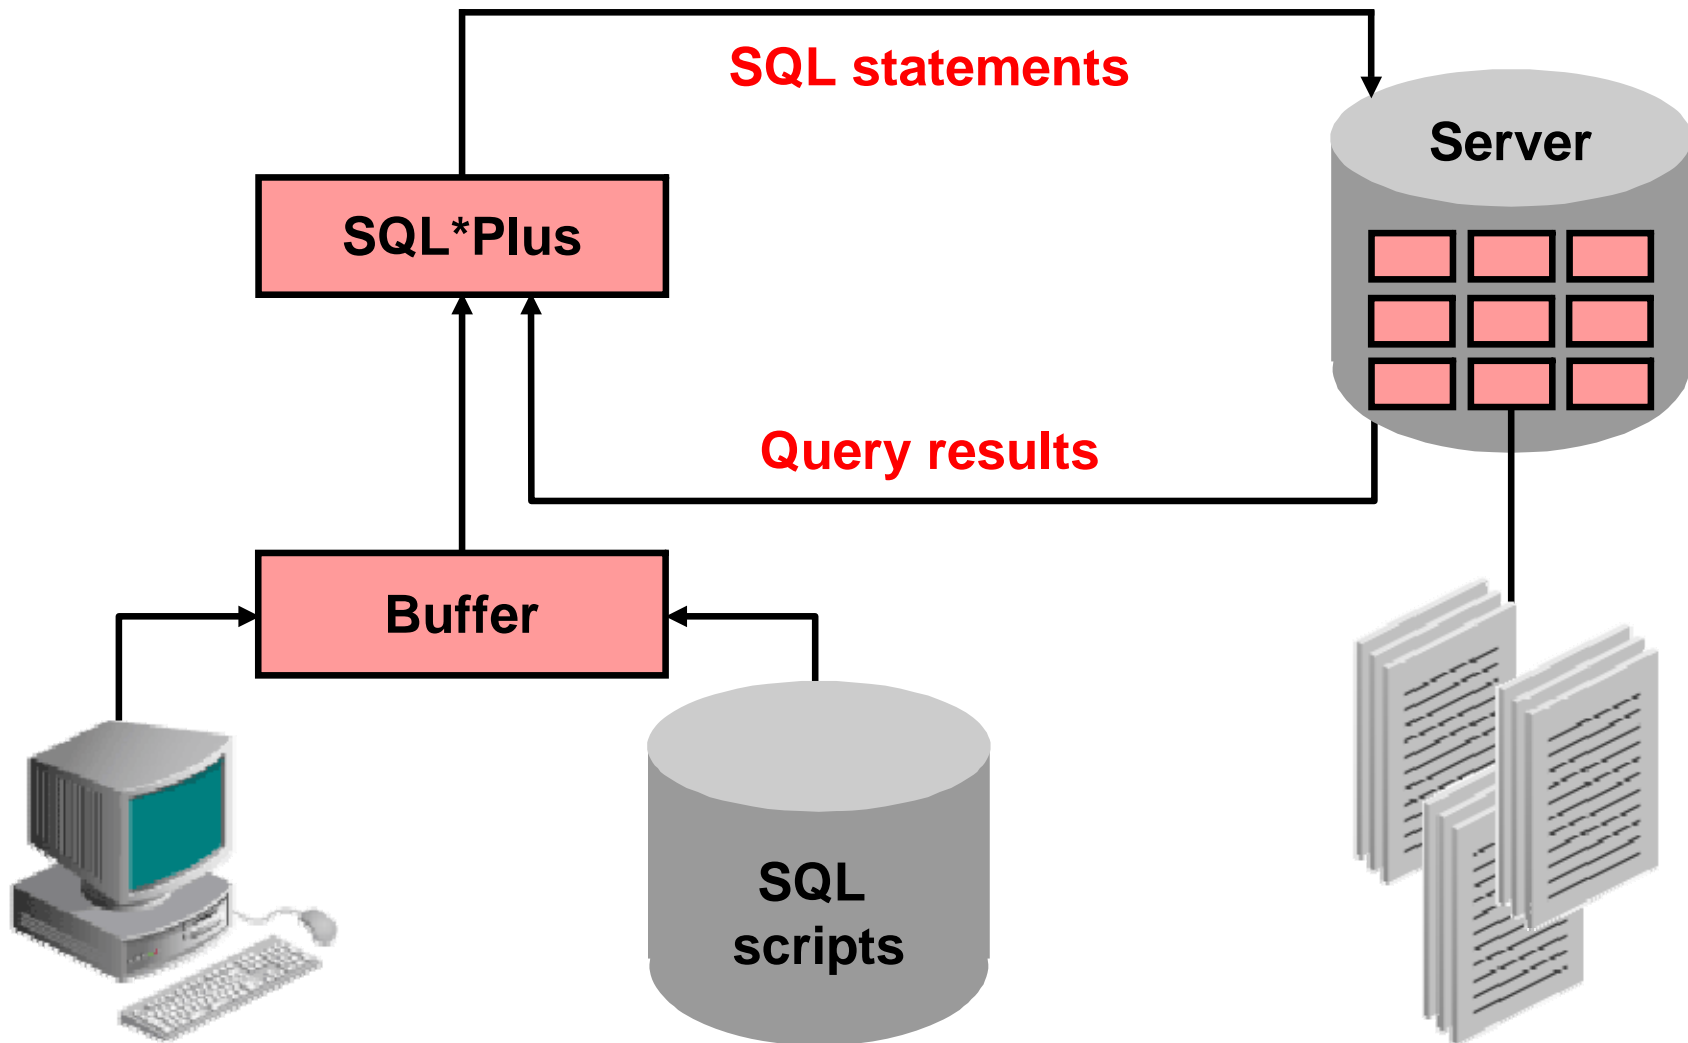
# Objectives

After completing this appendix, you should be able to do the following:

- Log in to SQL\*Plus
- Edit SQL commands
- Format the output using SQL\*Plus commands
- Interact with script files




# SQL and SQL\*Plus Interaction



# SQL Statements Versus SQL\*Plus Commands


## SQL

- A language
- ANSI-standard
- Keywords cannot be abbreviated.
- Statements manipulate data and table definitions in the database.



## SQL\*Plus

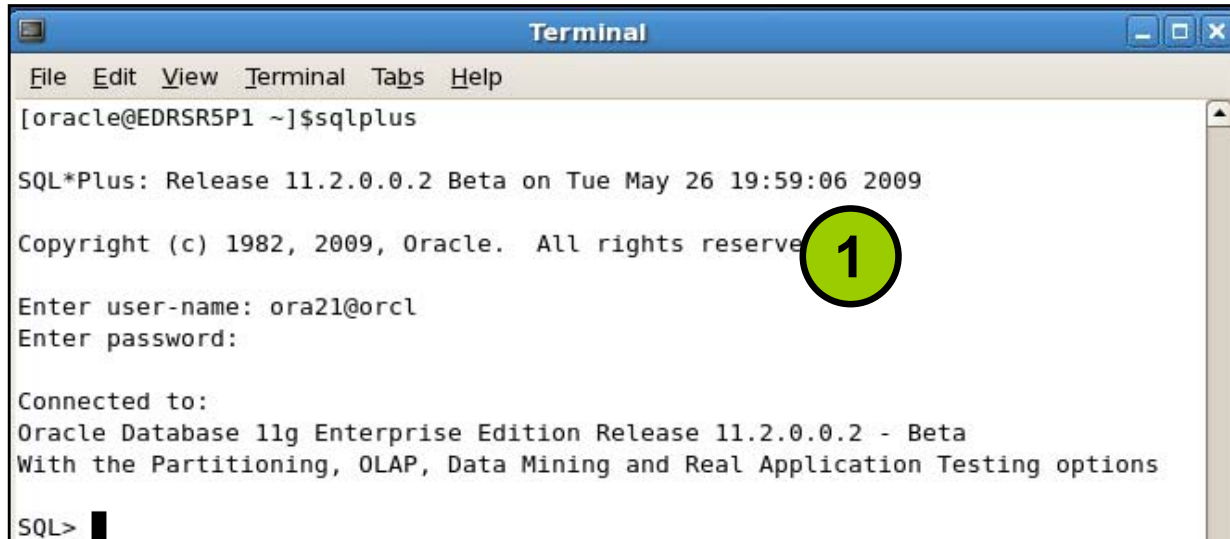
- An environment
- Oracle-proprietary
- Keywords can be abbreviated.
- Commands do not allow manipulation of values in the database.



# Overview of SQL\*Plus

- Log in to SQL\*Plus.
- Describe the table structure.
- Edit your SQL statement.
- Execute SQL from SQL\*Plus.
- Save SQL statements to files and append SQL statements to files.
- Execute saved files.
- Load commands from the file to buffer to edit.

# Logging In to SQL\*Plus

A terminal window titled "Terminal" with a menu bar (File, Edit, View, Terminal, Tabs, Help). The command prompt shows the user 'oracle' at host 'EDRSR5P1' running 'sqlplus'. The output displays the SQL\*Plus version (11.2.0.0.2 Beta), release date (Tue May 26 19:59:06 2009), and copyright notice. It then prompts for a username and password. The user enters 'ora21@orcl' and the password is masked. The terminal then shows the connection details for the Oracle Database 11g Enterprise Edition.

```
Terminal
File Edit View Terminal Tabs Help
[oracle@EDRSR5P1 ~]$sqlplus

SQL*Plus: Release 11.2.0.0.2 Beta on Tue May 26 19:59:06 2009

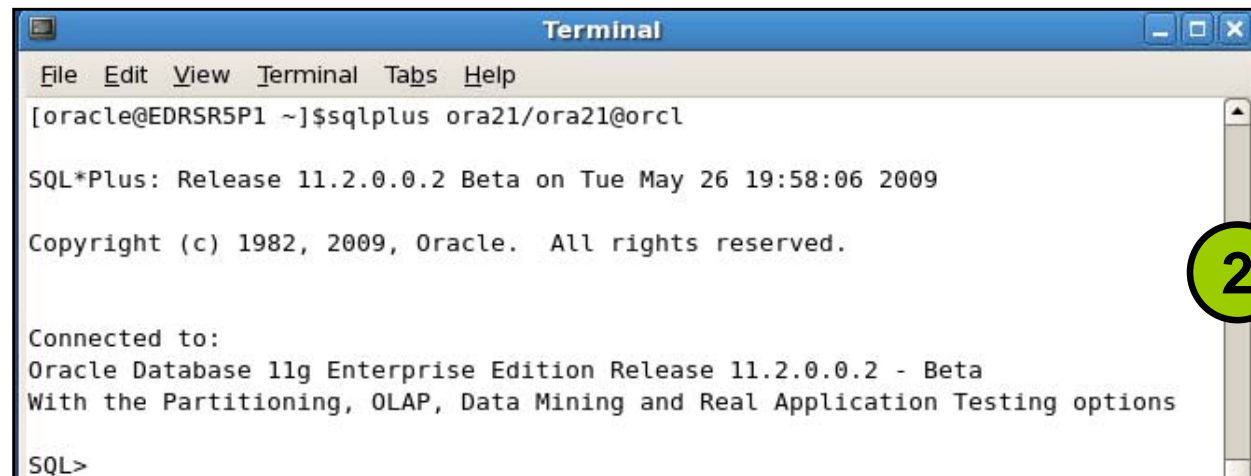
Copyright (c) 1982, 2009, Oracle. All rights reserved.

Enter user-name: ora21@orcl
Enter password:

Connected to:
Oracle Database 11g Enterprise Edition Release 11.2.0.0.2 - Beta
With the Partitioning, OLAP, Data Mining and Real Application Testing options

SQL>
```

```
sqlplus [username[/password[@database]]]
```

A terminal window titled "Terminal" with a menu bar (File, Edit, View, Terminal, Tabs, Help). The command prompt shows the user 'oracle' at host 'EDRSR5P1' running 'sqlplus ora21/ora21@orcl'. The output displays the SQL\*Plus version (11.2.0.0.2 Beta), release date (Tue May 26 19:58:06 2009), and copyright notice. It then shows the connection details for the Oracle Database 11g Enterprise Edition.

```
Terminal
File Edit View Terminal Tabs Help
[oracle@EDRSR5P1 ~]$sqlplus ora21/ora21@orcl

SQL*Plus: Release 11.2.0.0.2 Beta on Tue May 26 19:58:06 2009

Copyright (c) 1982, 2009, Oracle. All rights reserved.

Connected to:
Oracle Database 11g Enterprise Edition Release 11.2.0.0.2 - Beta
With the Partitioning, OLAP, Data Mining and Real Application Testing options

SQL>
```

# Displaying the Table Structure

Use the SQL\*Plus DESCRIBE command to display the structure of a table:

```
DESC[RIBE] tablename
```

# Displaying the Table Structure

```
DESCRIBE departments
```

Name	Null?	Type
-----	-----	-----
DEPARTMENT_ID	NOT NULL	NUMBER(4)
DEPARTMENT_NAME	NOT NULL	VARCHAR2(30)
MANAGER_ID		NUMBER(6)
LOCATION_ID		NUMBER(4)

# SQL\*Plus Editing Commands

- A[PPEND] *text*
- C[HANGE] / *old* / *new*
- C[HANGE] / *text* /
- CL[EAR] BUFF[ER]
- DEL
- DEL *n*
- DEL *m n*

# SQL\*Plus Editing Commands

- I [NPUT]
- I [NPUT] *text*
- L [IST]
- L [IST] *n*
- L [IST] *m n*
- R [UN]
- *n*
- *n text*
- 0 *text*



# Using LIST, n, and APPEND

```
LIST
  1  SELECT last_name
  2* FROM    employees
```

```
1
  1* SELECT last_name
```

```
A , job_id
  1* SELECT last_name, job_id
```

```
LIST
  1  SELECT last_name, job_id
  2* FROM    employees
```

# Using the CHANGE Command

```
LIST  
1* SELECT * from employees
```

```
c/employees/departments  
1* SELECT * from departments
```

```
LIST  
1* SELECT * from departments
```

# SQL\*Plus File Commands

- `SAVE filename`
- `GET filename`
- `START filename`
- `@ filename`
- `EDIT filename`
- `SPOOL filename`
- `EXIT`

# Using the SAVE and START Commands

```
LIST
```

```
1  SELECT last_name, manager_id, department_id
2*  FROM employees
```

```
SAVE my_query
```

```
Created file my_query
```

```
START my_query
```

LAST_NAME	MANAGER_ID	DEPARTMENT_ID
King		90
Kochhar	100	90
...		

107 rows selected.

# SERVEROUTPUT Command

- Use the `SET SERVEROUT [PUT]` command to control whether to display the output of stored procedures or PL/SQL blocks in SQL\*Plus.
- The `DBMS_OUTPUT` line length limit is increased from 255 bytes to 32767 bytes.
- The default size is now unlimited.
- Resources are not preallocated when `SERVEROUTPUT` is set.
- Because there is no performance penalty, use `UNLIMITED` unless you want to conserve physical memory.

```
SET SERVEROUT [PUT] {ON | OFF} [SIZE {n | UNL[IMITED]}]  
[FOR [MAT] {WRA[PPED] | WOR[D_WAPPED] | TRU[NCATED]}]
```

# Using the SQL\*Plus SPOOL Command

```
SPO[OL] [file_name[.ext] [CRE[ATE] | REP[LACE] |  
APP[END]] | OFF | OUT]
```

Option	Description
file_name[.ext]	Spools output to the specified file name
CRE[ATE]	Creates a new file with the name specified
REP[LACE]	Replaces the contents of an existing file. If the file does not exist, REPLACE creates the file.
APP[END]	Adds the contents of the buffer to the end of the file you specify
OFF	Stops spooling
OUT	Stops spooling and sends the file to your computer's standard (default) printer

# Using the AUTOTRACE Command

- It displays a report after the successful execution of SQL data manipulation statements (DML) statements such as SELECT, INSERT, UPDATE, or DELETE.
- The report can now include execution statistics and the query execution path.

```
SET AUTOT[RACE] {ON | OFF | TRACE[ONLY]} [EXP[LAIN]]  
[STAT[ISTICS]]
```

```
SET AUTOTRACE ON  
-- The AUTOTRACE report includes both the optimizer  
-- execution path and the SQL statement execution  
-- statistics
```

# Summary

In this appendix, you should have learned how to use SQL\*Plus as an environment to do the following:

- Execute SQL statements
- Edit SQL statements
- Format the output
- Interact with script files



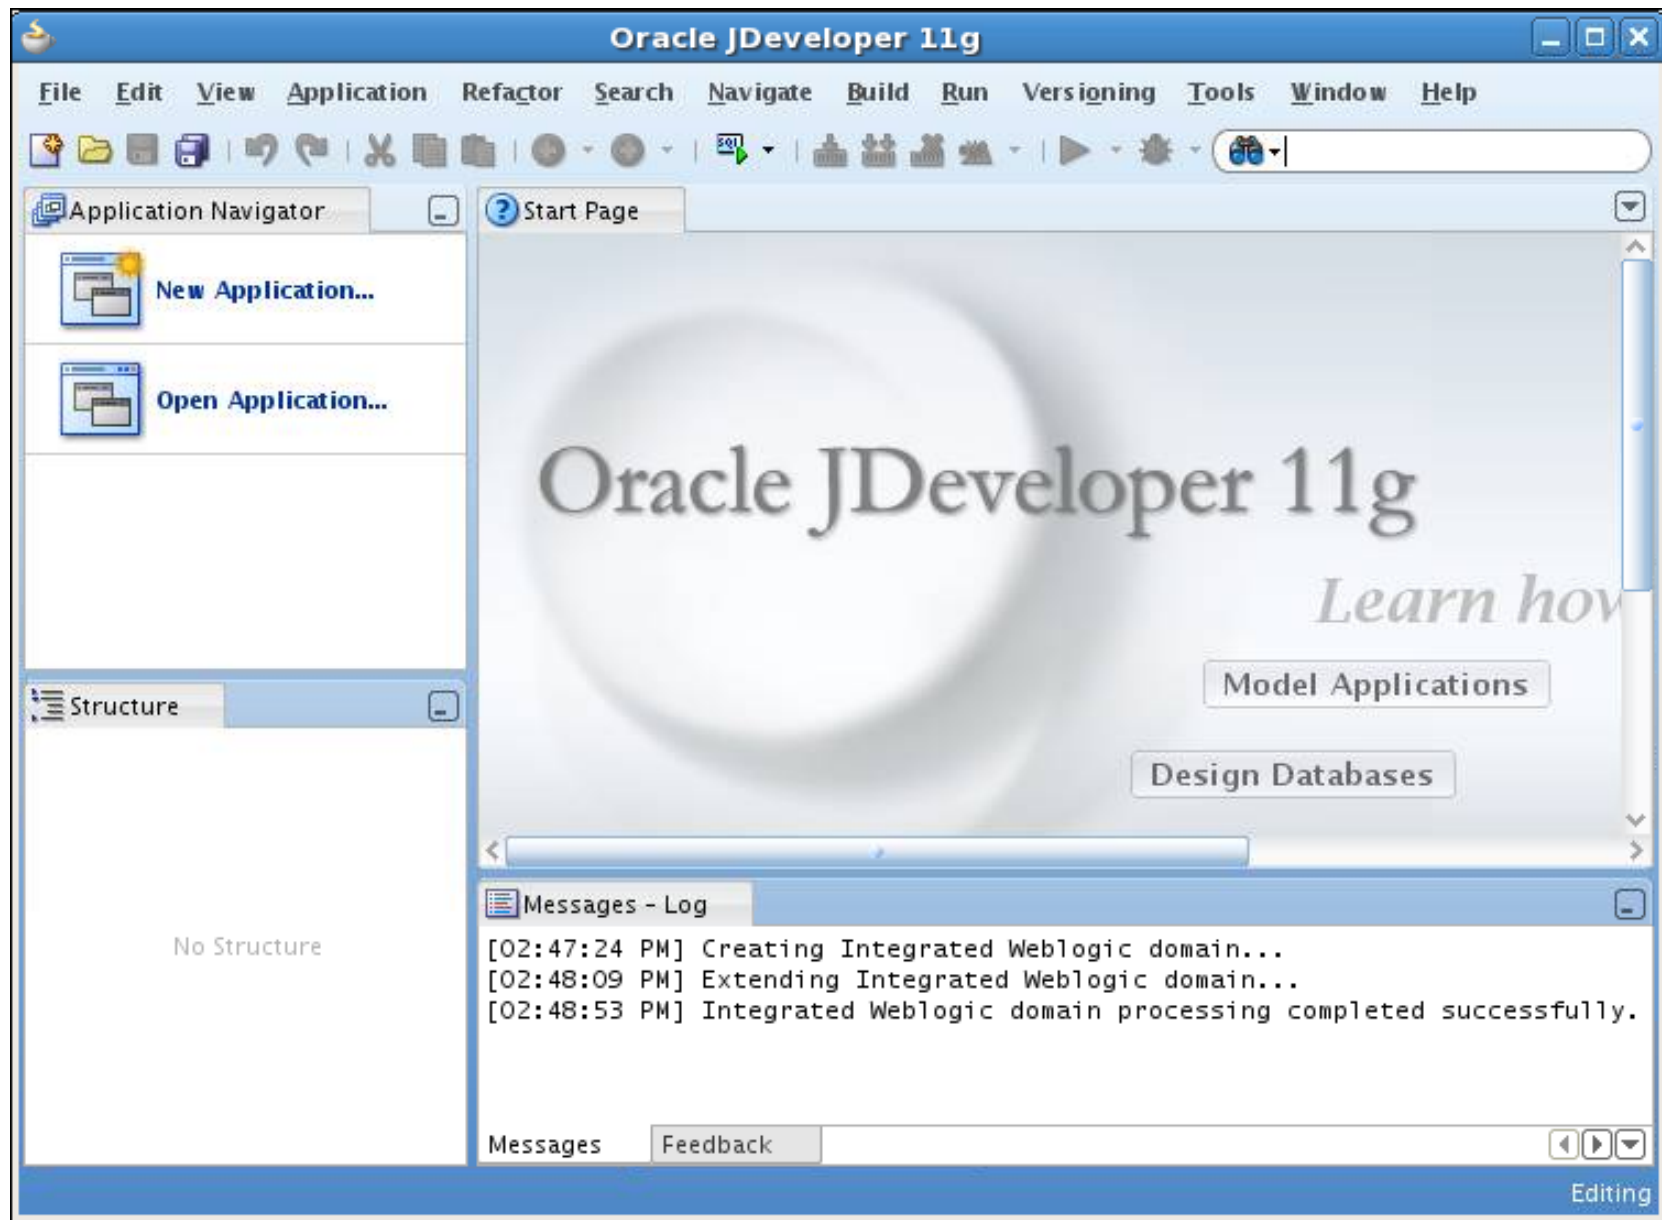
# Using JDeveloper

# Objectives

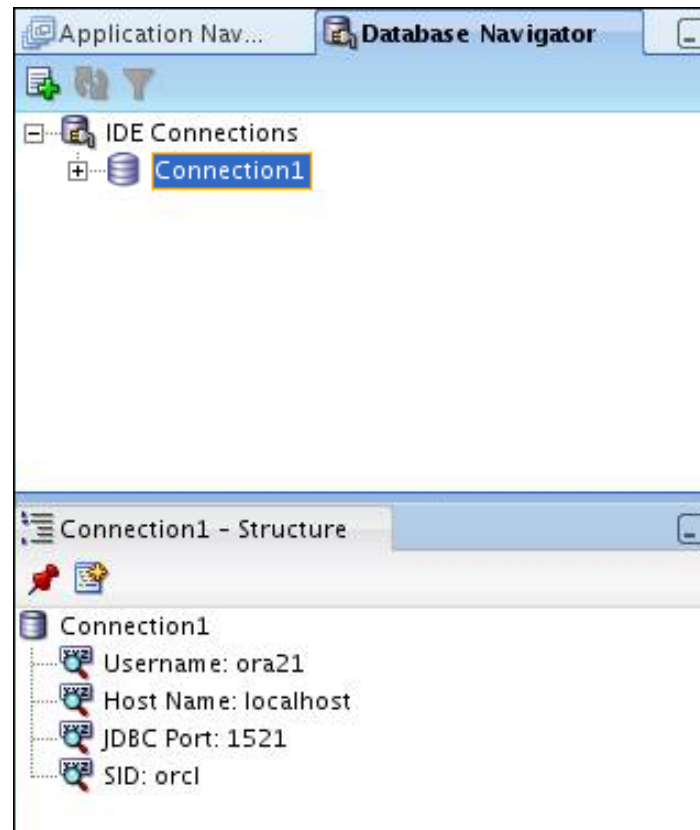
After completing this appendix, you should be able to do the following:

- List the key features of Oracle JDeveloper
- Create a database connection in JDeveloper
- Manage database objects in JDeveloper
- Use JDeveloper to execute SQL Commands
- Create and run PL/SQL Program Units


# Oracle JDeveloper



# Database Navigator



# Creating a Connection



1

Click the New Connection icon in the Database Navigator.

2

In the Create Database Connection window, enter the username, password and the SID.

3

Test the connection.

A screenshot of the 'Create Database Connection' dialog box. The 'Create Connection In:' dropdown is set to 'IDE Connections'. The 'Connection Name' is 'Connection1' and the 'Connection Type' is 'Oracle (JDBC)'. The 'Username' field contains 'ora21', the 'Password' field contains six dots, and the 'Role' dropdown is empty. The 'Save Password' checkbox is checked. Under the 'Oracle (JDBC) Settings' section, the 'Enter Custom JDBC URL' checkbox is unchecked. The 'Driver' dropdown is set to 'thin'. The 'Host Name' is 'localhost' and the 'JDBC Port' is '1521'. The 'SID' radio button is selected with the value 'orcl', and the 'Service Name' radio button is unselected with the value 'XE'. The 'Test Connection' button is highlighted with a red box. Below it, a text area shows 'Success!'. At the bottom, the 'OK' button is also highlighted with a red box, along with 'Help' and 'Cancel' buttons.

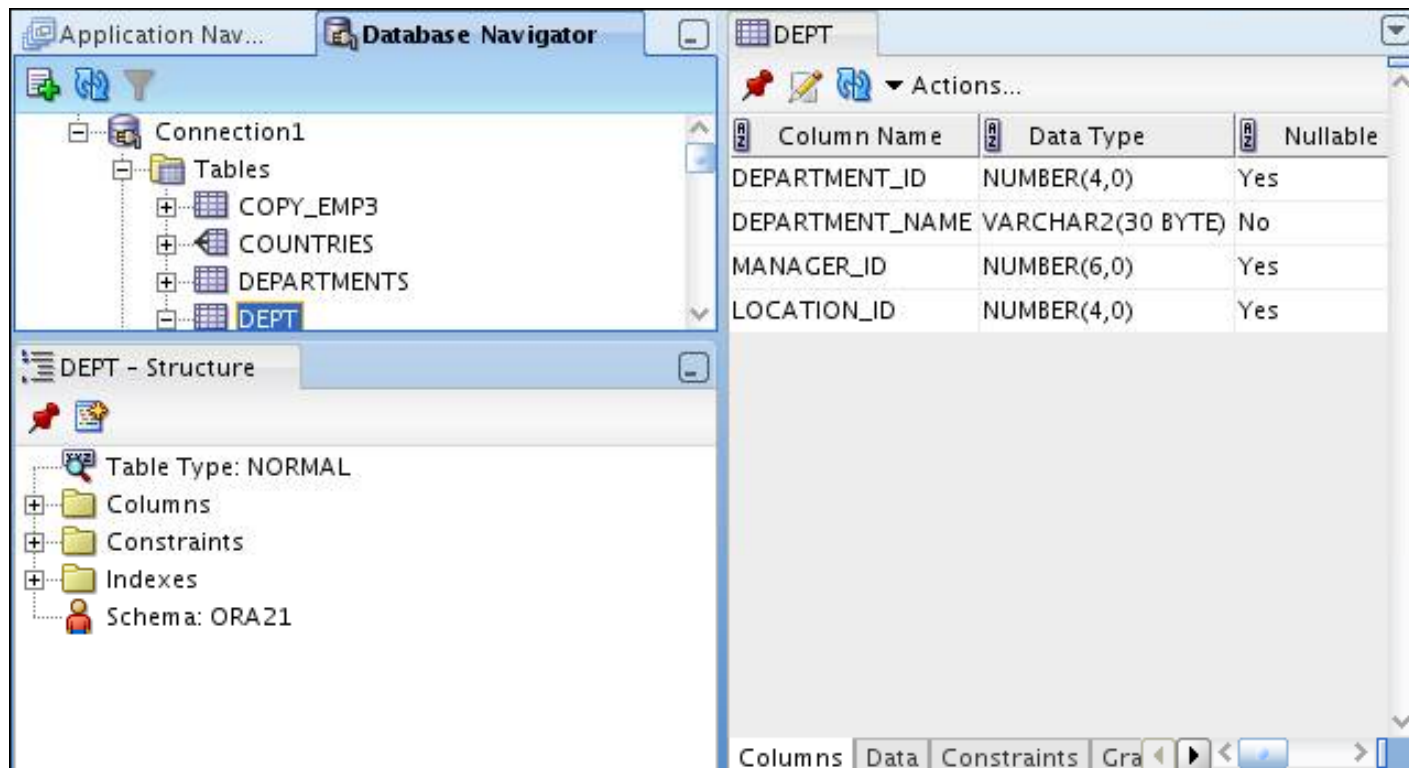
4

Click OK.

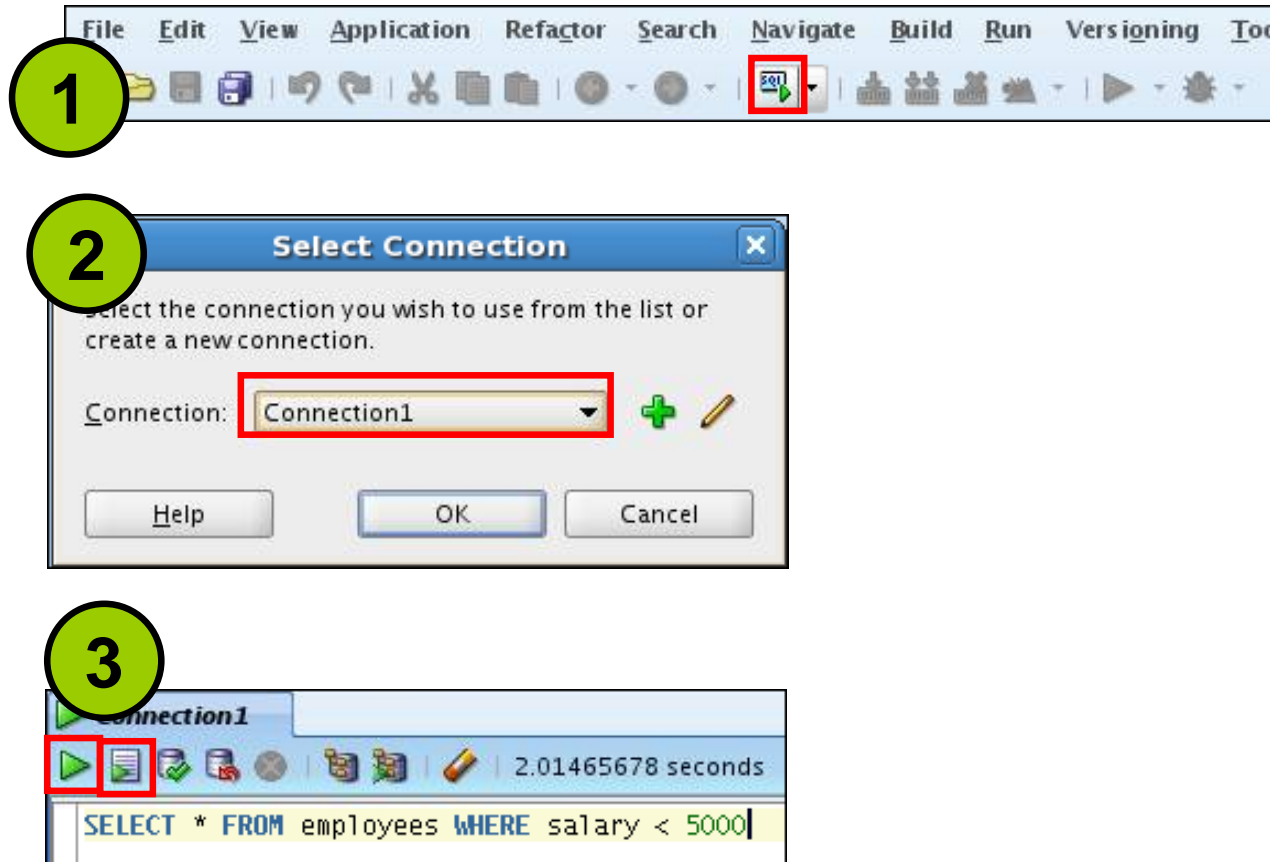
# Browsing Database Objects

Use the Database Navigator to:

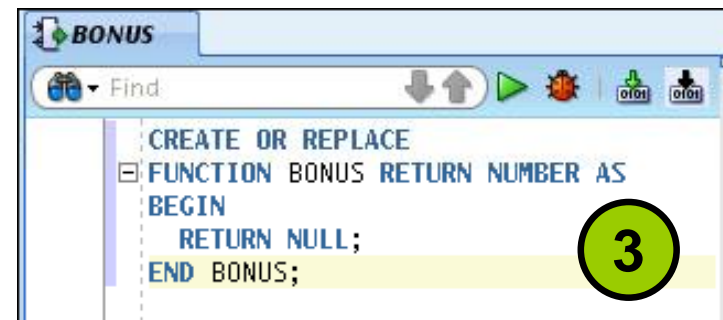
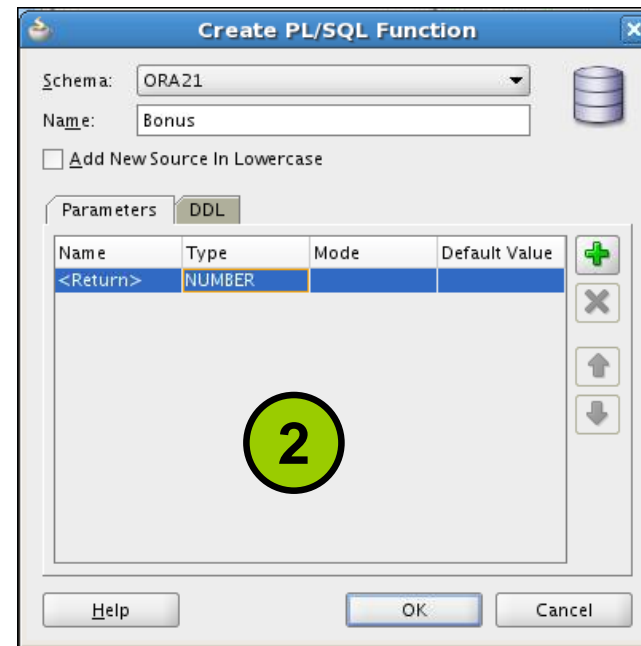
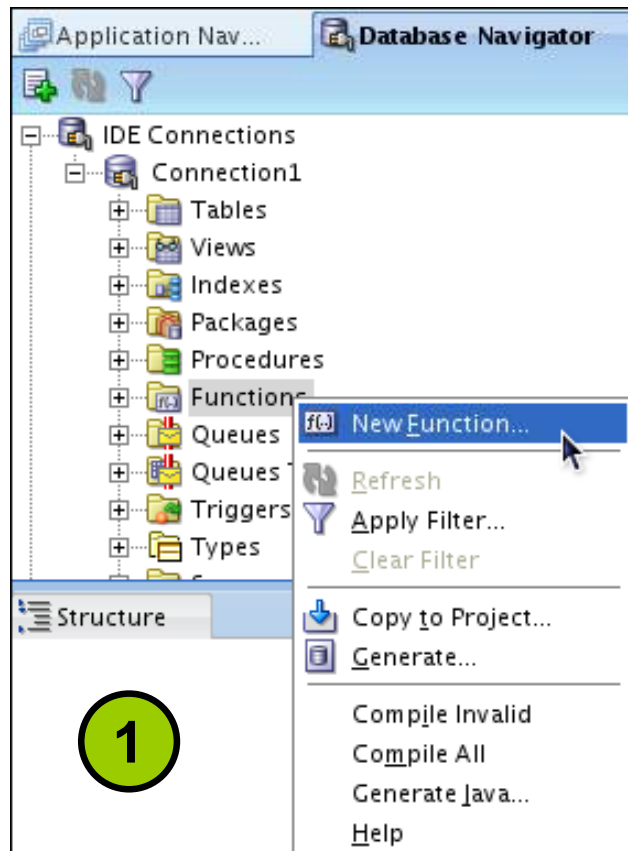
- Browse through many objects in a database schema
- Review the definitions of objects at a glance



# Executing SQL Statements



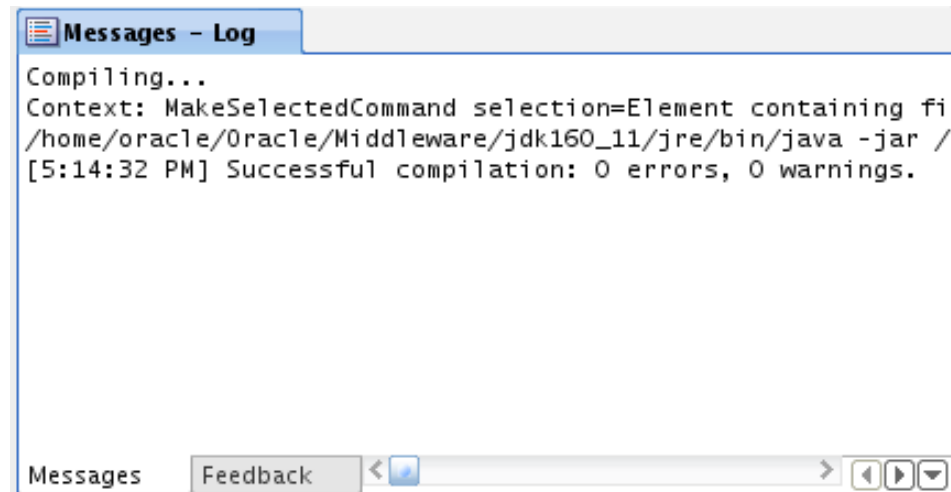
# Creating Program Units



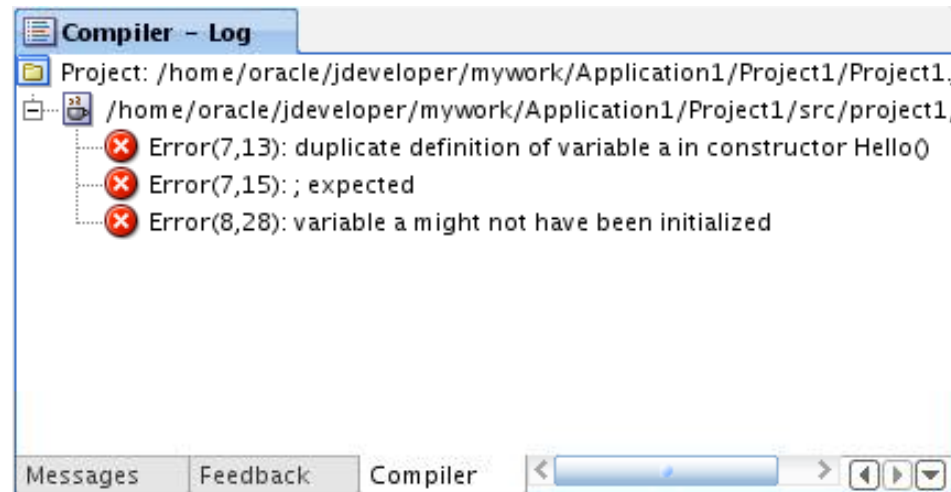
**Skeleton of the function**



# Compiling

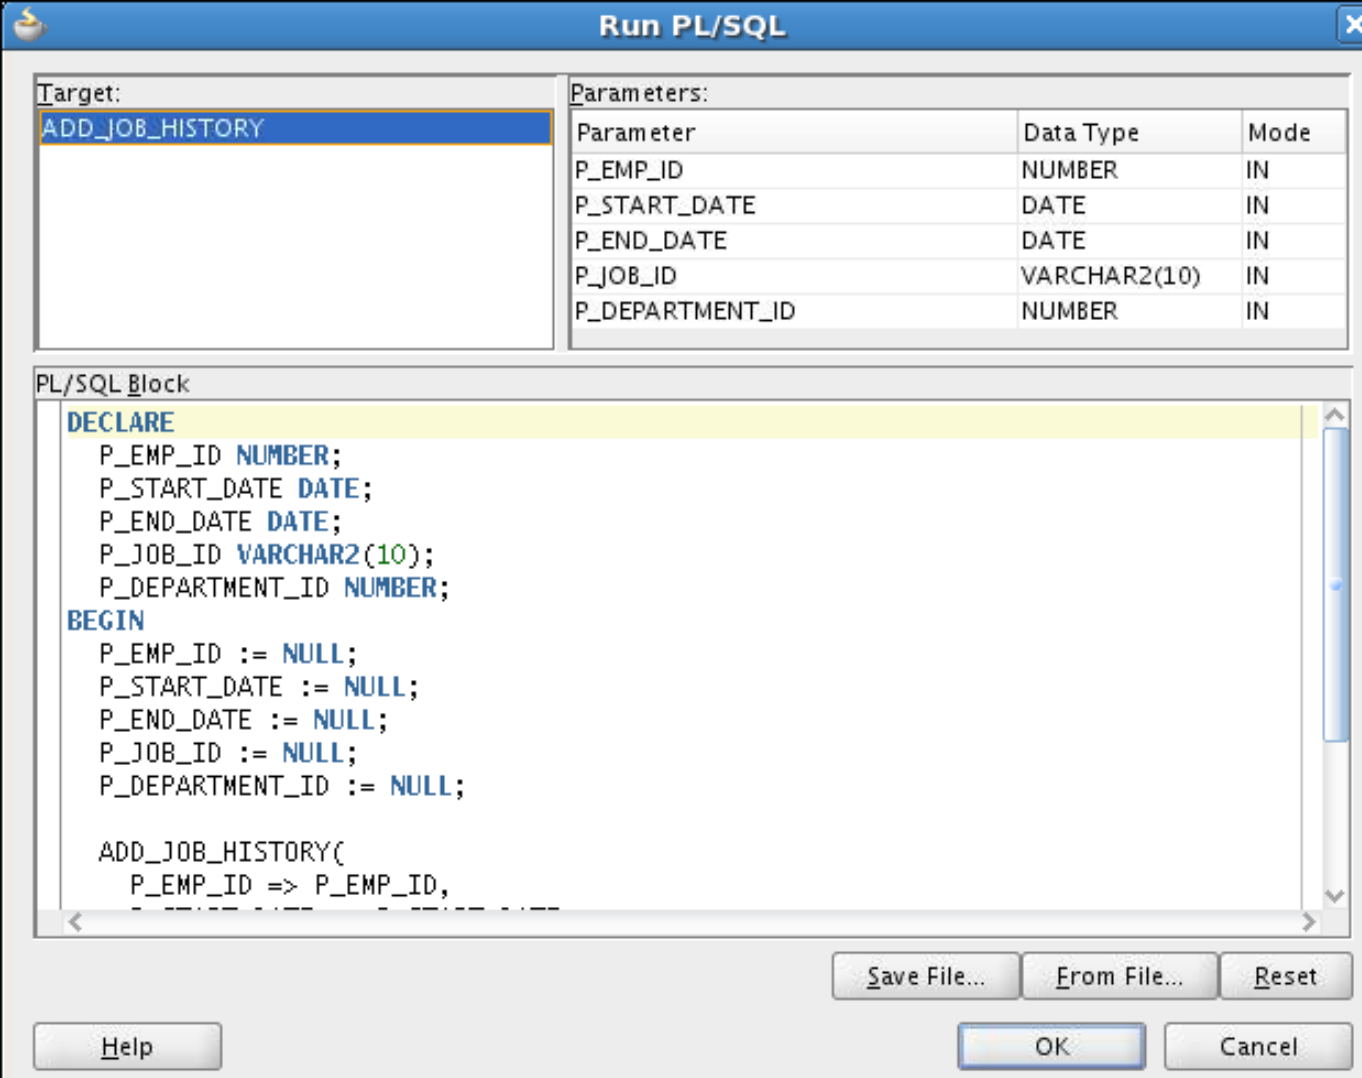


## Compilation with errors



## Compilation without errors

# Running a Program Unit



**Run PL/SQL**

Target:

ADD\_JOB\_HISTORY

Parameters:

Parameter	Data Type	Mode
P_EMP_ID	NUMBER	IN
P_START_DATE	DATE	IN
P_END_DATE	DATE	IN
P_JOB_ID	VARCHAR2(10)	IN
P_DEPARTMENT_ID	NUMBER	IN

PL/SQL Block

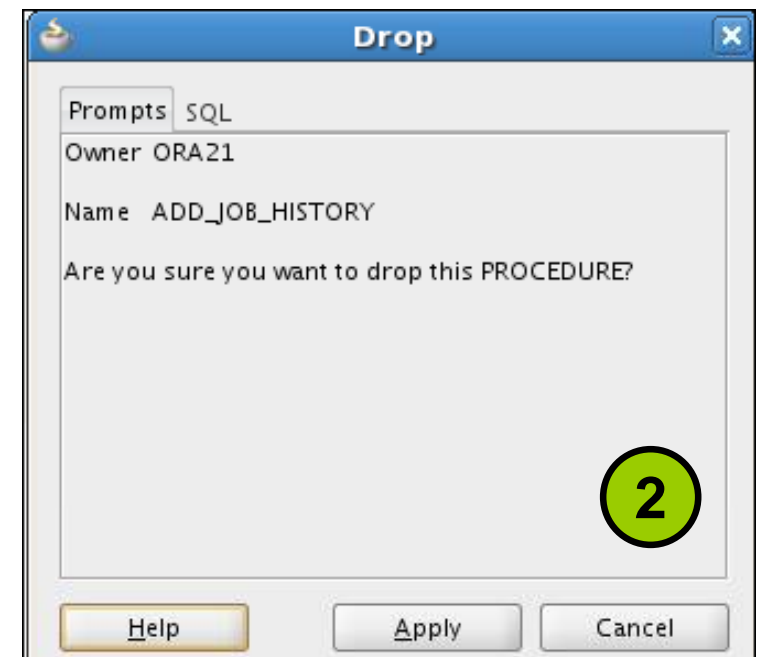
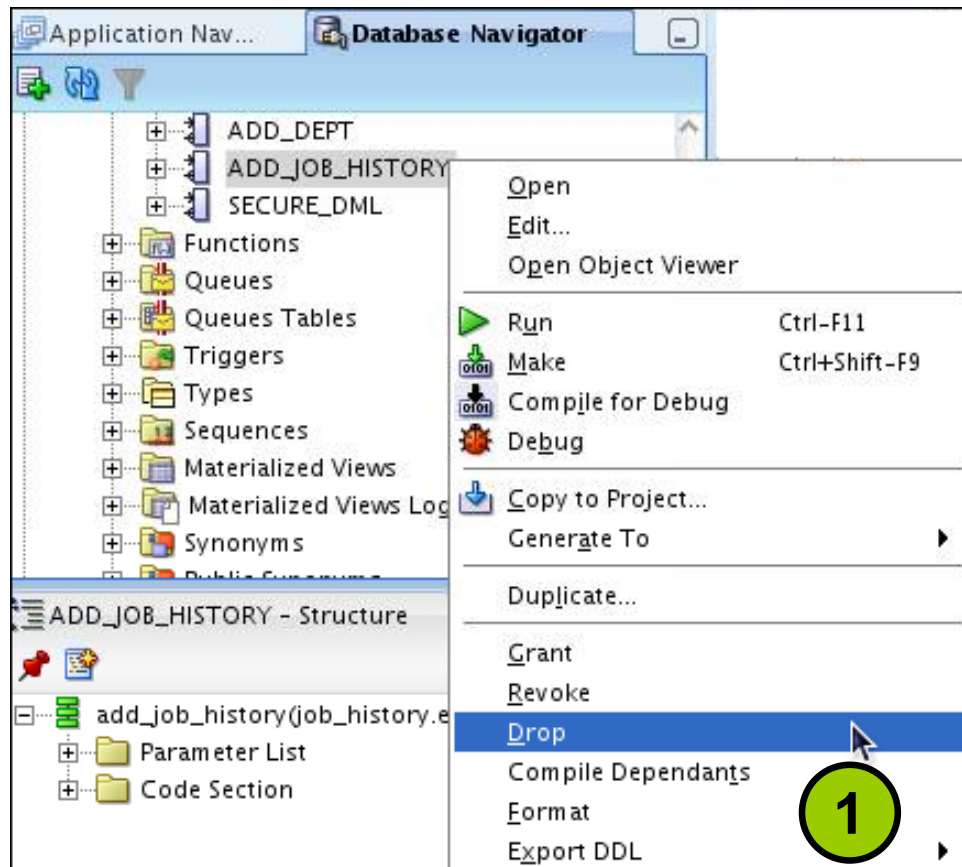
```
DECLARE
  P_EMP_ID NUMBER;
  P_START_DATE DATE;
  P_END_DATE DATE;
  P_JOB_ID VARCHAR2(10);
  P_DEPARTMENT_ID NUMBER;
BEGIN
  P_EMP_ID := NULL;
  P_START_DATE := NULL;
  P_END_DATE := NULL;
  P_JOB_ID := NULL;
  P_DEPARTMENT_ID := NULL;

  ADD_JOB_HISTORY(
    P_EMP_ID => P_EMP_ID,
```

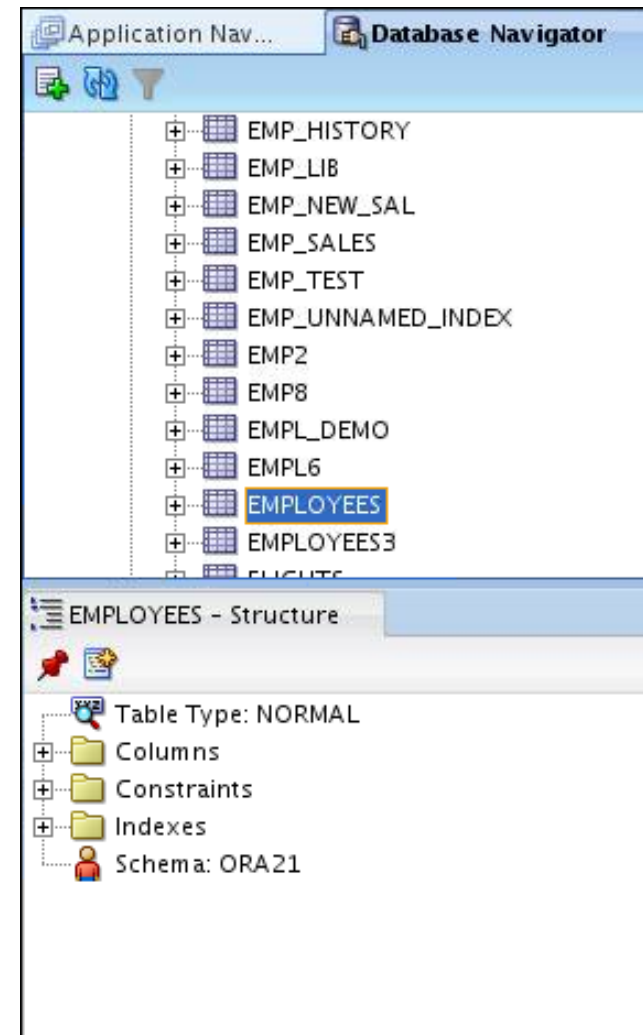
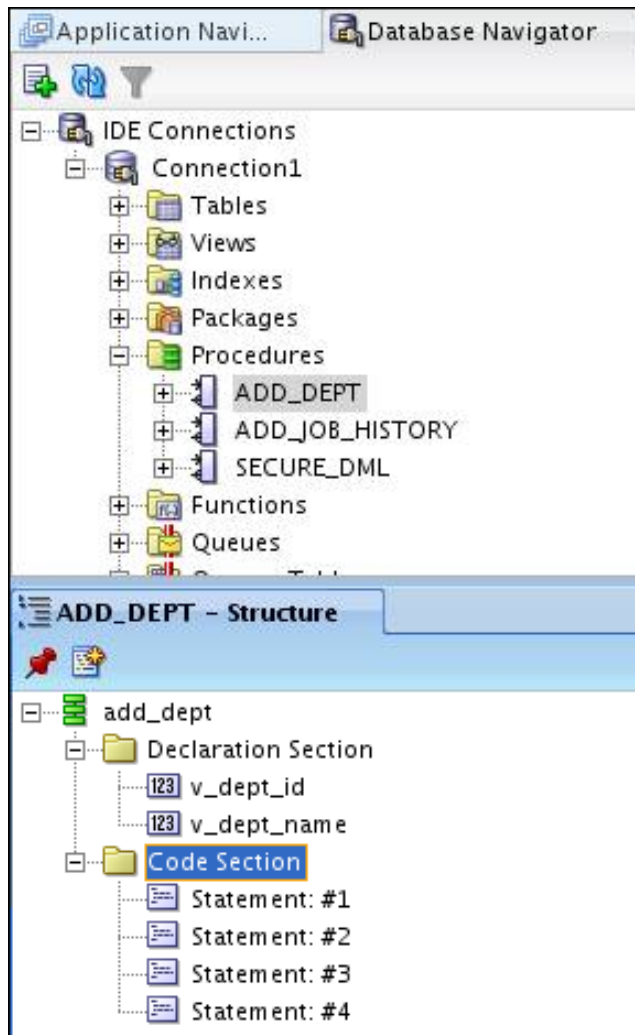
Save File... From File... Reset

Help OK Cancel

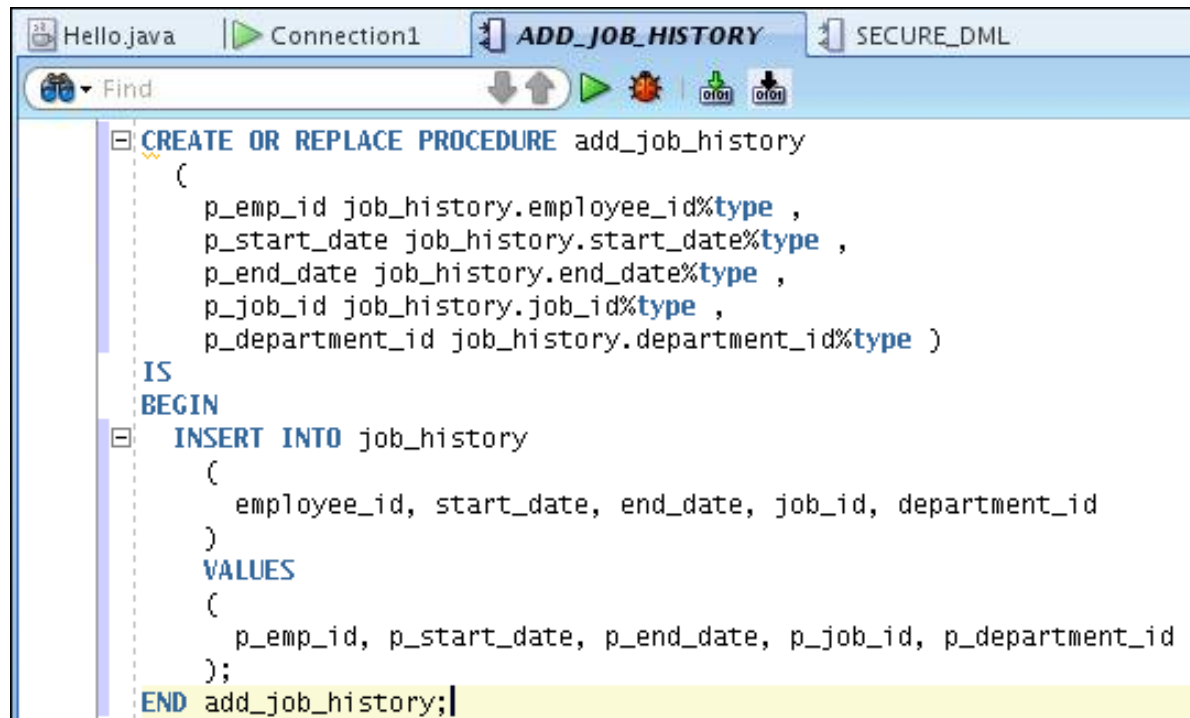
# Dropping a Program Unit



# Structure Window



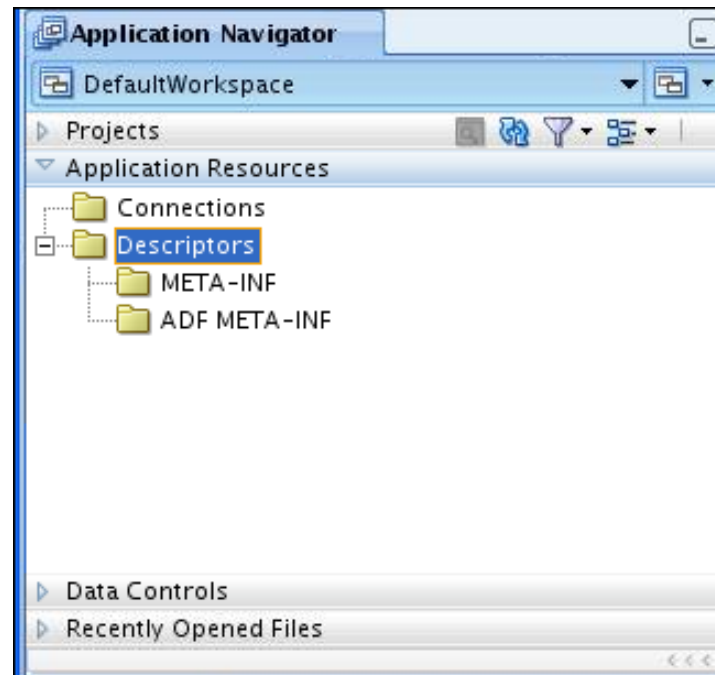
# Editor Window



The screenshot shows the Oracle SQL Developer Editor Window. The title bar includes tabs for 'Hello.java', 'Connection1', 'ADD\_JOB\_HISTORY', and 'SECURE\_DML'. Below the title bar is a toolbar with icons for Find, Run, and other functions. The main editor area displays the following PL/SQL code:

```
CREATE OR REPLACE PROCEDURE add_job_history
(
    p_emp_id job_history.employee_id%type ,
    p_start_date job_history.start_date%type ,
    p_end_date job_history.end_date%type ,
    p_job_id job_history.job_id%type ,
    p_department_id job_history.department_id%type )
IS
BEGIN
    INSERT INTO job_history
    (
        employee_id, start_date, end_date, job_id, department_id
    )
    VALUES
    (
        p_emp_id, p_start_date, p_end_date, p_job_id, p_department_id
    );
END add_job_history;
```


# Application Navigator



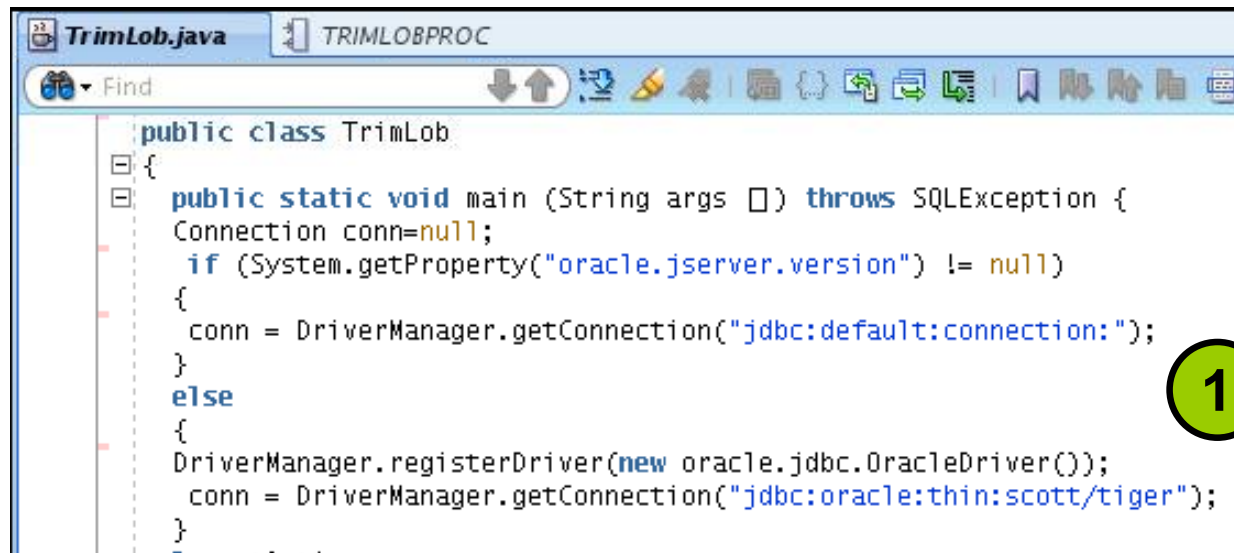
# Deploying Java Stored Procedures

Before deploying Java stored procedures, perform the following steps:

1. Create a database connection.
2. Create a deployment profile.
3. Deploy the objects.



# Publishing Java to PL/SQL



```
public class TrimLob
{
    public static void main (String args []) throws SQLException {
        Connection conn=null;
        if (System.getProperty("oracle.jserver.version") != null)
        {
            conn = DriverManager.getConnection("jdbc:default:connection:");
        }
        else
        {
            DriverManager.registerDriver(new oracle.jdbc.OracleDriver());
            conn = DriverManager.getConnection("jdbc:oracle:thin:scott/tiger");
        }
    }
}
```



```
CREATE OR REPLACE PROCEDURE TRIMLOBPROC
as language java
    name 'TrimLob.main(java.lang.String[])';
/
```



# How Can I Learn More About JDeveloper 11g ?

Topic	Web site
Oracle JDeveloper Product Page	<a href="http://www.oracle.com/technology/products/jdev/index.html">http://www.oracle.com/technology/products/jdev/index.html</a>
Oracle JDeveloper 11g Tutorials	<a href="http://www.oracle.com/technology/obe/obe11jdev/11/index.html">http://www.oracle.com/technology/obe/obe11jdev/11/index.html</a>
Oracle JDeveloper 11g Product Documentation	<a href="http://www.oracle.com/technology/documentation/jdev.html">http://www.oracle.com/technology/documentation/jdev.html</a>
Oracle JDeveloper 11g Discussion Forum	<a href="http://forums.oracle.com/forums/forum.jspa?forumID=83">http://forums.oracle.com/forums/forum.jspa?forumID=83</a>

# Summary

In this appendix, you should have learned how to use JDeveloper to do the following:

- List the key features of Oracle JDeveloper
- Create a database connection in JDeveloper
- Manage database objects in JDeveloper
- Use JDeveloper to execute SQL Commands
- Create and run PL/SQL Program Units



# **Generating Reports by Grouping Related Data**

# Objectives

After completing this appendix, you should be able to use the:

- ROLLUP operation to produce subtotal values
- CUBE operation to produce cross-tabulation values
- GROUPING function to identify the row values created by ROLLUP or CUBE
- GROUPING SETS to produce a single result set

# Review of Group Functions

- Group functions operate on sets of rows to give one result per group.

```
SELECT      [column,] group_function(column) . . .
FROM        table
[WHERE      condition]
[GROUP BY   group_by_expression]
[ORDER BY   column];
```

- Example:

```
SELECT AVG(salary), STDDEV(salary),
       COUNT(commission_pct), MAX(hire_date)
FROM   employees
WHERE  job_id LIKE 'SA%';
```

# Review of the GROUP BY Clause

- Syntax:

```
SELECT      [column,] group_function(column) . . .  
FROM        table  
[WHERE      condition]  
[GROUP BY   group_by_expression]  
[ORDER BY   column];
```

- Example:

```
SELECT      department_id, job_id, SUM(salary),  
            COUNT(employee_id)  
FROM        employees  
GROUP BY    department_id, job_id ;
```

# Review of the HAVING Clause

- Use the HAVING clause to specify which groups are to be displayed.
- You further restrict the groups on the basis of a limiting condition.

```
SELECT      [column,] group_function(column) ...  
FROM        table  
[WHERE      condition]  
[GROUP BY   group_by_expression]  
[HAVING     having_expression]  
[ORDER BY   column] ;
```

## **GROUP BY with ROLLUP and CUBE Operators**

- Use ROLLUP or CUBE with GROUP BY to produce superaggregate rows by cross-referencing columns.
- ROLLUP grouping produces a result set containing the regular grouped rows and the subtotal values.
- CUBE grouping produces a result set containing the rows from ROLLUP and cross-tabulation rows.



# ROLLUP Operator

- ROLLUP is an extension to the GROUP BY clause.
- Use the ROLLUP operation to produce cumulative aggregates, such as subtotals.

```
SELECT      [column,] group_function(column) . . .  
FROM        table  
[WHERE      condition]  
[GROUP BY   [ROLLUP] group_by_expression]  
[HAVING     having_expression];  
[ORDER BY   column];
```

# ROLLUP Operator: Example

```
SELECT  department_id, job_id, SUM(salary)
FROM    employees
WHERE   department_id < 60
GROUP BY ROLLUP(department_id, job_id);
```

	DEPARTMENT_ID	JOB_ID	SUM(SALARY)
1	10	AD_ASST	4400
2	10	(null)	4400
3	20	MK_MAN	13000
4	20	MK_REP	6000
5	20	(null)	19000
6	30	PU_MAN	11000
7	30	PU_CLERK	13900
8	30	(null)	24900
9	40	HR_REP	6500
10	40	(null)	6500
11	50	ST_MAN	36400
12	50	SH_CLERK	64300
13	50	ST_CLERK	55700
14	50	(null)	156400
15	(null)	(null)	211200

1

2

3

# CUBE Operator

- CUBE is an extension to the GROUP BY clause.
- You can use the CUBE operator to produce cross-tabulation values with a single SELECT statement.

```
SELECT      [column,] group_function(column) ...  
FROM        table  
[WHERE      condition]  
[GROUP BY   [CUBE] group_by_expression]  
[HAVING     having_expression]  
[ORDER BY   column];
```

# CUBE Operator: Example

```
SELECT    department_id, job_id, SUM(salary)
FROM      employees
WHERE     department_id < 60
GROUP BY  CUBE (department_id, job_id) ;
```

	DEPARTMENT_ID	JOB_ID	SUM(SALARY)
1	(null)	(null)	211200
2	(null)	HR_REP	6500
3	(null)	MK_MAN	13000
4	(null)	MK_REP	6000
5	(null)	PU_MAN	11000
6	(null)	ST_MAN	36400
7	(null)	AD_ASST	4400
8	(null)	PU_CLERK	13900
9	(null)	SH_CLERK	64300
10	(null)	ST_CLERK	55700
11	10	(null)	4400
12	10	AD_ASST	4400
13	20	(null)	19000
14	20	MK_MAN	13000
15	20	MK_REP	6000
16	30	(null)	24900

1

2

3

4

# GROUPING Function

The GROUPING function:

- Is used with either the CUBE or ROLLUP operator
- Is used to find the groups forming the subtotal in a row
- Is used to differentiate stored NULL values from NULL values created by ROLLUP or CUBE
- Returns 0 or 1

```
SELECT      [column,] group_function(column) .. ,  
            GROUPING(expr)  
FROM        table  
[WHERE      condition]  
[GROUP BY  [ROLLUP] [CUBE] group_by_expression]  
[HAVING    having_expression]  
[ORDER BY  column];
```

# GROUPING Function: Example

```
SELECT    department_id DEPTID, job_id JOB,
          SUM(salary),
          GROUPING(department_id) GRP_DEPT,
          GROUPING(job_id) GRP_JOB
FROM      employees
WHERE     department_id < 50
GROUP BY  ROLLUP(department_id, job_id);
```

	DEPTID	JOB	SUM(SALARY)	GRP_DEPT	GRP_JOB
1	10	AD_ASST	4400	0	0
2	10	(null)	4400	0	1
3	20	MK_MAN	13000	0	0
4	20	MK_REP	6000	0	0
5	20	(null)	19000	0	1
6	30	PU_MAN	11000	0	0
7	30	PU_CLERK	13900	0	0
8	30	(null)	24900	0	1
9	40	HR_REP	6500	0	0
10	40	(null)	6500	0	1
11	(null)	(null)	54800	1	1

# GROUPING SETS

- The GROUPING SETS syntax is used to define multiple groupings in the same query.
- All groupings specified in the GROUPING SETS clause are computed and the results of individual groupings are combined with a UNION ALL operation.
- Grouping set efficiency:
  - Only one pass over the base table is required.
  - There is no need to write complex UNION statements.
  - The more elements GROUPING SETS has, the greater is the performance benefit.

# GROUPING SETS: Example

```
SELECT    department_id, job_id,  
          manager_id,AVG(salary)  
FROM      employees  
GROUP BY GROUPING SETS  
        ((department_id,job_id), (job_id,manager_id));
```

	DEPARTMENT_ID	JOB_ID	MANAGER_ID	AVG(SALARY)
1	(null)	SH_CLERK	122	3200
2	(null)	AC_MGR	101	12000
3	(null)	ST_MAN	100	7280
4	...	ST_CLERK	121	2675

1

	DEPARTMENT_ID	JOB_ID	MANAGER_ID	AVG(SALARY)
39	110	AC_MGR	(null)	12000
40	90	AD_PRES	(null)	24000
41	60	IT_PROG	(null)	5760
42	100	FI_MGR	(null)	12000

2

...



# Composite Columns

- A composite column is a collection of columns that are treated as a unit.

```
ROLLUP (a, (b, c), d)
```

- Use parentheses within the GROUP BY clause to group columns, so that they are treated as a unit while computing ROLLUP or CUBE operations.
- When used with ROLLUP or CUBE, composite columns would require skipping aggregation across certain levels.

# Composite Columns: Example

```
SELECT    department_id, job_id, manager_id,
          SUM(salary)
FROM      employees
GROUP BY ROLLUP( department_id, (job_id, manager_id));
```

	DEPARTMENT_ID	JOB_ID	MANAGER_ID	SUM(SALARY)
1	(null)	SA_REP	149	7000
2	(null)	(null)	(null)	7000
3	10	AD_ASST	101	4400
4	10	(null)	(null)	4400
5	20	MK_MAN	100	13000
6	20	MK_REP	201	6000
7	20	(null)	(null)	19000
...				
40	100	FI_MGR	101	12000
41	100	FI_ACCOUNT	108	39600
42	100	(null)	(null)	51600
43	110	AC_MGR	101	12000
44	110	AC_ACCOUNT	205	8300
45	110	(null)	(null)	20300
46	(null)	(null)	(null)	691400

# Concatenated Groupings

- Concatenated groupings offer a concise way to generate useful combinations of groupings.
- To specify concatenated grouping sets, you separate multiple grouping sets, ROLLUP and CUBE operations with commas so that the Oracle server combines them into a single GROUP BY clause.
- The result is a cross-product of groupings from each GROUPING SET.

```
GROUP BY GROUPING SETS (a, b), GROUPING SETS (c, d)
```

# Concatenated Groupings: Example

```
SELECT  department_id, job_id, manager_id,
        SUM(salary)
FROM    employees
GROUP BY department_id,
        ROLLUP (job_id),
        CUBE (manager_id);
```

	DEPARTMENT_ID	JOB_ID	MANAGER_ID	SUM(SALARY)
1	(null)	SA_REP	149	7000
2	10	AD_ASST	101	4400
3	20	MK_MAN	100	13000
4	20	MK_REP	201	6000

1

...

	90	AD_VP	100	34000
	90	AD_PREP	(null)	24000

...

	(null)	SA_REP	(null)	7000
	10	AD_ASST	(null)	4400

...

2

	110	(null)	101	12000
	110	(null)	205	8300
	110	(null)	(null)	20300

3

# Summary

In this appendix, you should have learned how to use the:

- ROLLUP operation to produce subtotal values
- CUBE operation to produce cross-tabulation values
- GROUPING function to identify the row values created by ROLLUP or CUBE
- GROUPING SETS syntax to define multiple groupings in the same query
- GROUP BY clause to combine expressions in various ways:
  - Composite columns
  - Concatenated grouping sets



# **Hierarchical Retrieval**

# Objectives

After completing this appendix, you should be able to do the following:

- Interpret the concept of a hierarchical query
- Create a tree-structured report
- Format hierarchical data
- Exclude branches from the tree structure

# Sample Data from the EMPLOYEES Table

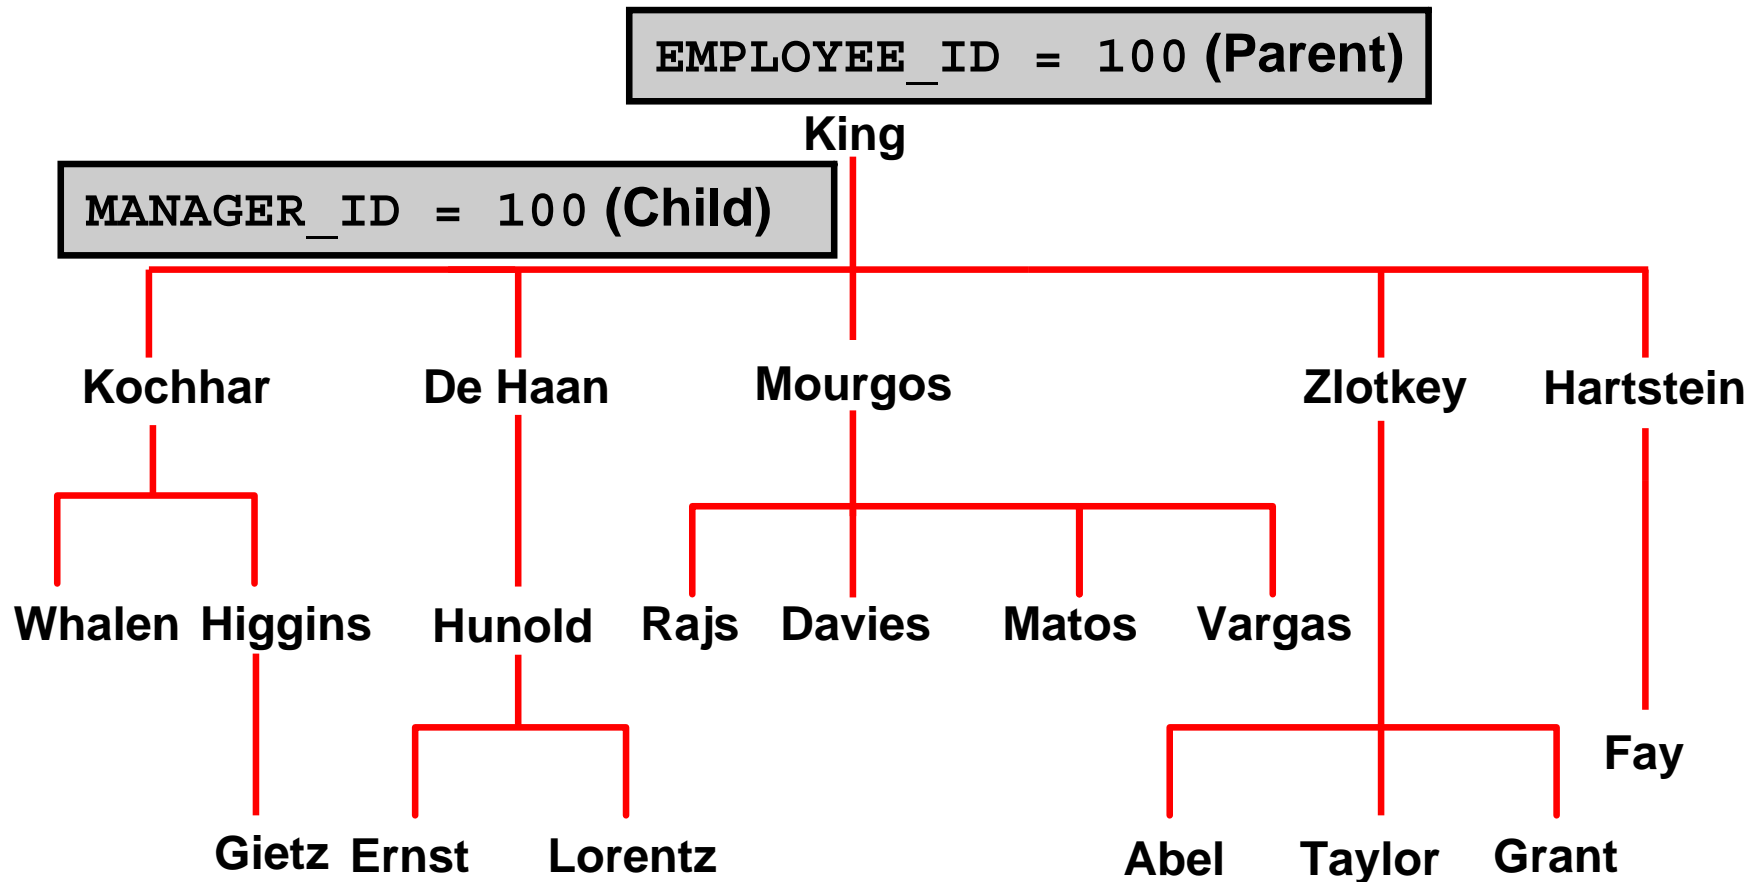
	EMPLOYEE_ID	LAST_NAME	JOB_ID	MANAGER_ID
1	100	King	AD_PRES	(null)
2	101	Kochhar	AD_VP	100
3	102	De Haan	AD_VP	100
4	103	Hunold	IT_PROG	102
5	104	Ernst	IT_PROG	103
6	107	Lorentz	IT_PROG	103

...

16	200	Whalen	AD_ASST	101
17	201	Hartstein	MK_MAN	100
18	202	Fay	MK_REP	201
19	205	Higgins	AC_MGR	101
20	206	Gietz	AC_ACCOUNT	205



# Natural Tree Structure



# Hierarchical Queries

```
SELECT [LEVEL], column, expr...  
FROM   table  
[WHERE condition(s)]  
[START WITH condition(s)]  
[CONNECT BY PRIOR condition(s)] ;
```

*condition:*

```
expr comparison_operator expr
```

# Walking the Tree

## Starting Point

- Specifies the condition that must be met
- Accepts any valid condition

```
START WITH column1 = value
```

Using the EMPLOYEES table, start with the employee whose last name is Kochhar.

```
...START WITH last_name = 'Kochhar'
```

# Walking the Tree

CONNECT BY PRIOR *column1* = *column2*

Walk from the top down, using the `EMPLOYEES` table.

```
... CONNECT BY PRIOR employee id = manager id
```

## Direction

Top down      →      Column1 = Parent Key  
                                Column2 = Child Key

Bottom up  $\longrightarrow$  Column1 = Child Key  
Column2 = Parent Key

# Walking the Tree: From the Bottom Up

```
SELECT employee_id, last_name, job_id, manager_id
FROM   employees
START WITH employee_id = 101
CONNECT BY PRIOR manager_id = employee_id ;
```

	EMPLOYEE_ID	LAST_NAME	JOB_ID	MANAGER_ID
1	101	Kochhar	AD_VP	100
2	100	King	AD_PRES	(null)

# Walking the Tree: From the Top Down

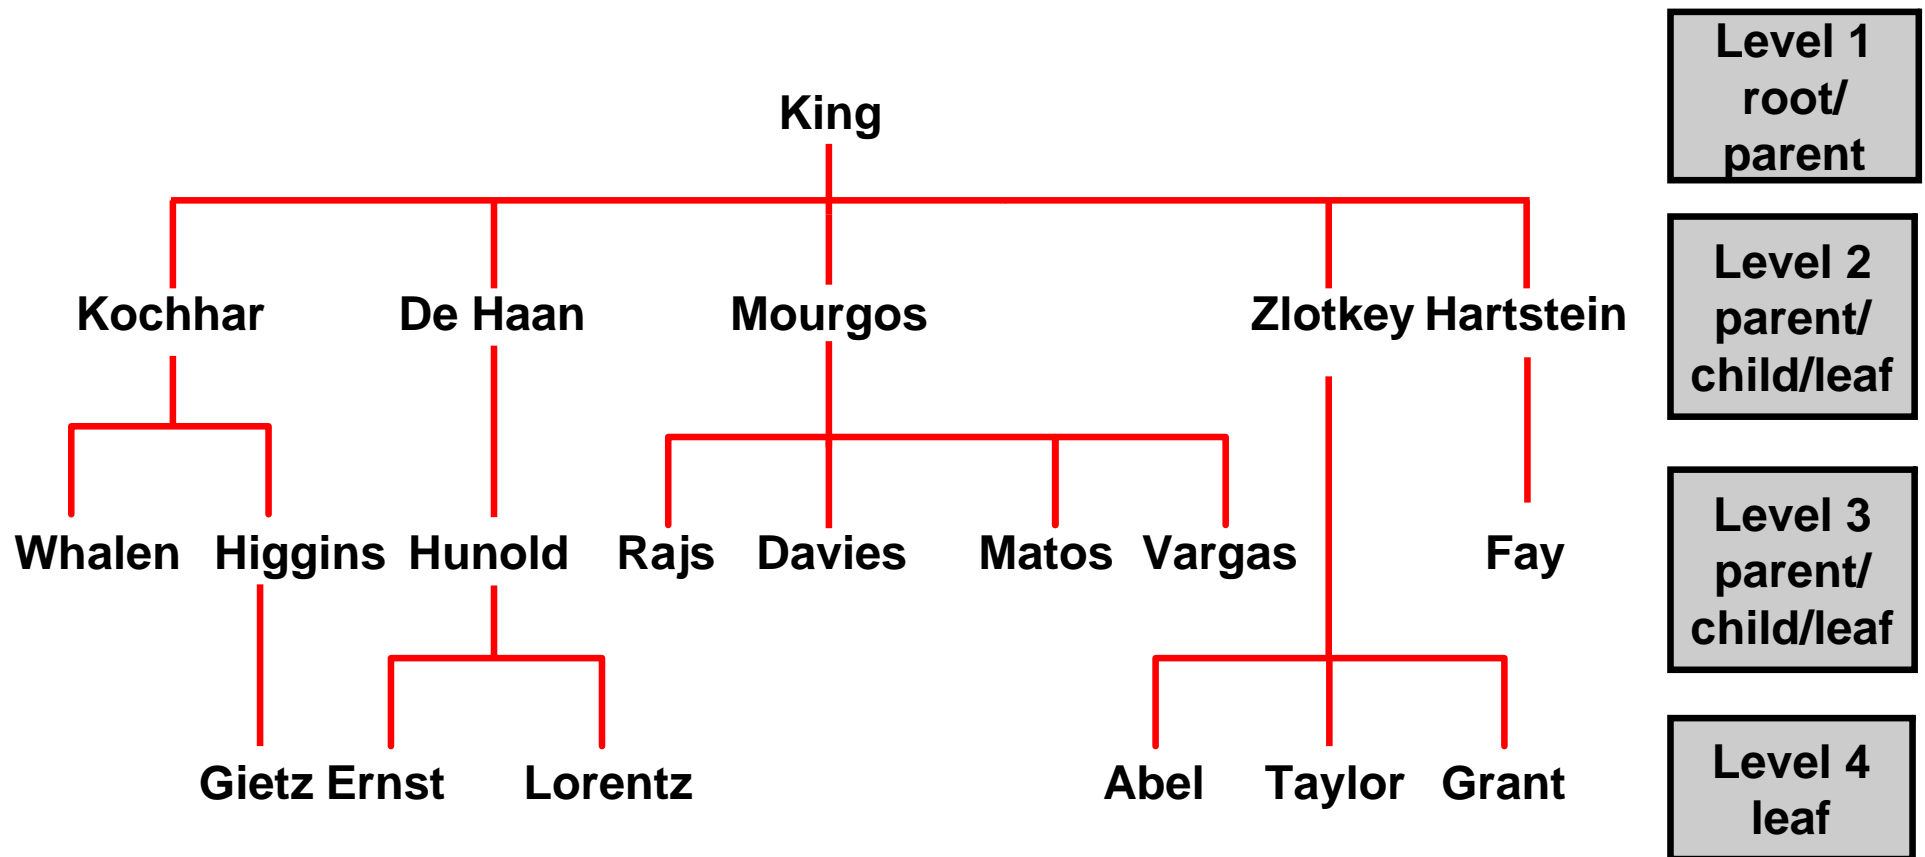
```
SELECT last_name || ' reports to ' ||  
PRIOR last_name "Walk Top Down"  
FROM employees  
START WITH last_name = 'King'  
CONNECT BY PRIOR employee_id = manager_id ;
```

Walk Top Down	
1	King reports to
2	King reports to
3	Kochhar reports to King
4	Greenberg reports to Kochhar
5	Faviet reports to Greenberg

...

105	Grant reports to Zlotkey
106	Johnson reports to Zlotkey
107	Hartstein reports to King
108	Fay reports to Hartstein

# Ranking Rows with the `LEVEL` Pseudocolumn



# Formatting Hierarchical Reports Using LEVEL and LPAD

Create a report displaying company management levels, beginning with the highest level and indenting each of the following levels.

```
COLUMN org_chart FORMAT A12
SELECT LPAD(last_name, LENGTH(last_name) + (LEVEL*2) - 2, '_')
        AS org_chart
FROM    employees
START WITH first_name='Steven' AND last_name='King'
CONNECT BY PRIOR employee_id=manager_id
```





# Pruning Branches

Use the **WHERE** clause  
to eliminate a node.

Use the **CONNECT BY** clause  
to eliminate a branch.

```
WHERE last_name != 'Higgins' CONNECT BY PRIOR  
employee_id = manager_id  
AND last_name != 'Higgins'
```



# Summary

In this appendix, you should have learned that you can:

- Use hierarchical queries to view a hierarchical relationship between rows in a table
- Specify the direction and starting point of the query
- Eliminate nodes or branches by pruning

# Writing Advanced Scripts


# Objectives

After completing this appendix, you should be able to do the following:

- Describe the type of problems that are solved by using SQL to generate SQL
- Write a script that generates a script of `DROP TABLE` statements
- Write a script that generates a script of `INSERT INTO` statements

# Using SQL to Generate SQL

- SQL can be used to generate scripts in SQL.
- The data dictionary is:
  - A collection of tables and views that contain database information
  - Created and maintained by the Oracle server



# Creating a Basic Script

```
SELECT 'CREATE TABLE ' || table_name ||  
      '_test ' || 'AS SELECT * FROM '  
      || table_name || ' WHERE 1=2;'  
      AS "Create Table Script"  
FROM   user tables;
```

	Create Table Script
1	CREATE TABLE REGIONS_test AS SELECT * FROM REGIONS WHERE 1=2;
2	CREATE TABLE LOCATIONS_test AS SELECT * FROM LOCATIONS WHERE 1=2;
3	CREATE TABLE DEPARTMENTS_test AS SELECT * FROM DEPARTMENTS WHERE 1=2;
4	CREATE TABLE JOBS_test AS SELECT * FROM JOBS WHERE 1=2;
5	CREATE TABLE EMPLOYEES_test AS SELECT * FROM EMPLOYEES WHERE 1=2;
6	CREATE TABLE JOB_HISTORY_test AS SELECT * FROM JOB_HISTORY WHERE 1=2;

# Controlling the Environment

```
SET ECHO OFF  
SET FEEDBACK OFF  
SET PAGESIZE 0
```

Set system variables  
to appropriate values.



The diagram consists of a large light-gray rectangular box with a black border. Inside the box, the text is organized into two distinct sections. The top section contains three lines of SQL code: 'SET ECHO OFF', 'SET FEEDBACK OFF', and 'SET PAGESIZE 0'. Below these, there is a line of text 'SQL statement' centered horizontally. The bottom section contains three lines of SQL code: 'SET FEEDBACK ON', 'SET PAGESIZE 24', and 'SET ECHO ON'. To the right of the box, there are two lines of explanatory text. The first line, 'Set system variables to appropriate values.', has a black arrow pointing from the text to the top section of the box. The second line, 'Set system variables back to the default value.', has a black arrow pointing from the text to the bottom section of the box.

**SQL statement**

```
SET FEEDBACK ON  
SET PAGESIZE 24  
SET ECHO ON
```

Set system variables  
back to the default  
value.

# The Complete Picture

```
SET ECHO OFF
SET FEEDBACK OFF
SET PAGESIZE 0

SELECT 'DROP TABLE ' || object_name || ';'
FROM   user_objects
WHERE  object_type = 'TABLE'
/

SET FEEDBACK ON
SET PAGESIZE 24
SET ECHO ON
```



# Dumping the Contents of a Table to a File

```
SET HEADING OFF ECHO OFF FEEDBACK OFF
SET PAGESIZE 0

SELECT
  'INSERT INTO departments_test VALUES
    (' || department_id || ', ''' || department_name ||
    ''', ''' || location_id || ''');'
  AS "Insert Statements Script"
FROM   departments
/

SET PAGESIZE 24
SET HEADING ON ECHO ON FEEDBACK ON
```

# Dumping the Contents of a Table to a File

Source	Result
<code>'''X'''</code>	<code>'X'</code>
<code>'''</code>	<code>'</code>
<code>'''    department_name    '''</code>	<code>'Administration'</code>
<code>''', '''</code>	<code>','</code>
<code>''' ) ; '</code>	<code>' ) ;</code>

# Generating a Dynamic Predicate

```
COLUMN my_col NEW_VALUE dyn_where_clause

SELECT DECODE('&&deptno', null,
DECODE ('&&hiredate', null, ' ',
'WHERE hire_date=TO_DATE('' || '&&hiredate'', ''DD-MON-YYYY'')'),
DECODE ('&&hiredate', null,
'WHERE department_id = ' || '&&deptno',
'WHERE department_id = ' || '&&deptno' ||
' AND hire_date = TO_DATE('' || '&&hiredate'', ''DD-MON-YYYY'')'))
AS my_col FROM dual;
```

```
SELECT last_name FROM employees &dyn_where_clause;
```

# Summary

In this appendix, you should have learned that:

- You can write a SQL script to generate another SQL script
- Script files often use the data dictionary
- You can capture the output in a file




# Oracle Database Architectural Components

# Objectives

After completing this appendix, you should be able to do the following:

- List the major database architectural components
- Describe the background processes
- Explain the memory structures
- Correlate the logical and physical storage structures




# Oracle Database Architecture: Overview

The Oracle Relational Database Management System (RDBMS) is a database management system that provides an open, comprehensive, integrated approach to information management.



# Oracle Database Server Structures




ORACLE




# Connecting to the Database

- Connection: Communication pathway between a user process and a database instance
- Session: A specific connection of a user to a database instance through a user process



# Interacting with an Oracle Database



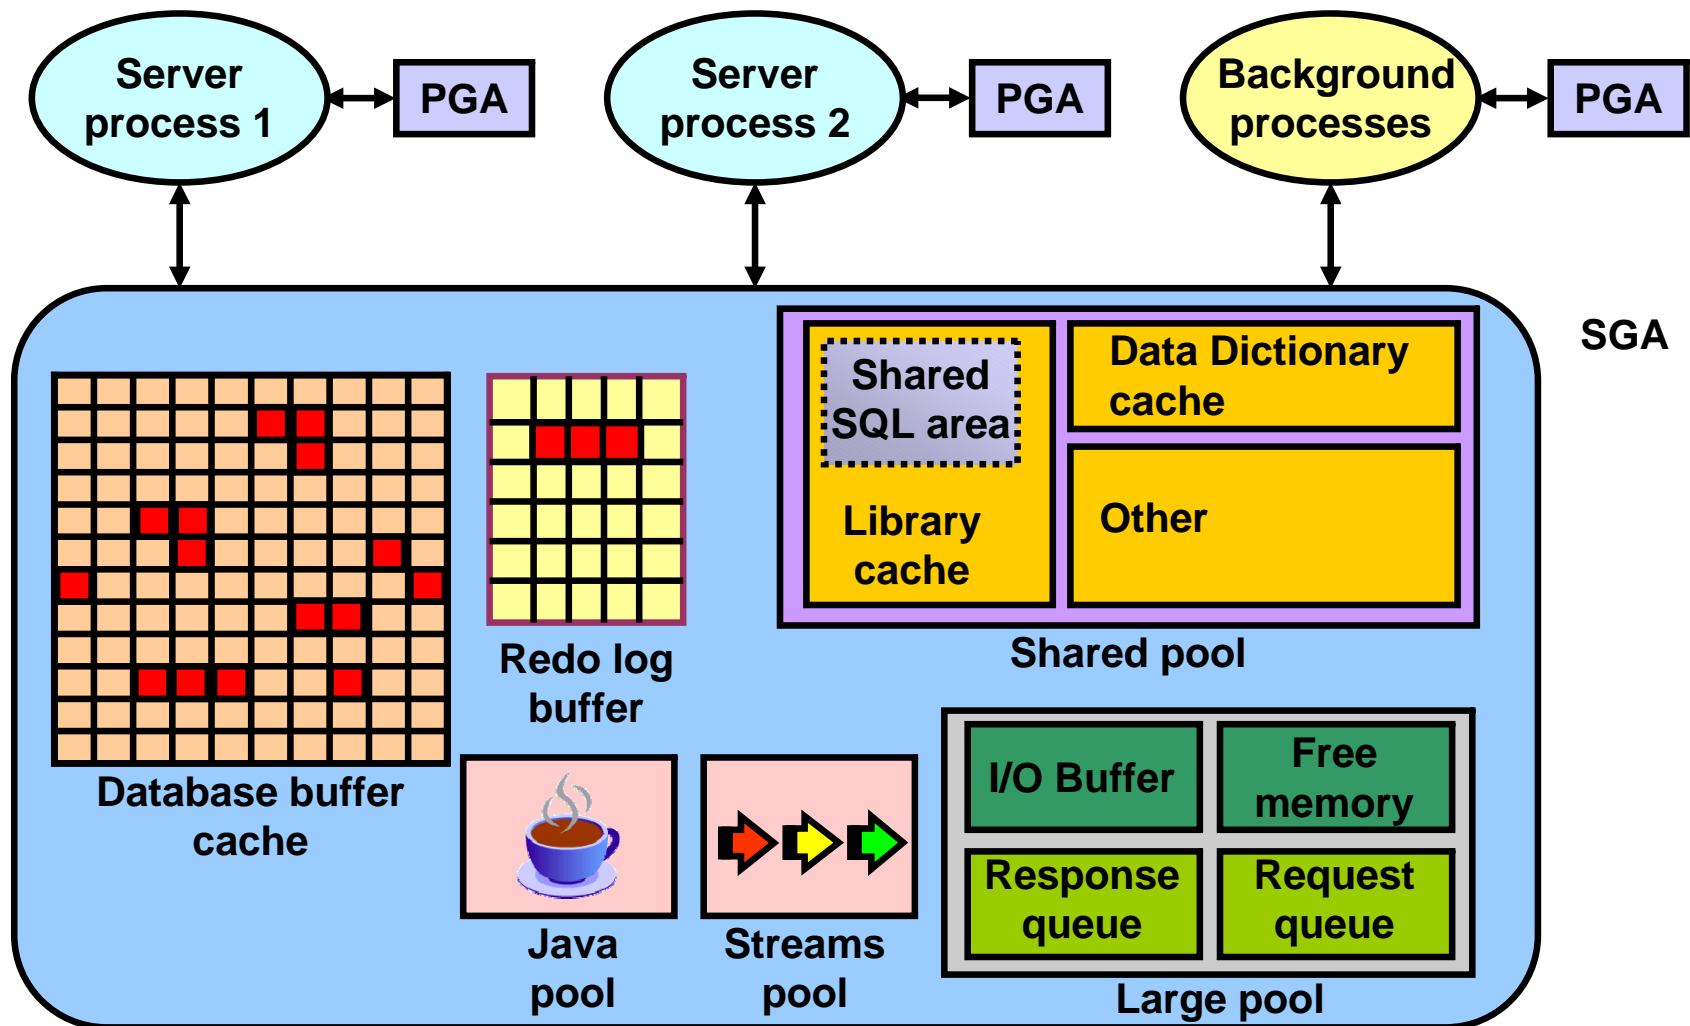
# Oracle Memory Architecture

DB structures

→ Memory

- Process

- Storage




# Process Architecture

## DB structures

- Memory
- **Process**
- Storage

- User process:
  - Is started when a database user or a batch process connects to the Oracle Database
- Database processes:
  - Server process: Connects to the Oracle instance and is started when a user establishes a session
  - Background processes: Are started when an Oracle instance is started




ORACLE

# Database Writer Process


Writes modified (dirty) buffers in the database buffer cache to disk:

- Asynchronously while performing other processing
- Periodically to advance the checkpoint



# Log Writer Process


- Writes the redo log buffer to a redo log file on disk
- LGWR writes:
  - A process commits a transaction
  - When the redo log buffer is one-third full
  - Before a DBWn process writes modified buffers to disk



# Checkpoint Process


Records checkpoint information in:

- The control file
- Each datafile header



# System Monitor Process


- Performs recovery at instance startup
- Cleans up unused temporary segments





# Process Monitor Process


- Performs process recovery when a user process fails:
  - Cleans up the database buffer cache
  - Frees resources used by the user process
- Monitors sessions for idle session timeout
- Dynamically registers database services with listeners





# Oracle Database Storage Architecture

## DB structures


- Memory
- Process
- **Storage**



Parameter file  
Password file  
Network files  
Alert and trace files



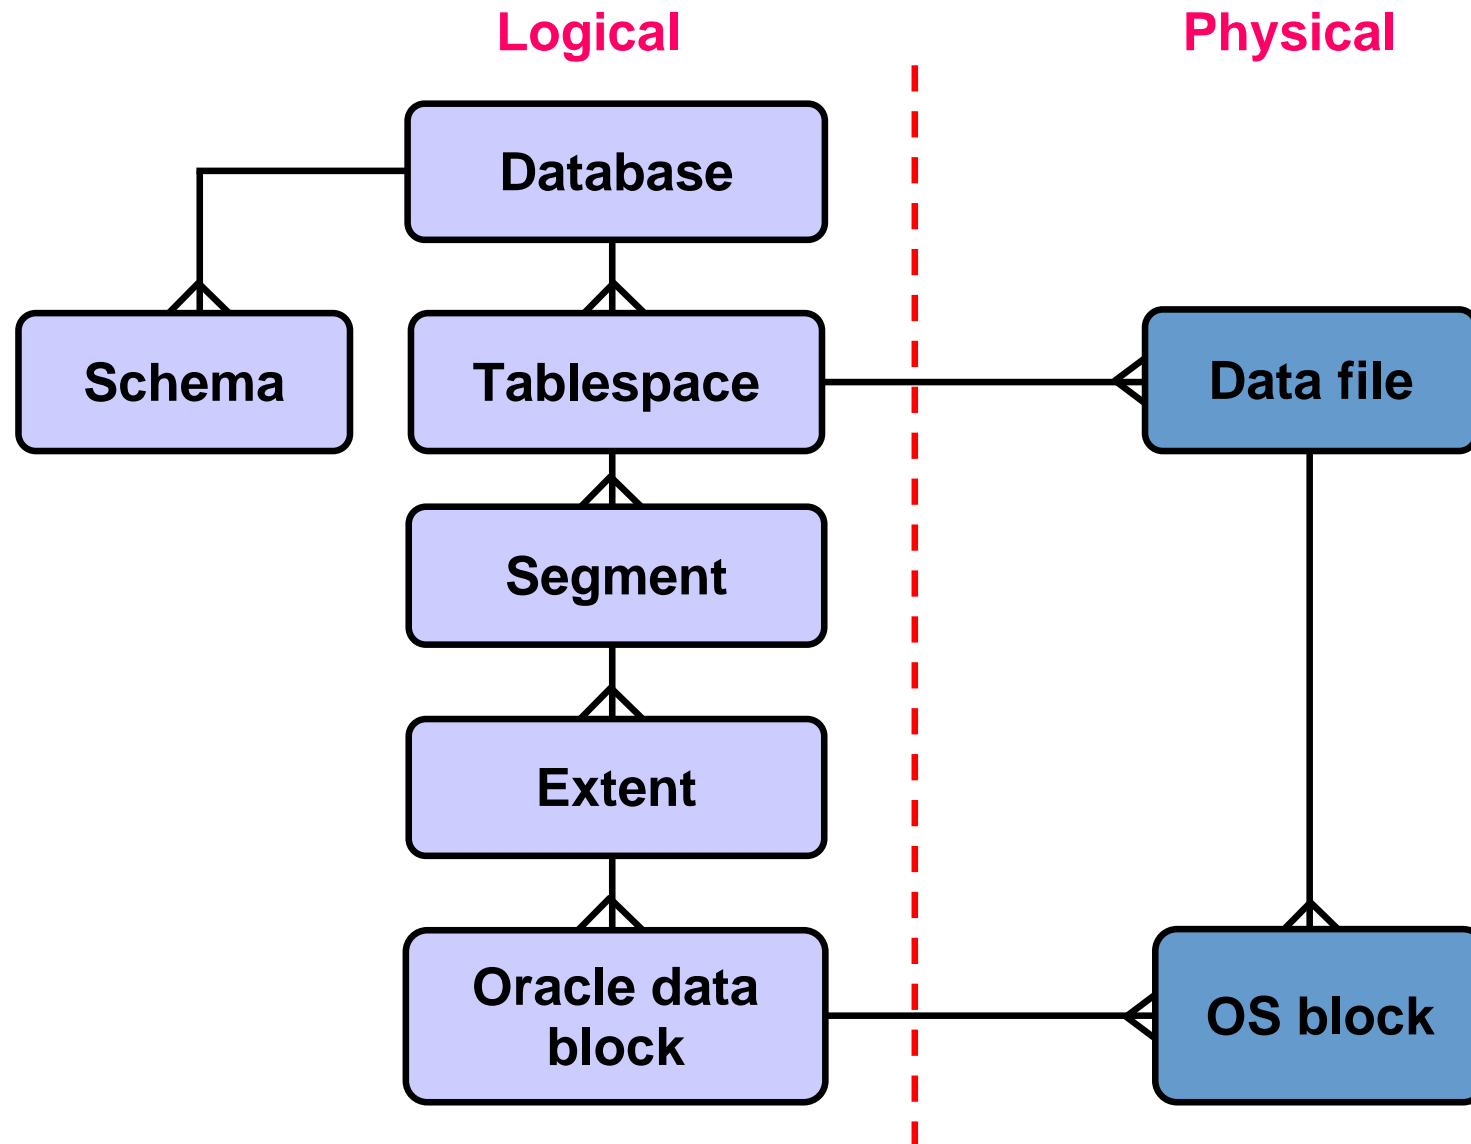
Backup files



Archived log files

ORACLE

# Logical and Physical Database Structures



# Processing a SQL Statement


- Connect to an instance using:
  - The user process
  - The server process
- The Oracle server components that are used depend on the type of SQL statement:
  - Queries return rows.
  - Data manipulation language (DML) statements log changes.
  - Commit ensures transaction recovery.
- Some Oracle server components do not participate in SQL statement processing.

# Processing a Query

- Parse:
  - Search for an identical statement.
  - Check the syntax, object names, and privileges.
  - Lock the objects used during parse.
  - Create and store the execution plan.
- Execute: Identify the rows selected.
- Fetch: Return the rows to the user process.

# Shared Pool

- The library cache contains the SQL statement text, parsed code, and execution plan.
- The data dictionary cache contains table, column, and other object definitions and privileges.
- The shared pool is sized by `SHARED_POOL_SIZE`.




# Database Buffer Cache

- The database buffer cache stores the most recently used blocks.
- The size of a buffer is based on `DB_BLOCK_SIZE`.
- The number of buffers is defined by `DB_BLOCK_BUFFERS`.



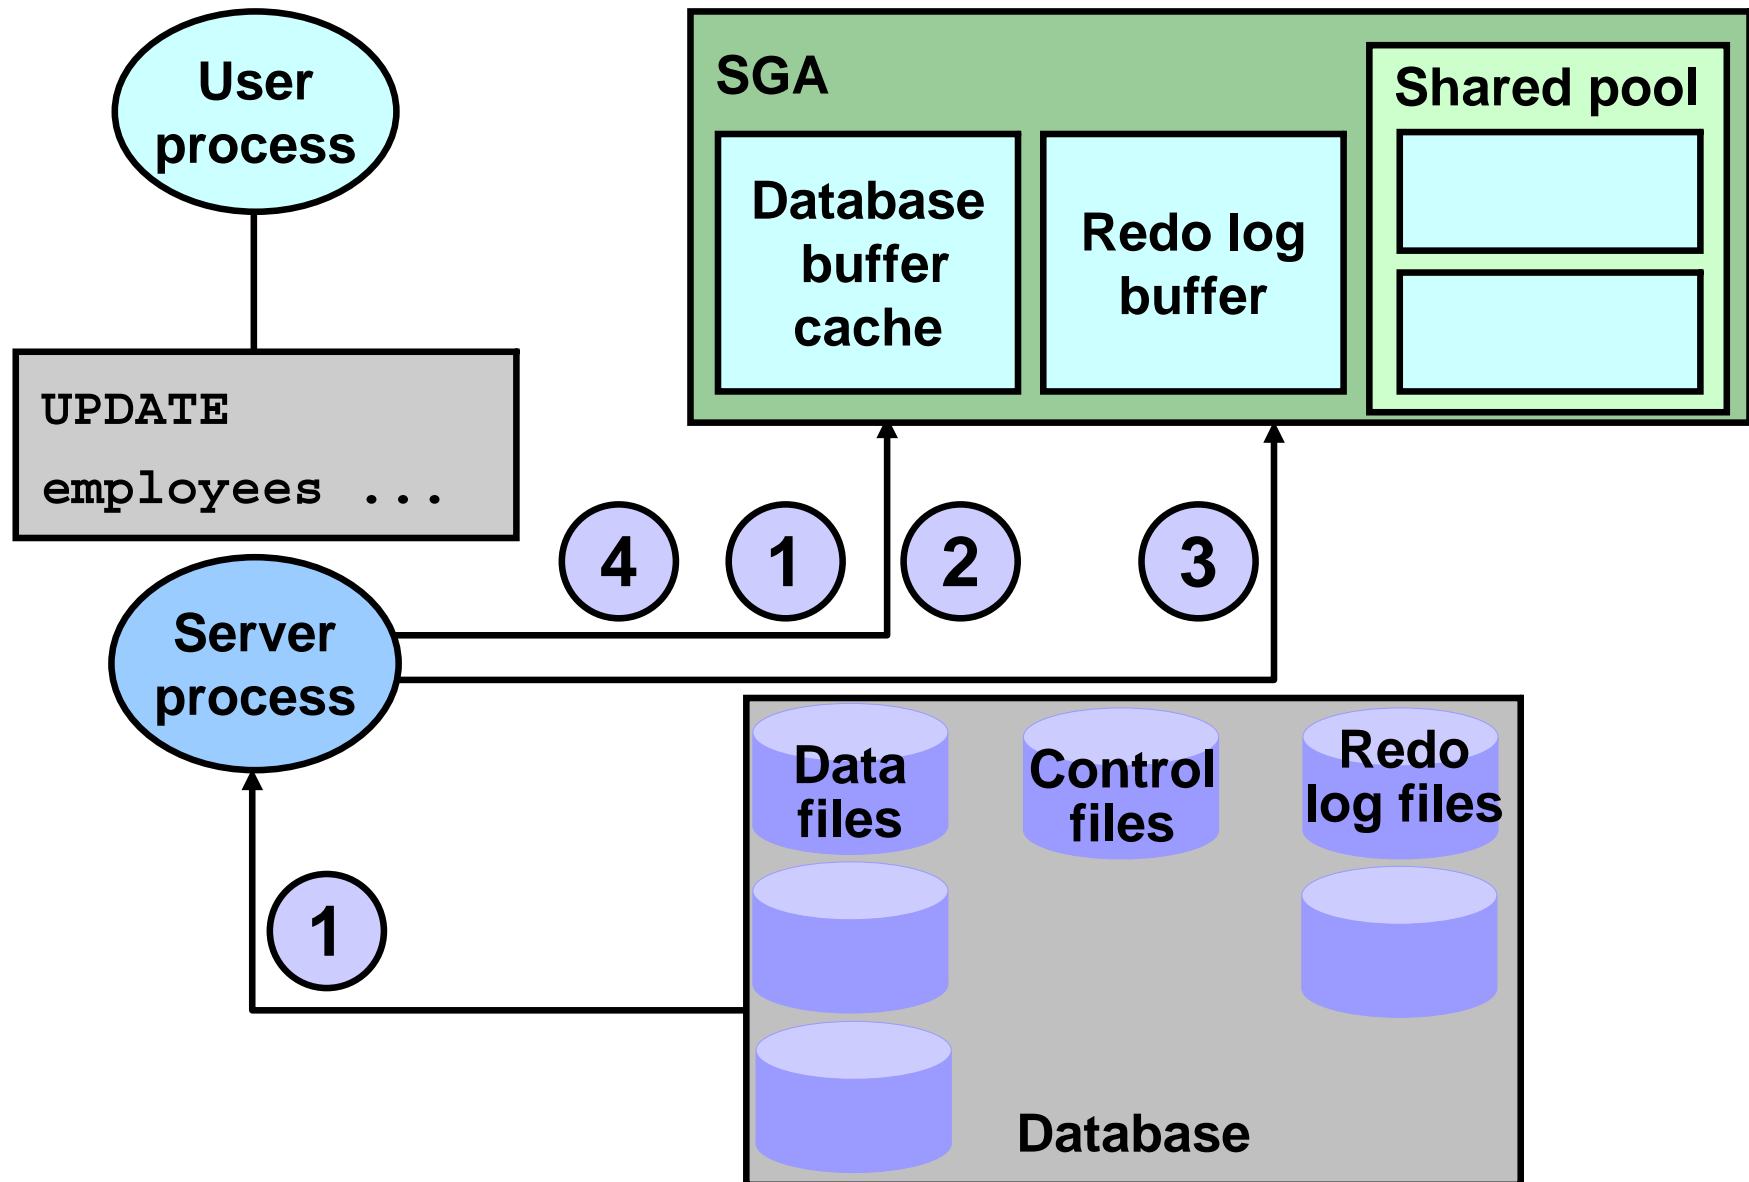
# Program Global Area (PGA)

- Is not shared
- Is writable only by the server process
- Contains:
  - Sort area
  - Session information
  - Cursor state
  - Stack space





# Processing a DML Statement




# Redo Log Buffer

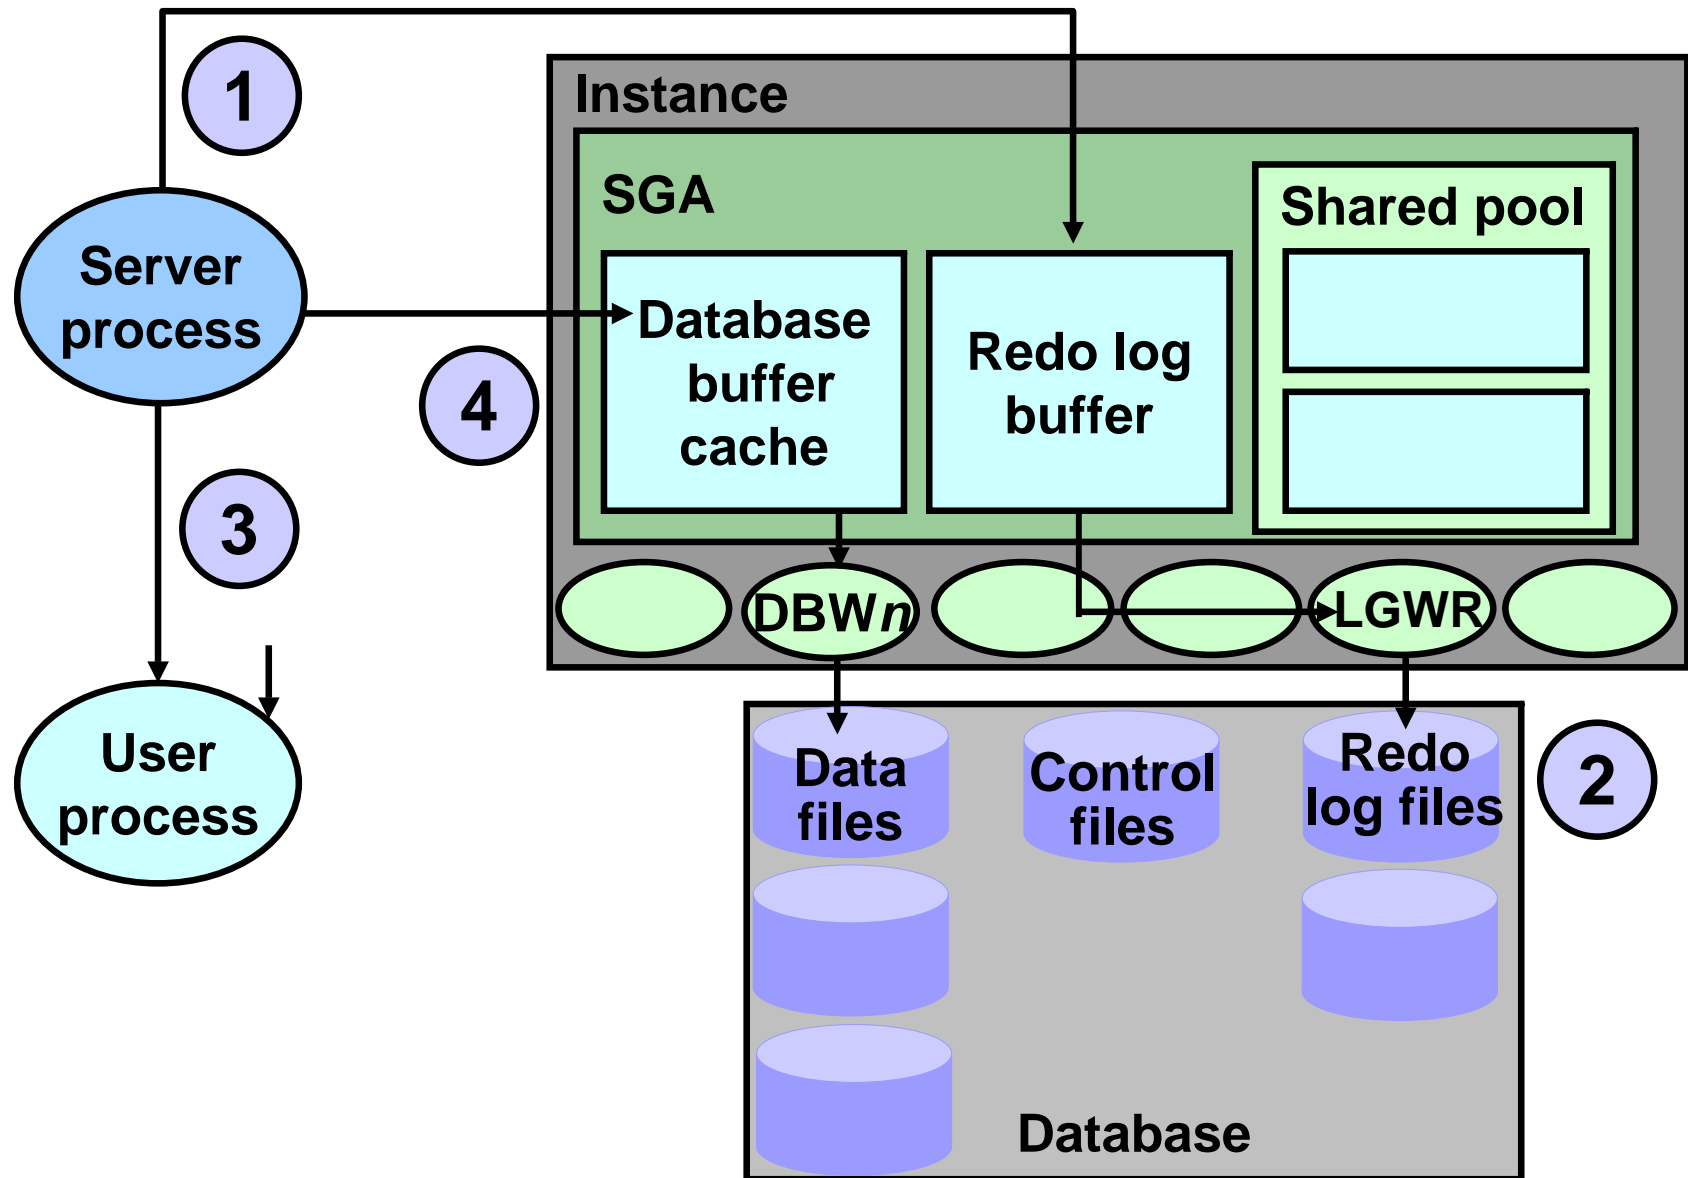
- Has its size defined by `LOG_BUFFER`
- Records changes made through the instance
- Is used sequentially
- Is a circular buffer



# Rollback Segment



# COMMIT Processing



# Summary of the Oracle Database Architecture

