



Manipulating Large Data Sets

Objectives

After completing this lesson, you should be able to do the following:

- Manipulate data by using subqueries
- Specify explicit default values in the `INSERT` and `UPDATE` statements
- Describe the features of multitable `INSERTs`
- Use the following types of multitable `INSERTs`:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merge rows in a table
- Track the changes to data over a period of time

Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

Using Subqueries to Manipulate Data

You can use subqueries in data manipulation language (DML) statements to:

- Retrieve data by using an inline view
- Copy data from one table to another
- Update data in one table based on the values of another table
- Delete rows from one table based on rows in another table

Retrieving Data by Using a Subquery as Source

```
SELECT department_name, city
FROM departments
NATURAL JOIN (SELECT l.location_id, l.city, l.country_id
              FROM loc l
              JOIN countries c
              ON(l.country_id = c.country_id)
              JOIN regions USING(region_id)
              WHERE region_name = 'Europe');
```

	DEPARTMENT_NAME	CITY
1	Human Resources	London
2	Sales	Oxford
3	Public Relations	Munich

Inserting by Using a Subquery as a Target

```
INSERT INTO (SELECT l.location_id, l.city, l.country_id
              FROM   locations l
              JOIN   countries c
              ON(l.country_id = c.country_id)
              JOIN   regions USING(region_id)
              WHERE  region_name = 'Europe')
VALUES (3300, 'Cardiff', 'UK');
```

1 rows inserted

Inserting by Using a Subquery as a Target

Verify the results.


```
SELECT location_id, city, country_id  
FROM    loc
```

	LOCATION_ID	CITY	COUNTRY_ID
20	2900	Geneva	CH
21	3000	Bern	CH
22	3100	Utrecht	NL
23	3200	Mexico City	MX
24	3300	Cardiff	UK

Using the WITH CHECK OPTION Keyword on DML Statements

The WITH CHECK OPTION keyword prohibits you from changing rows that are not in the subquery.

```
INSERT INTO ( SELECT location_id, city, country_id
              FROM loc
              WHERE country_id IN
              (SELECT country_id
               FROM countries
               NATURAL JOIN regions
               WHERE region name = 'Europe')
              WITH CHECK OPTION )
VALUES (3600, 'Washington', 'US');
```



Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

Overview of the Explicit Default Feature

- Use the `DEFAULT` keyword as a column value where the default column value is desired.
- This allows the user to control where and when the default value should be applied to data.
- Explicit defaults can be used in `INSERT` and `UPDATE` statements.

Using Explicit Default Values

- DEFAULT with INSERT:

```
INSERT INTO deptm3  
  (department_id, department_name, manager_id)  
VALUES (300, 'Engineering', DEFAULT);
```

- DEFAULT with UPDATE:

```
UPDATE deptm3  
SET manager_id = DEFAULT  
WHERE department_id = 10;
```

Copying Rows from Another Table

- Write your INSERT statement with a subquery.

```
INSERT INTO sales_reps(id, name, salary, commission_pct)
  SELECT employee_id, last_name, salary, commission_pct
 FROM   employees
WHERE  job_id LIKE '%REP%';
```


33 rows inserted

- Do not use the VALUES clause.
- Match the number of columns in the INSERT clause with that in the subquery.



Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time


Overview of Multitable INSERT Statements




Sourcetab



Target_a



Target_b



Target_c

```
INSERT ALL  
  INTO target_a VALUES (...,...)  
  INTO target_b VALUES (...,...)  
  INTO target_c VALUES (...,...)  
  SELECT ...  
  FROM sourcetab  
  WHERE ...;
```

Subquery

Overview of Multitable `INSERT` Statements

- Use the `INSERT...SELECT` statement to insert rows into multiple tables as part of a single DML statement.
- Multitable `INSERT` statements are used in data warehousing systems to transfer data from one or more operational sources to a set of target tables.
- They provide significant performance improvement over:
 - Single DML versus multiple `INSERT...SELECT` statements
 - Single DML versus a procedure to perform multiple inserts by using the `IF . . . THEN` syntax

Types of Multitable INSERT Statements

The different types of multitable INSERT statements are:

- Unconditional INSERT
- Conditional INSERT ALL
- Pivoting INSERT
- Conditional INSERT FIRST

Multitable INSERT Statements

- Syntax for multitable INSERT:

```
INSERT [conditional_insert_clause]  
[insert_into_clause values_clause] (subquery)
```

- conditional_insert_clause:

```
[ALL|FIRST]  
[WHEN condition THEN] [insert_into_clause values_clause]  
[ELSE] [insert_into_clause values_clause]
```

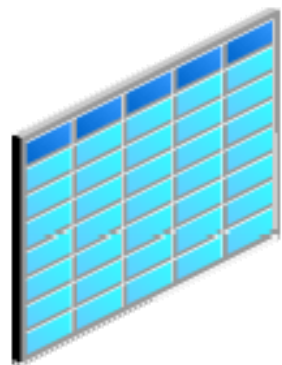
Unconditional INSERT ALL

- Select the EMPLOYEE_ID, HIRE_DATE, SALARY, and MANAGER_ID values from the EMPLOYEES table for those employees whose EMPLOYEE_ID is greater than 200.
- Insert these values into the SAL_HISTORY and MGR_HISTORY tables by using a multitable INSERT.

```
INSERT ALL
  INTO sal_history VALUES (EMPID, HIREDATE, SAL)
  INTO mgr_history VALUES (EMPID, MGR, SAL)
  SELECT employee_id EMPID, hire_date HIREDATE,
         salary SAL, manager_id MGR
  FROM   employees
 WHERE  employee_id > 200;
```

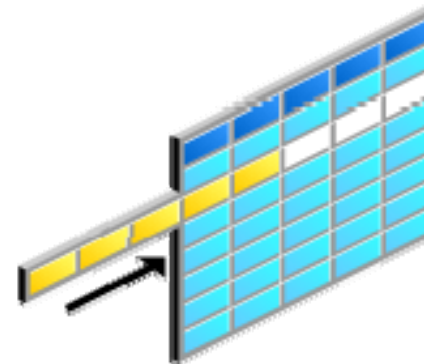

12 rows inserted

Conditional INSERT ALL: Example

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue.

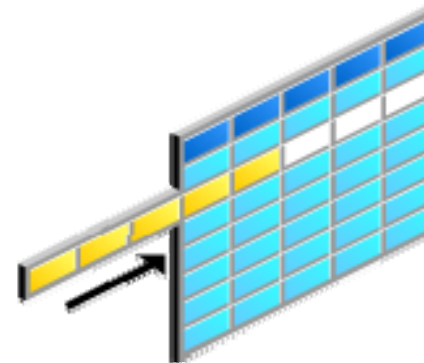
Employees

**Hired before
1995**

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue. A yellow bar highlights the first column, and a black arrow points to it.

EMP_HISTORY

**With sales
commission**

A 3D isometric icon of a table with 5 columns and 10 rows. The top row is dark blue, and the other rows are light blue. A yellow bar highlights the first column, and a black arrow points to it.

EMP_SALES

Conditional INSERT ALL

```
INSERT ALL
```

```
  WHEN HIREDATE < '01-JAN-95' THEN
```

```
    INTO emp_history VALUES (EMPID, HIREDATE, SAL)
```

```
  WHEN COMM IS NOT NULL THEN
```

```
    INTO emp_sales VALUES (EMPID, COMM, SAL)
```

```
  SELECT employee_id EMPID, hire_date HIREDATE,
```

```
         salary SAL, commission_pct COMM
```



```
  FROM employees
```

```
48 rows inserted
```



Conditional INSERT FIRST: Example

Scenario: If an employee salary is 2,000, the record is inserted into the SAL_LOW table only.

Salary < 5,000




SAL_LOW





EMPLOYEES

5000 <= Salary
<= 10,000



SAL_MID

Otherwise



SAL_HIGH

ORACLE

Conditional INSERT FIRST

```
INSERT FIRST
```

```
WHEN salary < 5000 THEN
```

```
    INTO sal_low VALUES (employee_id, last_name, salary)
```

```
WHEN salary between 5000 and 10000 THEN
```

```
    INTO sal_mid VALUES (employee_id, last_name, salary)
```

```
ELSE
```

```
    INTO sal_high VALUES (employee_id, last_name, salary)
```

```
SELECT employee_id, last_name, salary
```

```
FROM employees
```

```
107 rows inserted
```

Pivoting INSERT

Convert the set of sales records from the nonrelational database table to relational format.

Emp_ID	Week_ID	MON	TUES	WED	THUR	FRI
176	6	2000	3000	4000	5000	6000



Employee_ID	WEEK	SALES
176	6	2000
176	6	3000
176	6	4000
176	6	5000
176	6	6000

Pivoting INSERT

```
INSERT ALL
  INTO sales_info VALUES (employee_id, week_id, sales_MON)
  INTO sales_info VALUES (employee_id, week_id, sales_TUE)
  INTO sales_info VALUES (employee_id, week_id, sales_WED)
  INTO sales_info VALUES (employee_id, week_id, sales_THUR)
  INTO sales_info VALUES (employee_id, week_id, sales_FRI)
SELECT EMPLOYEE_ID, week_id, sales_MON, sales_TUE,
       sales_WED, sales_THUR, sales_FRI
FROM sales_source_data;
```

5 rows inserted

Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

MERGE Statement

- Provides the ability to conditionally update, insert, or delete data into a database table
- Performs an `UPDATE` if the row exists, and an `INSERT` if it is a new row:
 - Avoids separate updates
 - Increases performance and ease of use
 - Is useful in data warehousing applications

MERGE Statement Syntax

You can conditionally insert, update, or delete rows in a table by using the MERGE statement.

```
MERGE INTO table_name table_alias
  USING (table/view/sub_query) alias
  ON (join condition)
  WHEN MATCHED THEN
    UPDATE SET
      col1 = col1_val,
      col2 = col2_val
  WHEN NOT MATCHED THEN
    INSERT (column_list)
    VALUES (column_values);
```

Merging Rows: Example

Insert or update rows in the COPY_EMP3 table to match the EMPLOYEES table.

```
MERGE INTO copy_emp3 c
USING (SELECT * FROM EMPLOYEES ) e
ON (c.employee_id = e.employee_id)
WHEN MATCHED THEN
UPDATE SET
c.first_name = e.first_name,
c.last_name = e.last_name,
...
DELETE WHERE (E.COMMISSION_PCT IS NOT NULL)
WHEN NOT MATCHED THEN
INSERT VALUES(e.employee_id, e.first_name, e.last_name,
e.email, e.phone_number, e.hire_date, e.job_id,
e.salary, e.commission_pct, e.manager_id,
e.department_id);
```

Merging Rows: Example

```
TRUNCATE TABLE copy_emp3;  
SELECT * FROM copy_emp3;  
0 rows selected
```

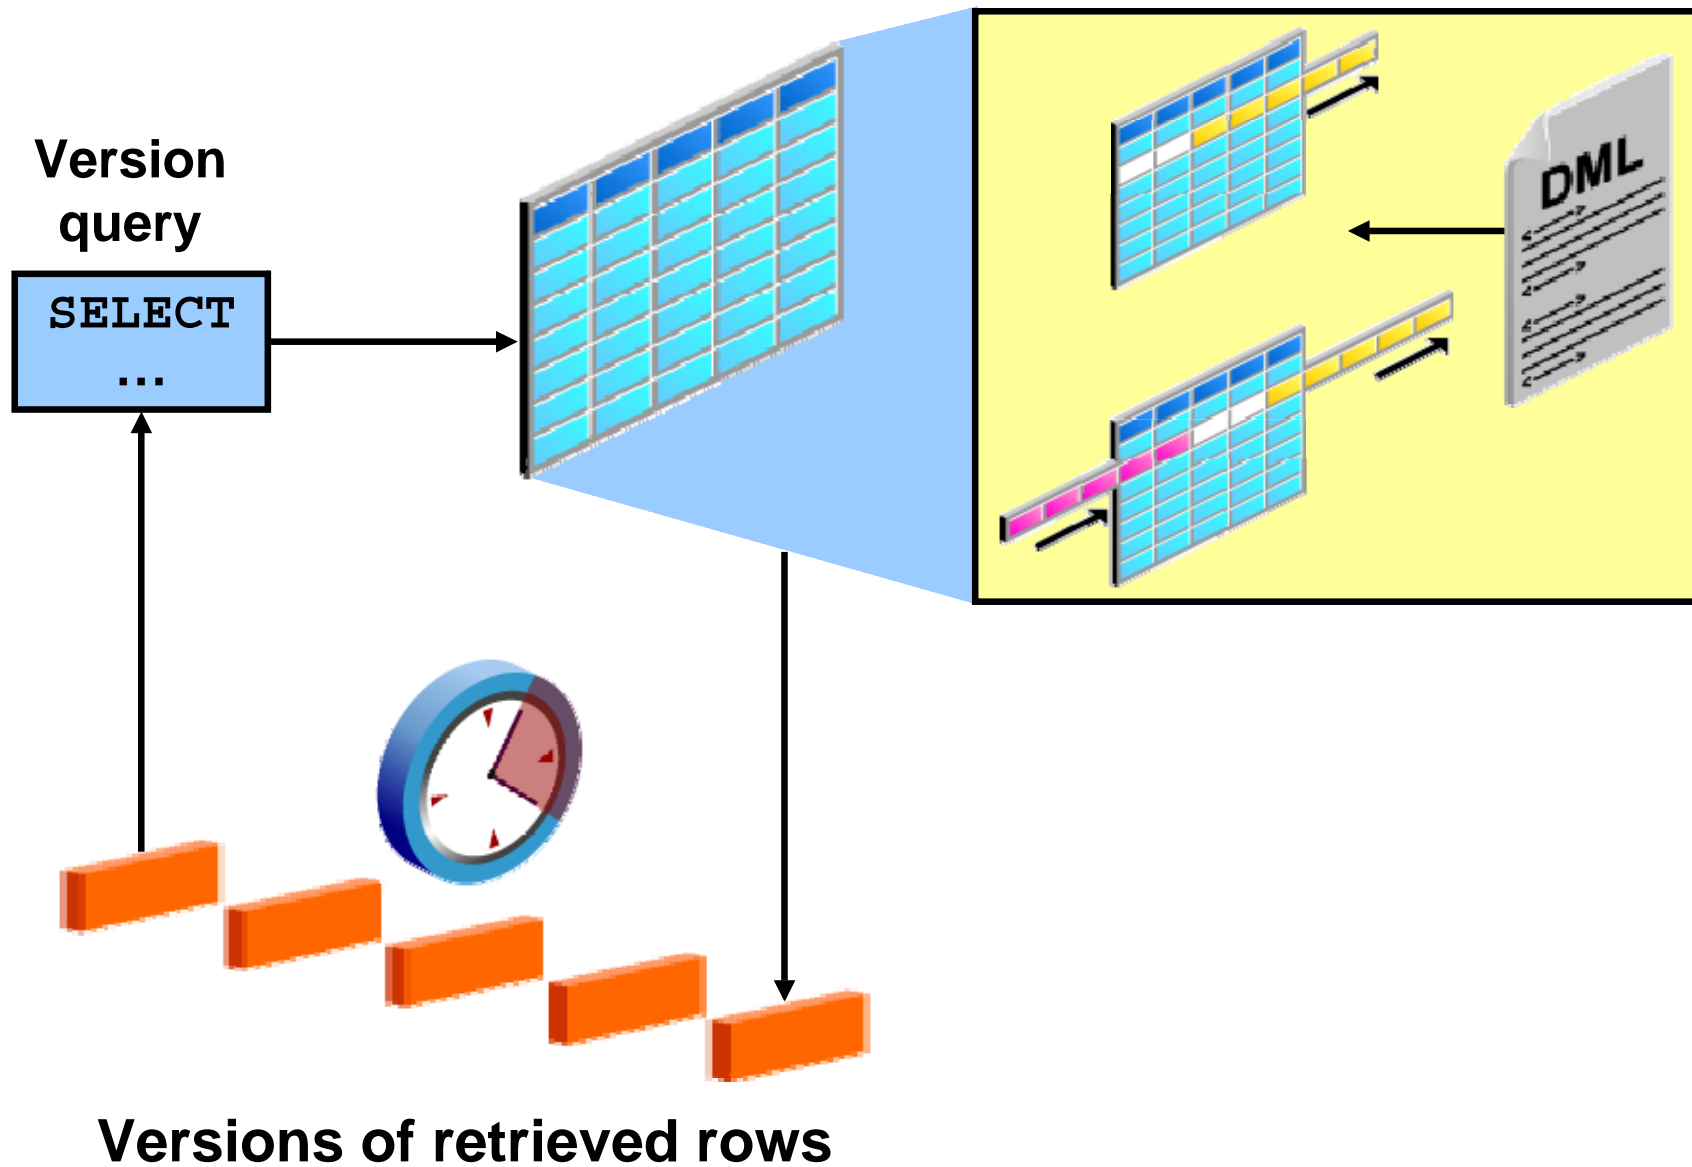
```
MERGE INTO copy_emp3 c  
USING (SELECT * FROM EMPLOYEES ) e  
ON (c.employee_id = e.employee_id)  
WHEN MATCHED THEN  
UPDATE SET  
c.first_name = e.first_name,  
c.last_name = e.last_name,  
...  
DELETE WHERE (E.COMMISSION_PCT IS NOT NULL)  
WHEN NOT MATCHED THEN  
INSERT VALUES(e.employee_id, e.first_name, ...
```

```
SELECT * FROM copy_emp3;  
107 rows selected.
```

Lesson Agenda

- Manipulating data by using subqueries
- Specifying explicit default values in the `INSERT` and `UPDATE` statements
- Using the following types of multitable `INSERT`s:
 - Unconditional `INSERT`
 - Pivoting `INSERT`
 - Conditional `INSERT ALL`
 - Conditional `INSERT FIRST`
- Merging rows in a table
- Tracking the changes to data over a period of time

Tracking Changes in Data



Example of the Flashback Version Query

```
SELECT salary FROM employees3  
WHERE employee_id = 107;
```

1

```
UPDATE employees3 SET salary = salary * 1.30  
WHERE employee_id = 107;
```

2

```
COMMIT;
```

```
SELECT salary FROM employees3  
  VERSIONS BETWEEN SCN MINVALUE AND MAXVALUE  
WHERE employee_id = 107;
```

3

1

	SALARY
1	4200

3

	SALARY
1	5460
2	4200

VERSIONS BETWEEN Clause

```
SELECT versions_starttime "START_DATE",  
       versions_endtime   "END_DATE",  
       salary  
FROM   employees  
       VERSIONS BETWEEN SCN MINVALUE  
       AND MAXVALUE  
WHERE  last_name = 'Lorentz';
```

	START_DATE	END_DATE	SALARY
1	18-JUN-09 05.07.10.0000000000 PM	(null)	5460
2	(null)	18-JUN-09 05.07.10.0000000000 PM	4200

Quiz

When you use the `INSERT` or `UPDATE` command, the `DEFAULT` keyword saves you from hard-coding the default value in your programs or querying the dictionary to find it.

1. True
2. False

Summary

In this lesson, you should have learned how to:

- Use DML statements and control transactions
- Describe the features of multitable INSERTs
- Use the following types of multitable INSERTs:
 - Unconditional INSERT
 - Pivoting INSERT
 - Conditional INSERT ALL
 - Conditional INSERT FIRST
- Merge rows in a table
- Manipulate data by using subqueries
- Track the changes to data over a period of time

Practice 4: Overview

This practice covers the following topics:

- Performing multitable `INSERTs`
- Performing `MERGE` operations
- Tracking row versions